

# Real-Time OS Frameworks

**Sandeep D'souza**

Lecture #7

# Outline of Lecture

- Approaches limiting Real-time and Non-real-time Task Interactions
  - Compliant Kernel Approach
  - Dual/Thin Kernel Approach
- Approaches that integrate Real-time and Non-real-time tasks
  - Core Kernel Approach
  - Resource Kernel Approach

# Approaches to Real-Time Linux

- Approaches limiting Real-time and Non-real-time Task Interactions
  - Compliant Kernel Approach
    - LynxOS/Blue Cat Linux
  - Dual Kernel Approach
    - RTLinux/RTAI
- Approaches that integrate Real-time and Non-real-time tasks
  - Core Kernel Approach
    - Monta Vista Linux, TimeSys Linux
  - Resource Kernel Approach
    - Linux/RK

# Linux Internals: Scheduling

- **Schedulable Entities**

- Processes

- Real-Time Class: `SCHED_FIFO` or `SCHED_RR`
    - Time-Sharing Class: `SCHED_OTHER`

- Real-time processes have

- Application-defined priority
    - Higher priority than time-sharing processes

- **Non-Schedulable Entities**

- Interrupt Handlers

- Have priorities, and can be nested

- “Bottom Halves” & Task Queues

- Run on schedule, ret from system call, ret from interrupt

# Linux and Real-Time: Traditional Problems

- **Timer Granularity**
  - Many real-time tasks are driven by timer interrupts
  - In Standard Linux, the timer was set to expire at 10 ms intervals
    - Beginning to change with usage of high-resolution timer and timestamp counters
- **Scheduler Predictability**
  - The Linux scheduler used to keep tasks in an unsorted list
  - Requires a scan of all tasks to make a scheduling decision
  - Scales poorly as number of tasks increases, and is especially poor for real-time performance
- **Various subsystems NOT designed for real-time use**
  - Network protocol stack
  - Filesystem
  - Windows manager

# Approaches to Real-Time Linux

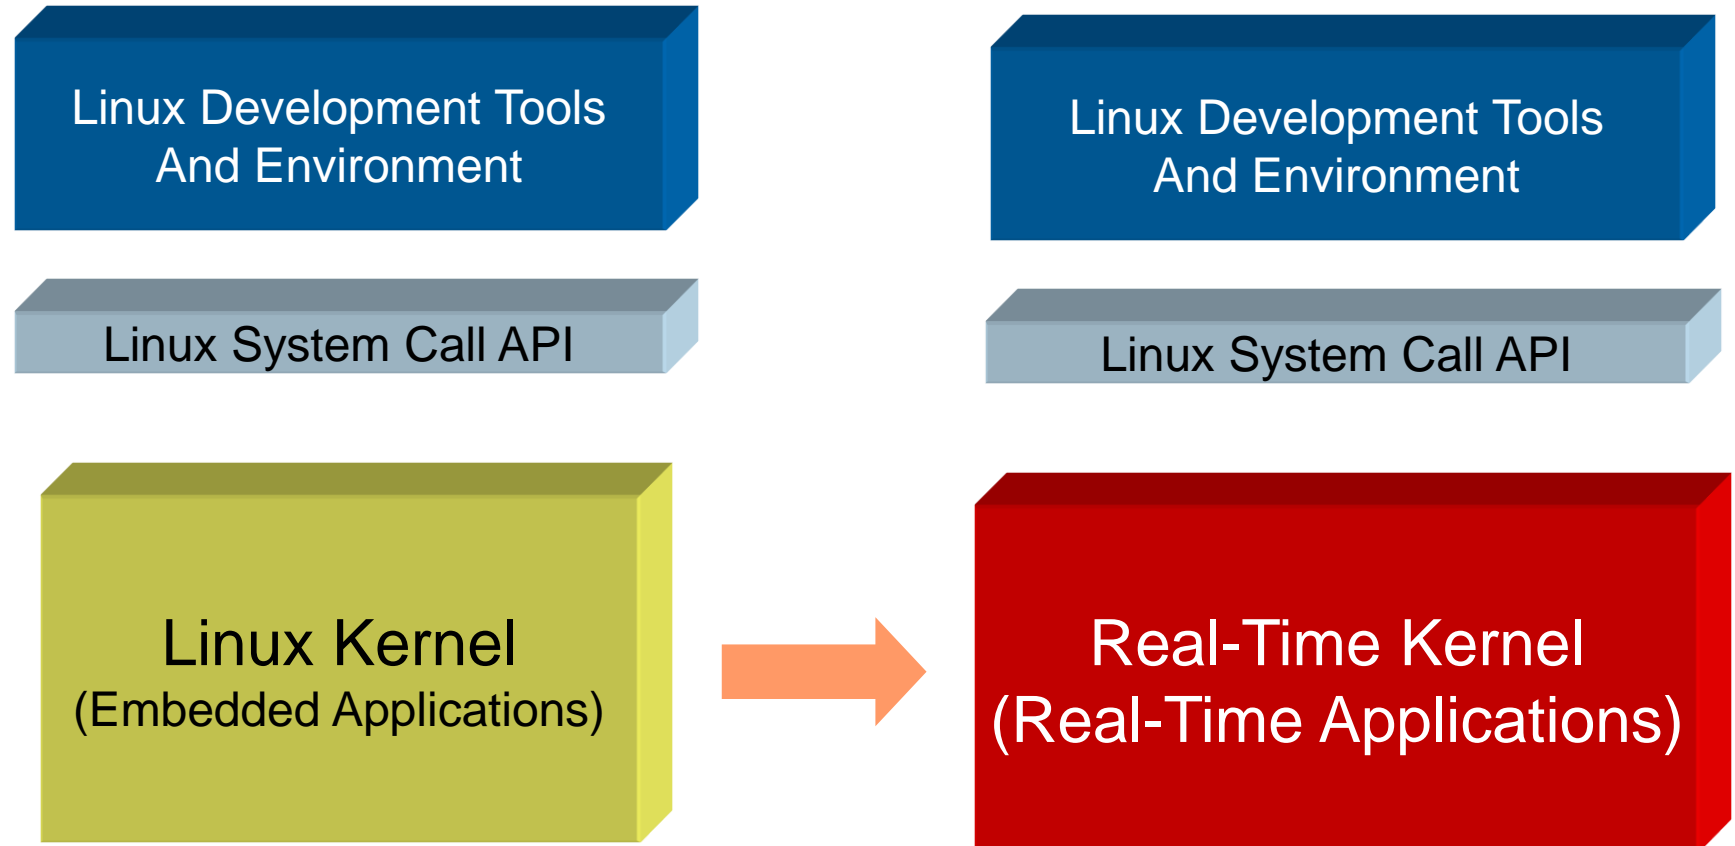
➤ Compliant Kernel Approach

➤ Dual Kernel Approach

➤ Core Kernel Approach

➤ Resource Kernel Approach

# Compliant Kernel Approach



# Compliant Kernel Approach

- Basic Claim

- Linux is defined by its API and not by its internal implementation
- The real-time kernel is a non-Linux kernel

- Implications

- No benefits from the Linux kernel
- Not possible to benefit from the Linux kernel evolution
- Not possible to use Linux hardware support
- Not (always) possible to use Linux device drivers



# Compliance

- 100% Linux API
  - Support all of Linux kernel API
- Implications
  - † Any Linux application can run on real-time kernel
    - Development can be done on a Linux host, with a rich set of host tools for development
  - † All Linux libraries are trivially available to run on a real-time kernel
    - Third-party software
  - Achieving 100% Linux API is non-trivial
    - Consider the amount of effort put into Linux kernel development

# Approaches to Real-Time Linux

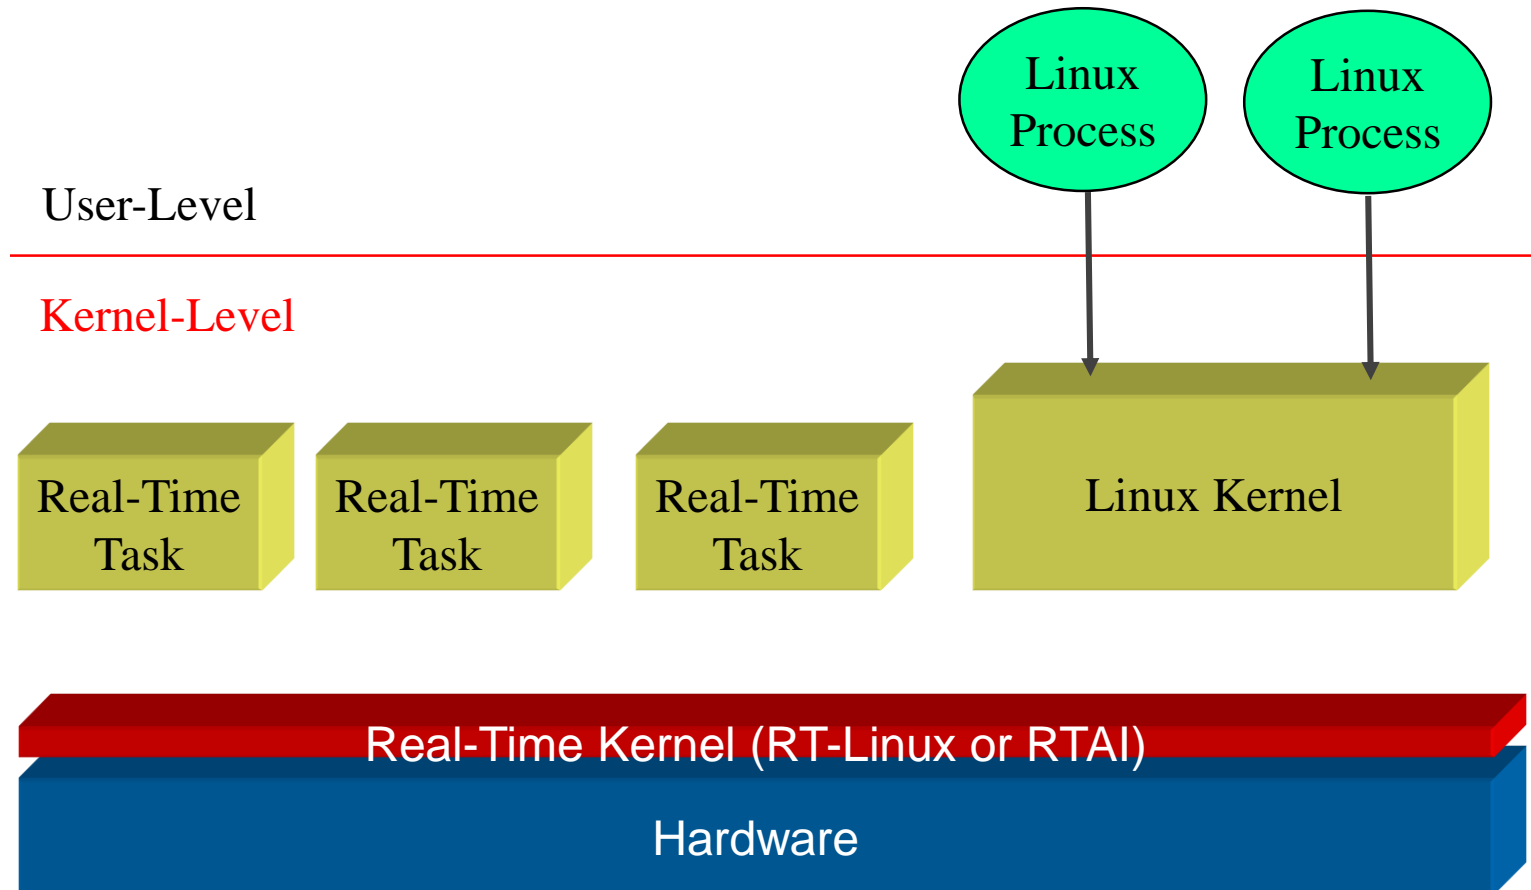
➤ Compliant Kernel Approach

➤ Dual Kernel Approach

➤ Core Kernel Approach

➤ Resource Kernel Approach

# The Thin Kernel Approach



- Real-time tasks do NOT use the Linux API or Linux facilities
- Failure in any real-time task crashes the entire system

# Approaches to Real-Time Linux

➤ Compliant Kernel Approach

➤ Dual Kernel Approach

➤ Core Kernel Approach

➤ Resource Kernel Approach

# Core Kernel Approach

- Basic Ideas

- Make the kernel more suitable for real-time
- Ensure that the impact of changes is localized so that
  - Kernel upgrades can be easily incorporated
  - Kernel reliability and scalability is not compromised

- Mechanisms

- Static Configuration
  - Can be configured at compile time
- Dynamic Configuration
  - Using loadable kernel modules

# Core Kernel Approach

- Allows the use of most, if not all, existing Linux primitives, applications, and tools.
  - Need to avoid primitives that can take extended time in the kernel
- Allows the use of most existing device drivers written to support Linux.
  - Need to avoid poorly written drivers that unfairly hog system resources
- Robustness and Reliability
  - Core kernel modifications can affect robustness, but source is available and extensive testing can be done.

# Approaches to Real-Time Linux

➤ Compliant Kernel Approach

➤ Dual Kernel Approach

➤ Core Kernel Approach

➤ Resource Kernel Approach

# Resource Kernel

- A Kernel that provides to Applications Timely, Guaranteed, and Enforced access to System Resources
- Allows Applications to specify only their Resource Demands, leaving the Kernel to satisfy those Demands using hidden management schemes

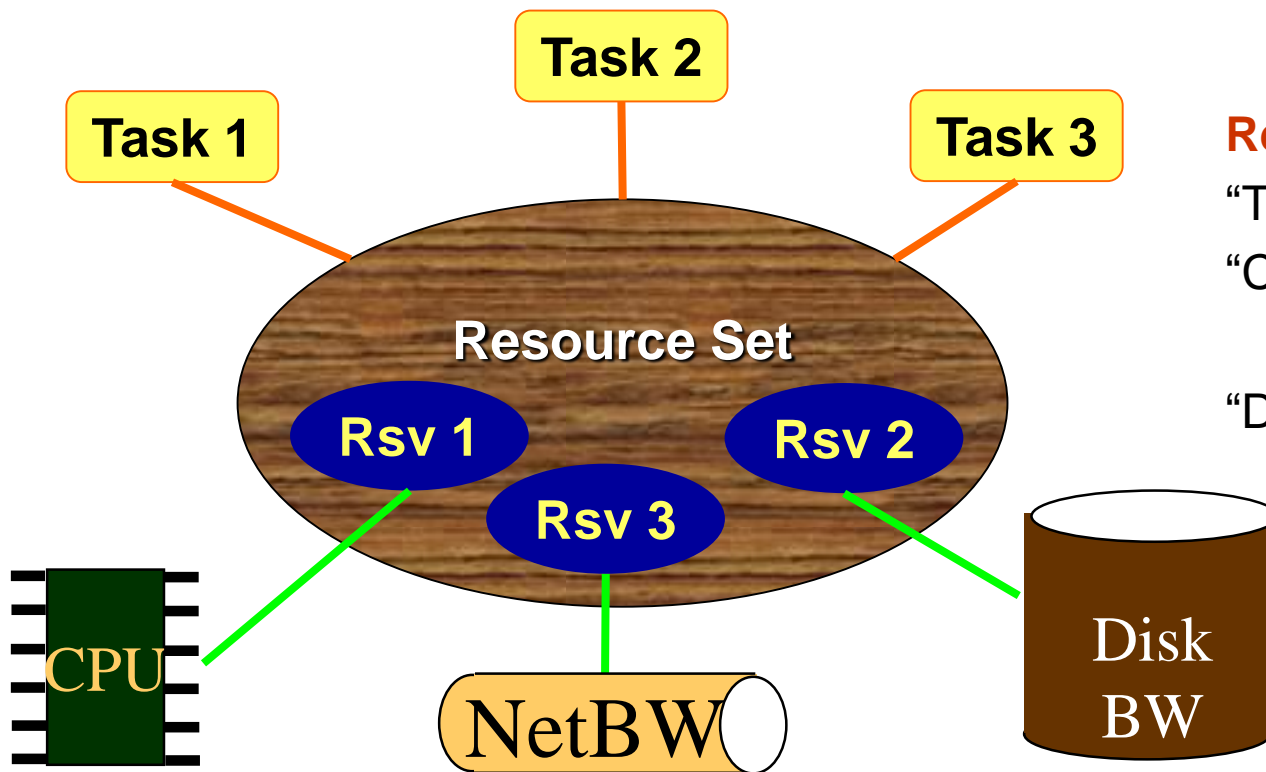


# Protection in Resource Kernels

- Each application (or a group of collaborating applications) operates in a “virtual machine”:
  - a machine which consists of a well-defined and guaranteed portion of system resources
    - CPU capacity, disk bandwidth, network bandwidth, and memory resource
- Multiple virtual machines can run simultaneously on the same physical machine
  - guarantees available to each resource set is valid despite the presence of other (potentially mis-behaving) applications using other resource sets

# Resource Kernel

- A Kernel that provides to applications Timely, Guaranteed, and Enforced access to System Resources
- Allows Applications to specify only their Resource Demands
  - leaving the Kernel to satisfy those Demands using hidden management schemes



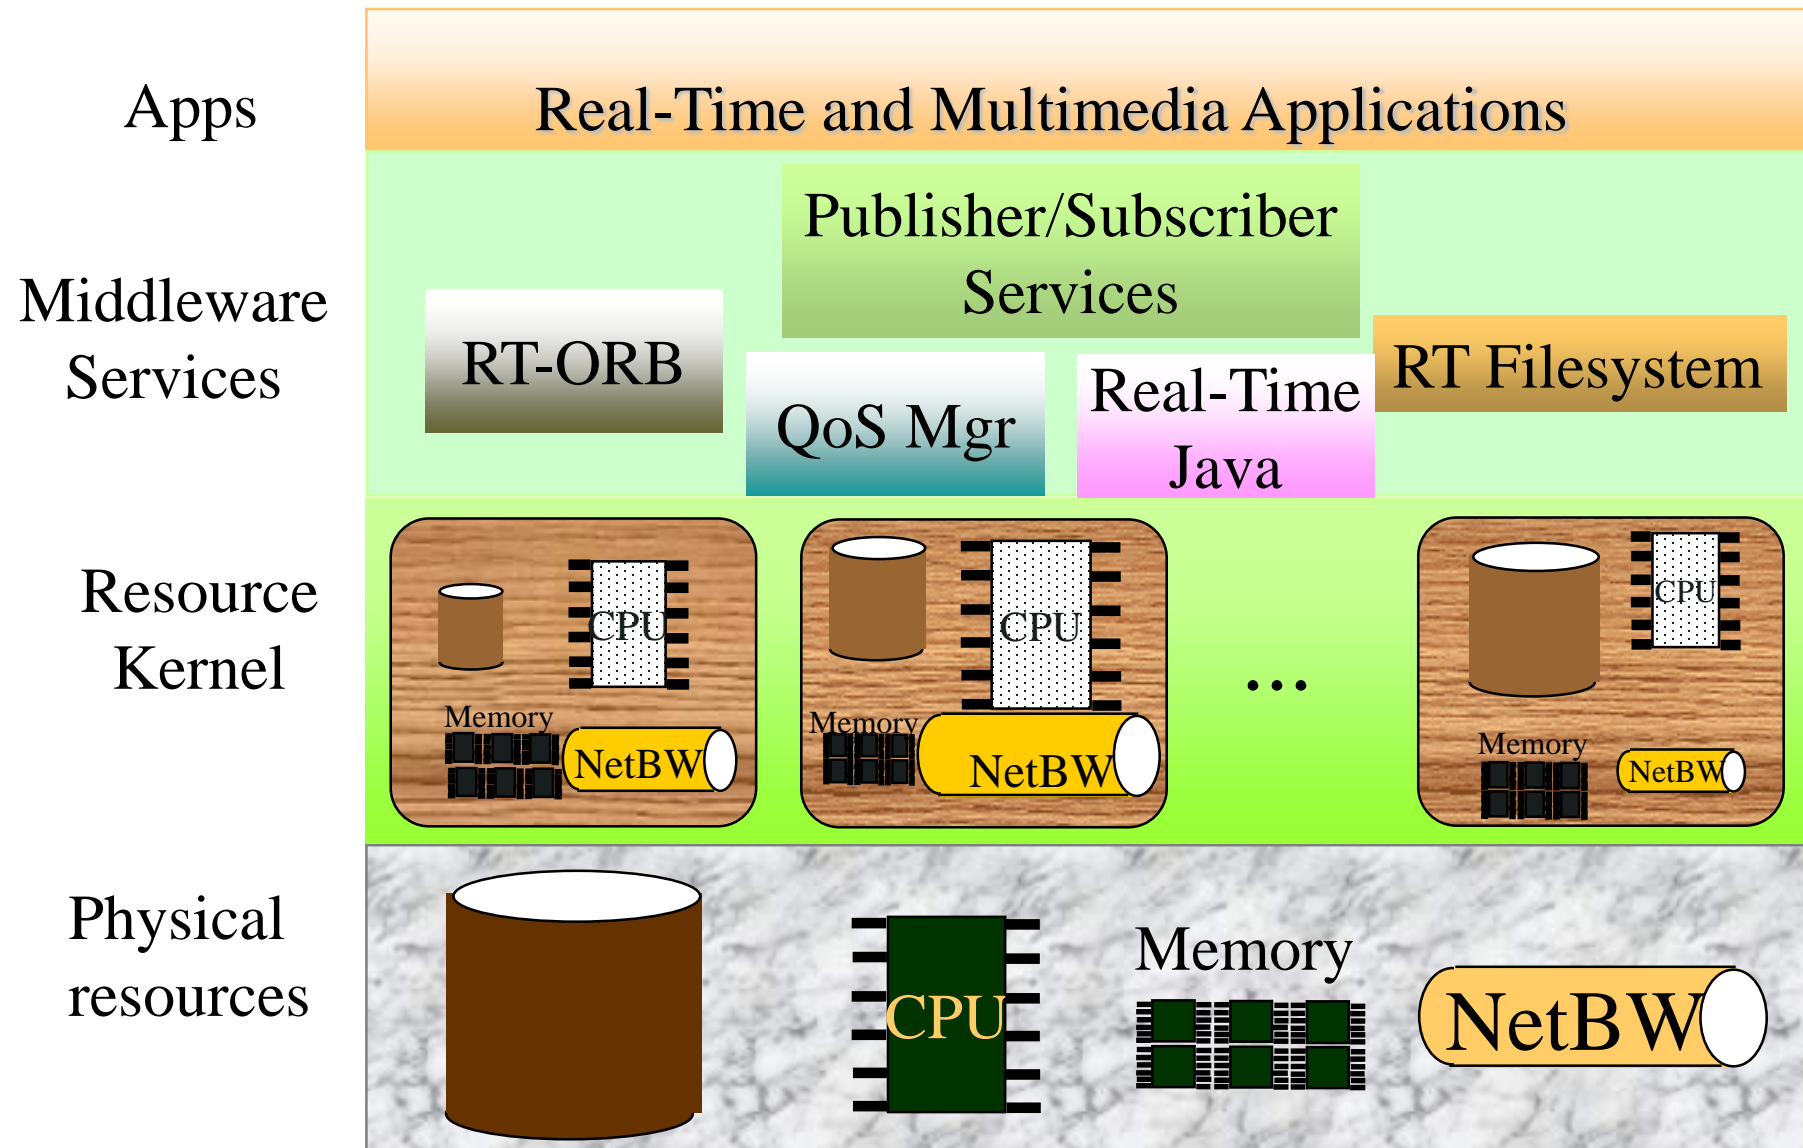
## Reservation Parameters

"T": Period ( $1/f$ )

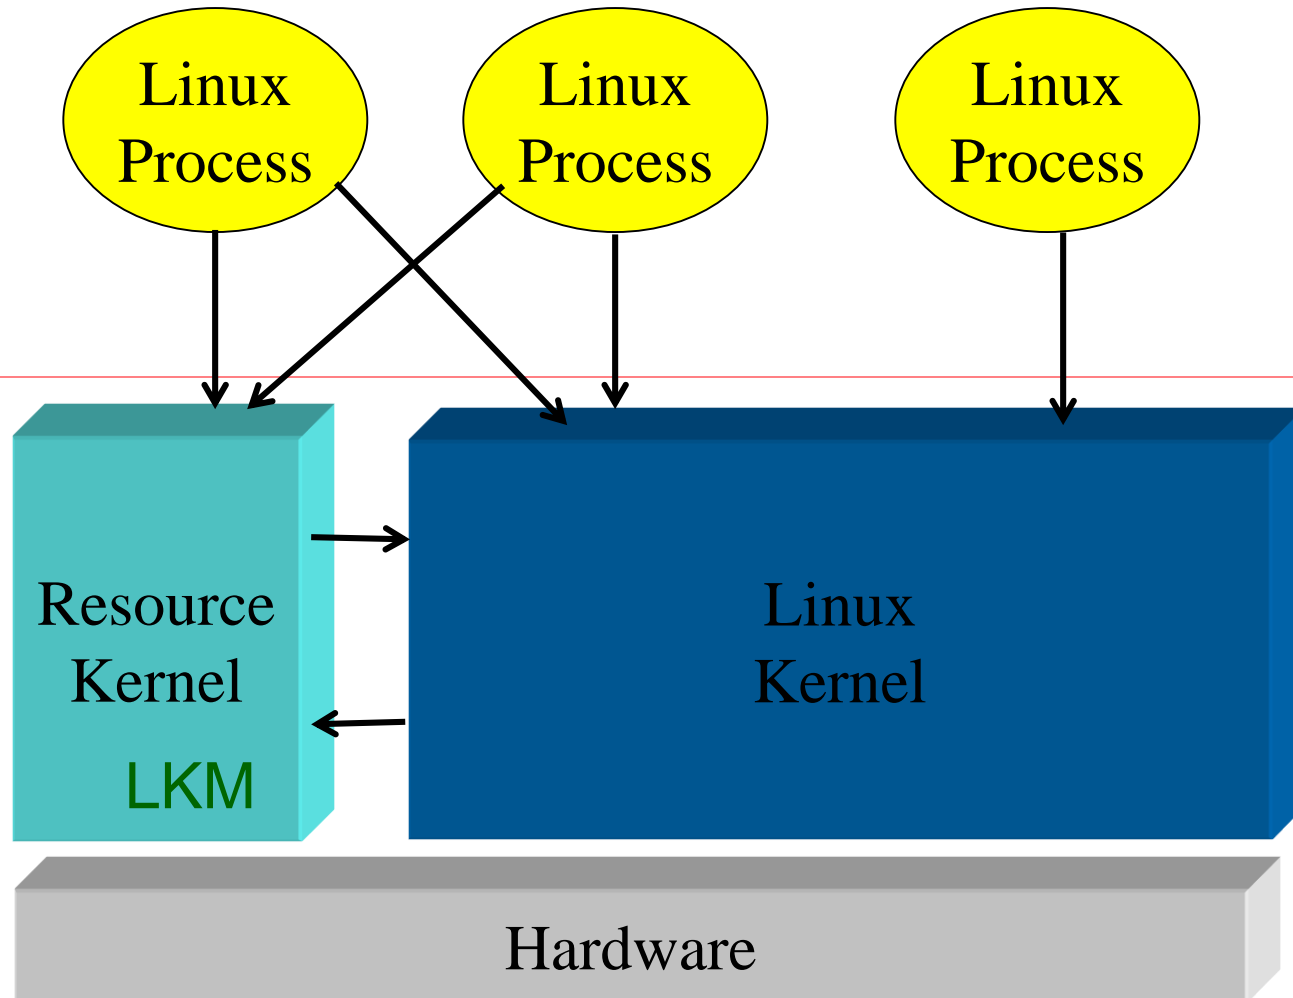
"C": Execution time within period

"D": Deadline within period

# “Resource Kernel” Architecture



# Linux Resource Kernel Architecture



# Reserves and Resource Sets

- Reserve

- A Share of a Single Resource
- Temporal Reserves
  - Parameters declare Portion and Timeframe of Resource Usage
    - E.g., CPU time, link bandwidth, disk bandwidth
- Spatial Reserves
  - Amount of space
    - E.g., memory pages, network buffers

- Resource Set

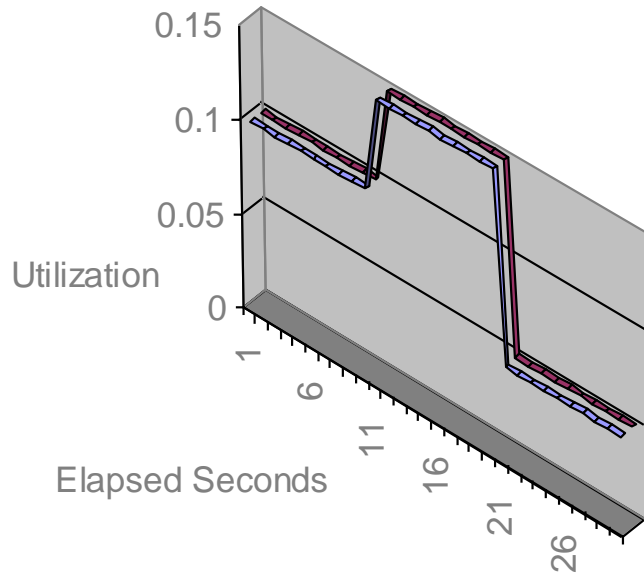
- A set of resource reserves
- Zero, one or more processes can be bound to a resource set.

# Linux/RK Abstractions

Linux/RK supports several abstractions and primitives for real-time scheduling of processes with real-time and QoS requirements:

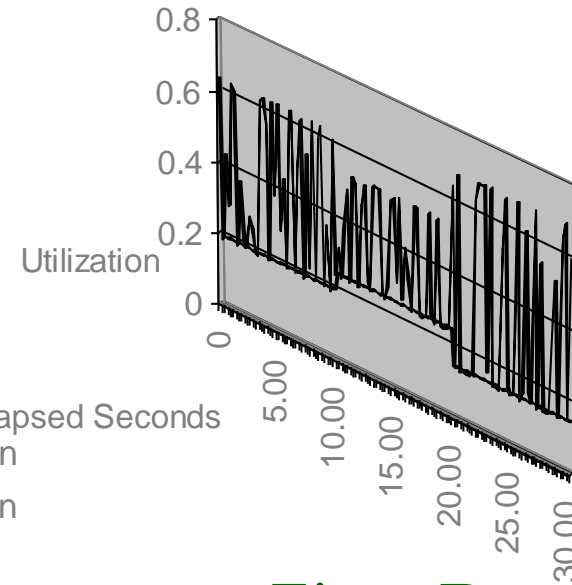
- **Resource reservations with latency guarantees**
  - CPU cycles
  - Network bandwidth
  - Disk bandwidth
- Support for **periodic tasks**.
- Support for 256 real-time fixed-priority levels.
- **High-resolution timers and clocks**.
- **Bounding of priority inversion** during synchronization operations.
- Wiring down of memory pages.

# Reservation Types



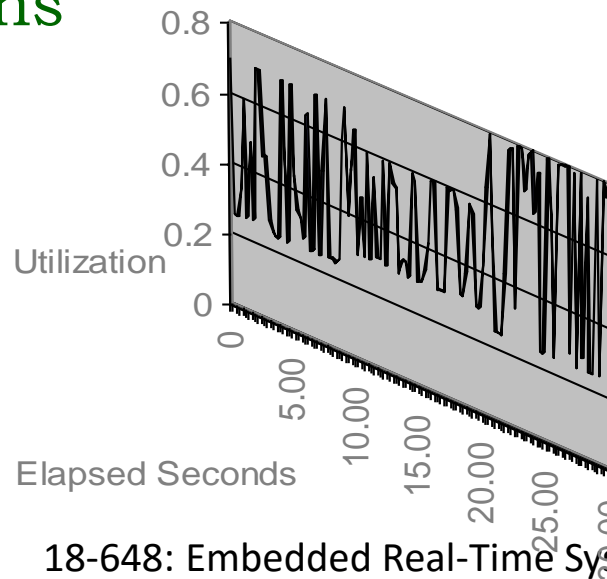
## Hard Reservations

(guarantees with **No** extras even if resource is idle)



## Firm Reservations

(guarantees with Extras only if no non-real-time)



**Soft Reservations**  
(guarantees with extras)

W/O Competition



# Performance Overhead of *Hierarchical* Reservations

- **Linear with the height of the hierarchy**
  - reservation replenishment and enforcement
- **Constant**
  - admission control (only *local* schedulability analysis)
  - scheduling (internal priority mapping and disabling/re-enabling of process eligibility to be scheduled).
- **Constraints**
  - reservation period must be greater than twice the parent's reservation period
- **Hidden overhead**
  - Higher degree of interrupts, because of replenishment, enforcement timers going off more frequently.



# Degrees of Temporal Isolation

- Different degrees of temporal isolation in the presence of resource-sharing
  - **Strict Isolation:** the timing behavior of an application is not affected by the timing misbehavior of *any* other application
    - RK applications in the absence of logical sharing of resources
  - **Non-Strict Isolation:** traditional priority-driven systems
  - **Weak Isolation:** timing behavior is not affected by the timing misbehavior of applications with which no logical resources are shared.



# Resource-Sharing Protocols in RK

- **Analogues to Priority Inheritance and Priority Ceiling Protocols** in Resource kernels
- Temporal isolation can only be weak
  - under logical resource-sharing using mutexes and client-server architectures
  - timeout and restart schemes may need to be applied.

# Priority Ceiling Protocol Equivalents

- **Single-Reserve PCP**


- Assign one reservation to the logical resource execution
- Very pessimistic allocation is required to maintain PCP semantics

- **Multi-Reserve PCP**

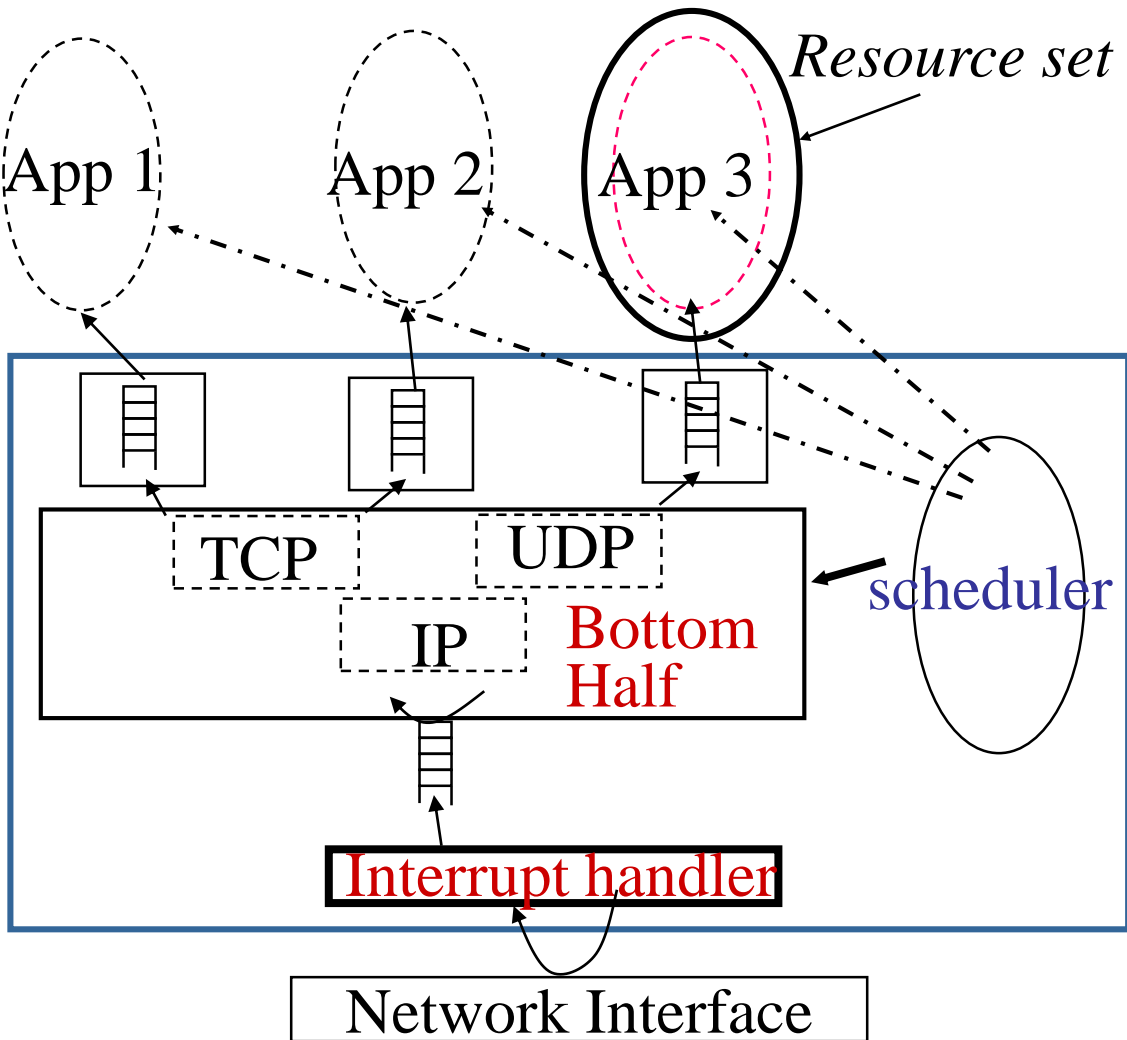
- Has the same schedulability analysis as traditional PCP
- Requires special support in RK
- Can be applied to client-server models
  - Pass client's reserve to server along with request - charges

# Priority Inheritance Protocol Equivalents

- **Priority Inheritance (PI)**
  - Server runs at the priority of the highest priority client waiting for server.
- **Priority Inheritance with Priority Inversion Enforcement (PIPIE)**
  - Enforce the duration of priority inversion encountered by any task
    - Can never exceed the amount specified at admission control
    - Need to track multiple tasks' priority inversions simultaneously
- **Reserve Inheritance (RI)**
  - Server usage is charged to client's reservation + inherit the highest priority of any client waiting for server
- **PIPIE + RI**




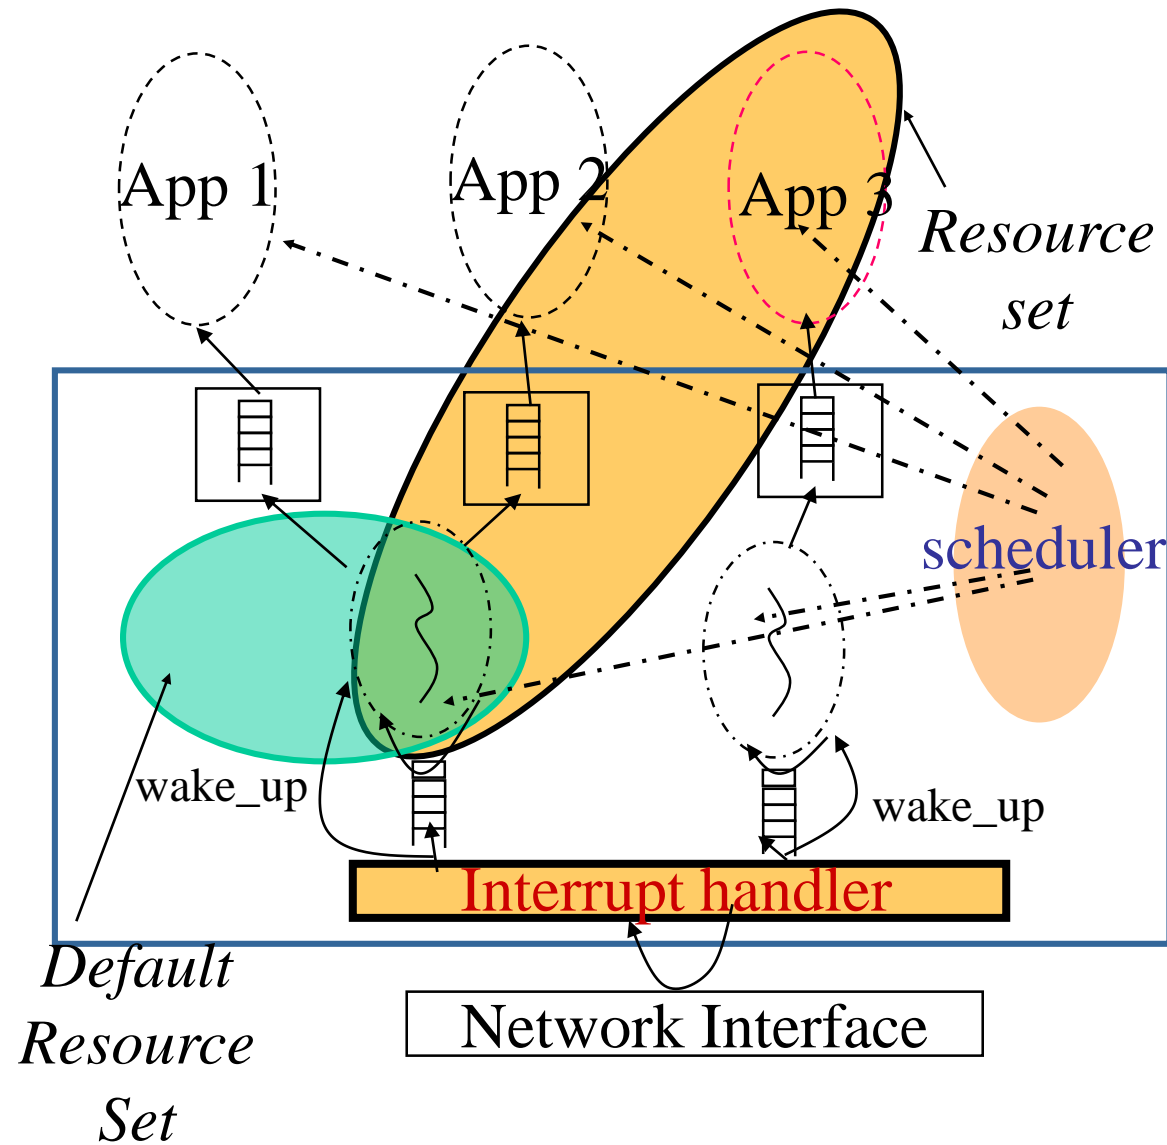
# Network Processing Architecture



## Problems

- Non-preemptive System calls
  - Priority Inversion
- “Eager” Receiver Processing
  - Interrupt-driven
    - Priority(capturing packet) > Priority(protocol processing) > Priority(application processing)
- Lack of effective load shedding
- Lack of traffic separation
- Inappropriate resource accounting

# Threaded Network Processing



Application domain extends into the kernel, and its activity is controlled

# Network and CPU Service Guarantees

- Reduction in non-preemptibility
- Control of receiver overload (*receive -livelock*)
- Prevention of scheduling disruption
- Separation of individual flows and proper resource accounting
- Packet scheduling for QoS (Quality of Service) guarantees

# Summary

- OS Approaches limiting Real-time and Non-real-time Task Interactions
  - Compliant Kernel Approach
  - Thin Kernel Approach
- OS Approaches that integrate Real-time and Non-real-time tasks
  - Core Kernel Approach
  - Resource Kernel Approach