

# Interprocess communication in UNIX

## Pipes

- The `pipe()` function:

```
#include <unistd.h>

int pipe(int fd[2]);

Returns: 0 if OK, -1 on error
```

- After calling `pipe()`:




Figure 15.2, APUE

- After calling `pipe()` and then `fork()`:




Figure 15.3, APUE

- Example: copying file to a pager program

```
#include "apue.h"
#include <sys/wait.h>

#define DEF_PAGER "/bin/more" /* default pager program */

int main(int argc, char *argv[])
{
    int          n;
    int          fd[2];
    pid_t        pid;
    char         *pager, *argv0;
    char         line[MAXLINE];
    FILE         *fp;

    if (argc != 2)
        err_quit("usage: a.out <pathname>");

    if ((fp = fopen(argv[1], "r")) == NULL)
        err_sys("can't open %s", argv[1]);
    if (pipe(fd) < 0)
        err_sys("pipe error");

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid > 0) {        /* parent */

        close(fd[0]);           /* close read end */

        /* parent copies argv[1] to pipe */
        while (fgets(line, MAXLINE, fp) != NULL) {
            n = strlen(line);
            if (write(fd[1], line, n) != n)
                err_sys("write error to pipe");
        }
        if (ferror(fp))
            err_sys("fgets error");

        close(fd[1]);           /* close write end of pipe for reader */

        if (waitpid(pid, NULL, 0) < 0)
            err_sys("waitpid error");
        exit(0);
    } else {                    /* child */

        close(fd[1]);           /* close write end */

        if (fd[0] != STDIN_FILENO) {
            if (dup2(fd[0], STDIN_FILENO) != STDIN_FILENO)
                err_sys("dup2 error to stdin");
        }
    }
}
```

```

        close(fd[0]);        /* don't need this after dup2 */
    }

    /* get arguments for execl() */
    if ((pager = getenv("PAGER")) == NULL)
        pager = DEF_PAGER;
    if ((argv0 = strrchr(pager, '/')) != NULL)
        argv0++;             /* step past rightmost slash */
    else
        argv0 = pager;       /* no slash in pager */

    if (execl(pager, argv0, (char *)0) < 0)
        err_sys("execl error for %s", pager);
}
exit(0);
}

```

## XSI IPC

They share common naming and interface scheme:

- XSI Message queues

```

int msgget(key_t key, int flag);
int msgctl(int msqid, int cmd, struct msqid_ds *buf);
int msgsnd(int msqid, const void *ptr, size_t nbytes, int flag);
ssize_t msgrcv(int msqid, void *ptr, size_t nbytes, long type, int flag);

```

- XSI Semaphores

```

int semget(key_t key, int nsems, int flag);
int semctl(int semid, int semnum, int cmd, ... /* union semun arg */ );
int semop(int semid, struct sembuf semoparray[], size_t nops);

```

- XSI Shared memory

```

int shmget(key_t key, size_t size, int flag);
int shmctl(int shmid, int cmd, struct shmid_ds *buf);
void *shmat(int shmid, const void *addr, int flag);
int shmdt(const void *addr);

```

And they all suck...

1. Hard to clean-up because there is no reference counting
  - pipes get automatically removed when last process terminates
  - data left in a FIFO is removed when last process terminates

## 2. Hard to use

- complex and inelegant interfaces that don't fit into UNIX file system paradigm
- stupid naming scheme: IPC identifiers, keys, and *project IDs* – *are you serious?*

They have been widely used for lack of alternatives. Fortunately we do have alternatives these days:

- Instead of XSI message queues, use:
  - UNIX domain sockets
  - POSIX message queues (still not widely available, so not covered in APUE; see `man 7 mq_overview`)
- Instead of XSI semaphores, use:
  - POSIX semaphores
- Instead of XSI shared memory, use:
  - memory mapping using `mmap()`

## Memory-mapped I/O

---

- `mmap()` function:

```
#include <sys/mman.h>
```

```
void *mmap(void *addr, size_t len, int prot, int flag, int fd, off_t off);
```

Returns: starting address of mapped region if OK, MAP\_FAILED on error

- Example of a memory-mapped file:

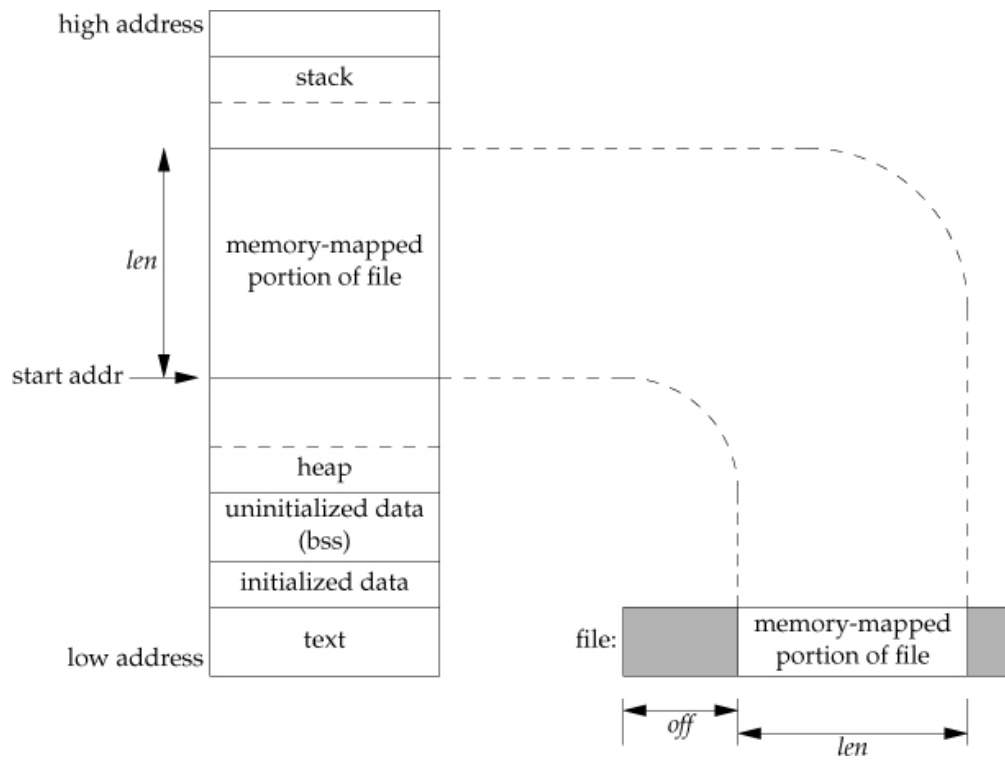


Figure 14.26, APUE

- Memory mapping `/dev/zero` for shared memory (Figure 15.33, APUE):

```

#include "apue.h"
#include <fcntl.h>
#include <sys/mman.h>

#define NLOOPS          1000
#define SIZE            sizeof(long)    /* size of shared memory area */

static int
update(long *ptr)
{
    return((*ptr)++);    /* return value before increment */
}

int
main(void)
{
    int          fd, i, counter;
    pid_t        pid;
    void          *area;

    if ((fd = open("/dev/zero", O_RDWR)) < 0)
        err_sys("open error");
    if ((area = mmap(0, SIZE, PROT_READ | PROT_WRITE, MAP_SHARED, fd, 0)) ==
        err_sys("mmap error");
    close(fd);    /* can close /dev/zero now that it's mapped */

    TELL_WAIT();

    if ((pid = fork()) < 0) {
        err_sys("fork error");
    } else if (pid > 0) {    /* parent */
        for (i = 0; i < NLOOPS; i += 2) {
            if ((counter = update((long *)area)) != i)
                err_quit("parent: expected %d, got %d", i, counter);

            TELL_CHILD(pid);
            WAIT_CHILD();
        }
    } else {    /* child */
        for (i = 1; i < NLOOPS + 1; i += 2) {
            WAIT_PARENT();

            if ((counter = update((long *)area)) != i)
                err_quit("child: expected %d, got %d", i, counter);

            TELL_PARENT(getppid());
        }
    }

    exit(0);
}

```

- Anonymous memory mapping

Same as `/dev/zero` mapping, but more portable and more convenient.

Change Figure 15.33 as follows:

1. remove `open("/dev/zero", ...)` and `close(fd)`
2. change `mmap` call to:

```
if ((area = mmap(0, SIZE, PROT_READ | PROT_WRITE,
                 MAP_ANON | MAP_SHARED, -1, 0)) == MAP_FAILED)
```

## POSIX Semaphores

- What is semaphore?
  - Binary vs. Counting semaphores
- Creating, opening, closing, and removing **named** POSIX semaphores:

```
#include <semaphore.h>

sem_t *sem_open(const char *name, int oflag, ... /* mode_t mode,
              unsigned int value */ );
    Returns: Pointer to semaphore if OK, SEM_FAILED on error

int sem_close(sem_t *sem);
    Returns: 0 if OK, -1 on error

int sem_unlink(const char *name);
    Returns: 0 if OK, -1 on error
```

- Initializing and destroying **unnamed** POSIX semaphores:

```
#include <semaphore.h>

int sem_init(sem_t *sem, int pshared, unsigned int value);
    Returns: 0 if OK, -1 on error

int sem_destroy(sem_t *sem);
    Returns: 0 if OK, -1 on error
```

- Using POSIX semaphores:

Decrement the value of semaphores:

```
#include <semaphore.h>

int sem_trywait(sem_t *sem);
int sem_wait(sem_t *sem);
    Both return: 0 if OK, -1 on error
```

Decrement with bounded waiting:

```
#include <semaphore.h>
#include <time.h>

int sem_timedwait(sem_t *restrict sem,
                  const struct timespec *restrict tsPtr);
    Returns: 0 if OK, -1 on error
```

Increment the value of semaphores:

```
#include <semaphore.h>

int sem_post(sem_t *sem);
    Returns: 0 if OK, -1 on error
```

---

*Last updated: 2014-02-11*