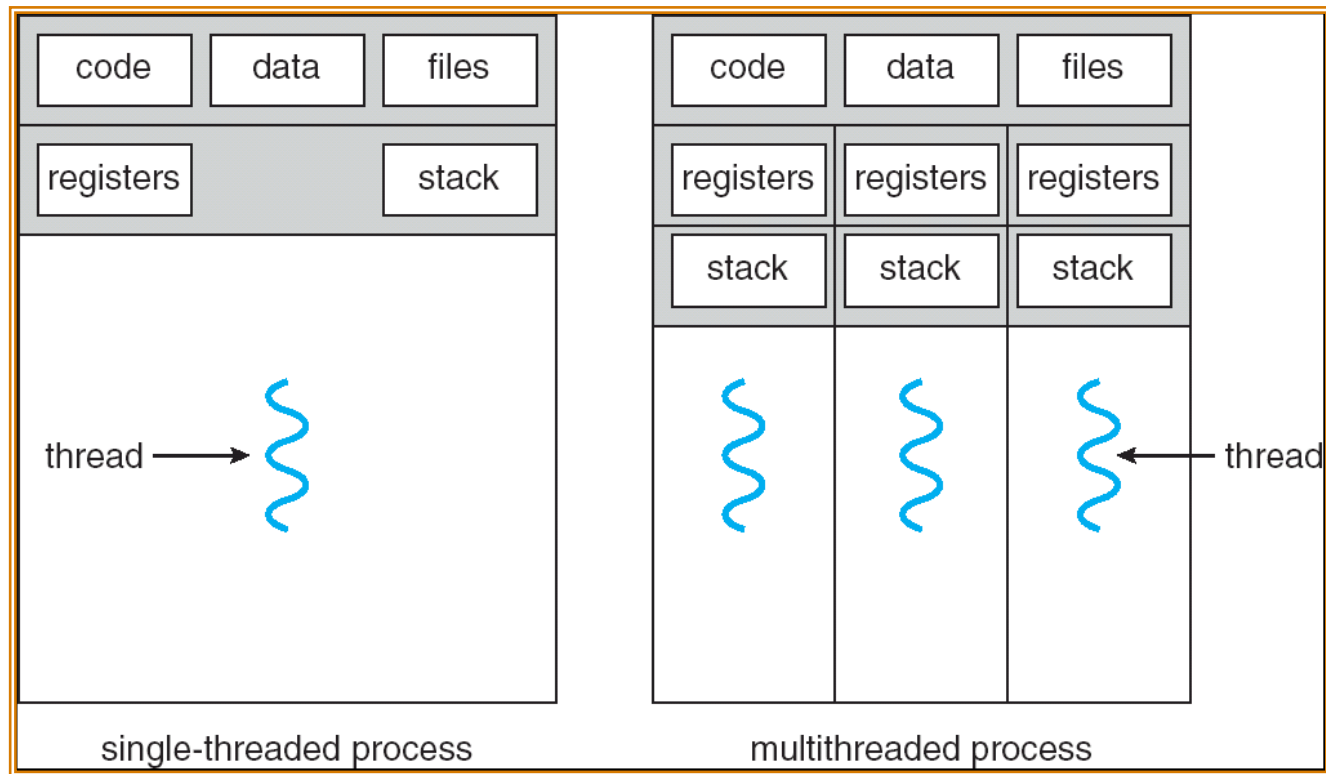


Single and multithreaded processes



Why threads?

- Express **concurrency**
 - Web server (multiple requests), Browser (GUI + network I/O + rendering), most GUI programs ...

```
for(;;) {  
    struct request *req = get_request();  
    create_thread(process_request, req);  
}
```
- **Efficient** communication
 - Using a separate process for each task can be heavyweight
- Leverage **multiple cores** (depends)
 - Unthreaded process can only run on a single CPU


Threads vs. Processes

- A thread has no data segment or heap
 - A thread cannot live on its own, it must live within a process
 - There can be more than one thread in a process, the first thread calls `main()` & has the process's stack
 - Inexpensive creation
 - Inexpensive context switching
 - Efficient communication
 - If a thread dies, its stack is reclaimed
- A process has code/data/heap & other segments
 - A process has at least one thread
 - Threads within a process share code/data/heap, share I/O, but each has its own stack & registers
 - Expensive creation
 - Expensive context switching
 - Interprocess communication can be expensive
 - If a process dies, its resources are reclaimed & all threads die


Multithreading models

- Where to support threads?
- **User threads**: thread management done by user-level threads library; kernel knows nothing
- **Kernel threads**: threads directly supported by the kernel
 - Virtually all modern OS support kernel threads

User vs. Kernel Threads



E.g., GreenThreads, any OS
(event ancient ones like DOS)



E.g., LinuxThreads, Solaris

Example from Tanenbaum, Modern Operating Systems 3 e,
(c) 2008 Prentice-Hall, Inc. All rights reserved. 0-13-6006639

Scheduling User Threads

- Non-preemptive Scheduling
 - No timer to make a thread yield the CPU
 - Threads must voluntarily yield control to let another thread run, e.g., `pthread_yield()`
 - Thread history isn't taken into account by scheduler
 - Threads are *co-operative*, not competitive
- Preemptive Scheduling
 - Can use signals to simulate interrupts, e.g., alarm
 - But then user code can't use alarm directly

User Thread Blocking

- What happens when a process does a read()?
 - Data needs to be fetched from disk
 - Kernel **blocks** the process (i.e., doesn't return) until disk read is done
 - Kernel unaware of thread structure: all user level threads will block as well!
- One solution: wrapper functions
 - Thread library contains alternate versions of syscalls
 - Check for blocking **before** calling the kernel
 - E.g., select() before read()
 - If the call will block, then schedule another thread
 - Complex – need to handle **all** blocking calls!

User vs. Kernel Threads (cont.)

User

- Pros: fast, no system call for creation, context switch
- Cons: kernel doesn't know → one thread blocks, all threads in the process blocks
- Cons: can't benefit from multicore or multiple CPUs

Kernel

- Cons: slow, kernel does creation, scheduling, etc
- Pros: kernel knows → one thread blocks, schedule another
- Pros: can fully utilize multiple cores/CPU's

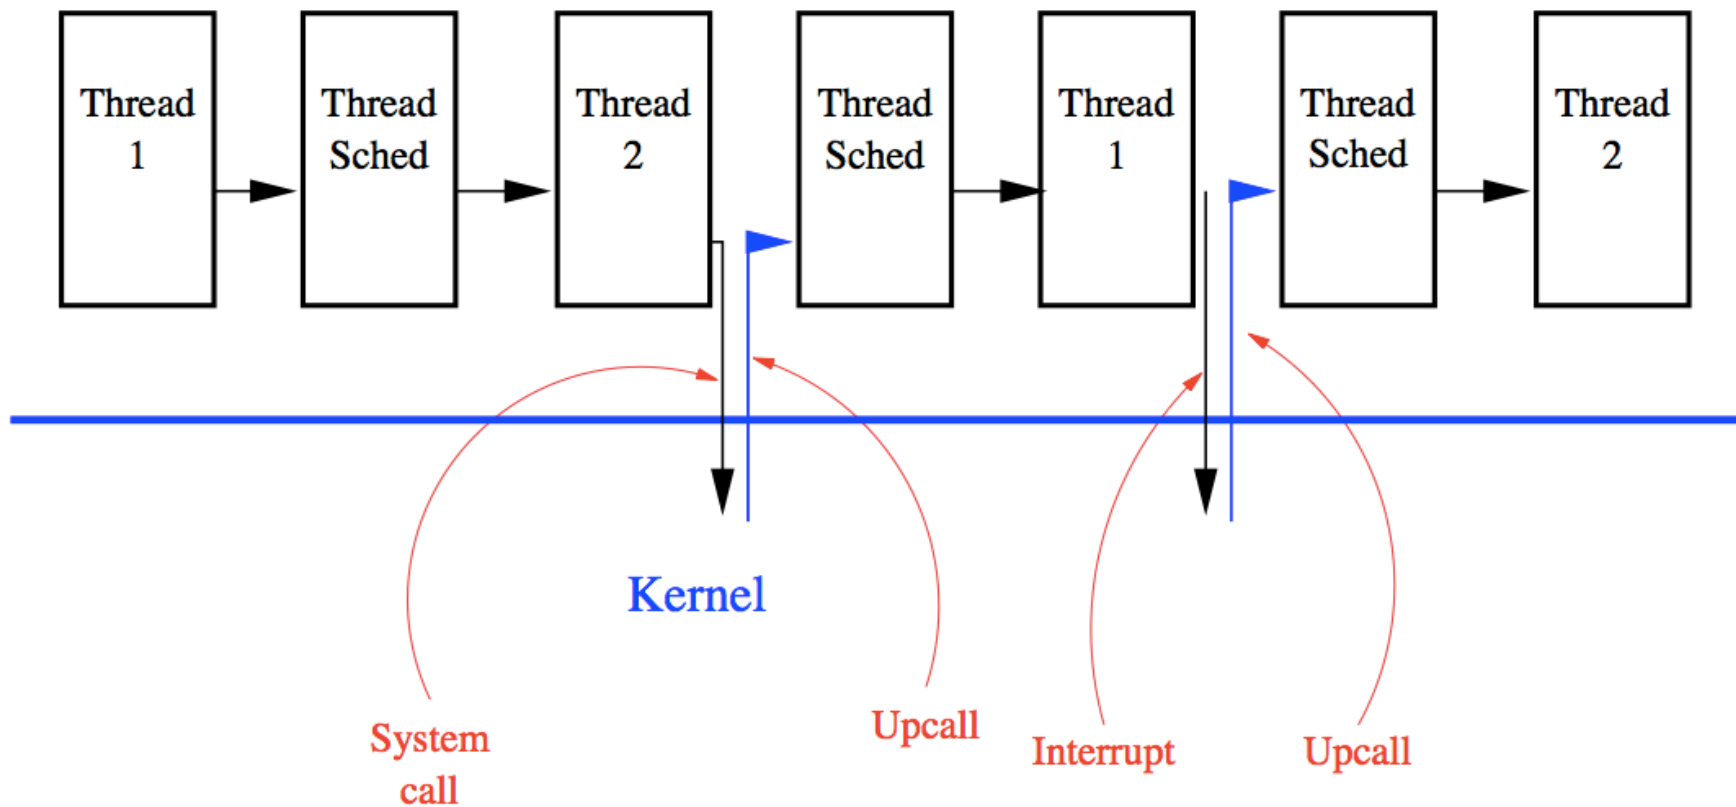
No free lunch!

Scheduler Activations

- Hybrid approach (Tru UNIX, NetBSD, some Mach, implementations for Linux)
 - Benefits of both user and kernel threads
 - Relies on **upcalls** (like signals)
- Scheduling done at user level
 - When a syscall is going to block, kernel informs user level thread manager via upcall
 - Thread manager can run another thread
 - When blocking call is done, kernel informs thread manager again

Reference: <http://homes.cs.washington.edu/~bershad/Papers/p53-anderson.pdf> (“Scheduler Activations: Effective Kernel Support for the User-Level Management of Parallelism”)

Scheduler Activations




Multiplexing User-Level Threads

- A thread library must map user threads to kernel threads
- Big picture:
 - kernel thread: **physical concurrency, how many cores?**
 - User thread: **application concurrency, how many tasks?**
- Different mappings exist, representing different tradeoffs
 - **Many-to-One**: many user threads map to one kernel thread, i.e. kernel sees a single process
 - **One-to-One**: one user thread maps to one kernel thread
 - **Many-to-Many**: many user threads map to many kernel threads


Many-to-One

- Many user-level threads map to one kernel thread
- Pros
 - **Fast**: no system calls required
 - **Portable**: few system dependencies
- Cons
 - **No parallel execution of threads**
 - All thread block when one waits for I/O




One-to-One

- One user-level thread maps to one kernel thread
- Pros: **more concurrency**
 - When one blocks, others can run
 - Better multicore or multiprocessor performance
- Cons: **expensive**
 - Thread operations involve kernel
 - Thread need kernel resources



Many-to-Many

- ❑ Many user-level threads map to many kernel threads ($U \geq K$)
 - Supported in some versions of BSD and Windows
- ❑ Pros: **flexible**
 - OS creates kernel threads for physical concurrency
 - Applications creates user threads for application concurrency
- ❑ Cons: **complex**
 - Most programs use 1:1 mapping anyway



Thread pool

- Problem:
 - Creating a thread for each request: **costly**
 - And, the created thread exits after serving a request
 - More user request → More threads, **server overload**
- Solution: **thread pool**
 - Pre-create a number of threads waiting for work
 - Wake up thread to serve user request --- **faster than thread creation**
 - When request done, don't exit --- go back to pool
 - **Limits the max number of threads**

Other thread design issues

- Semantics of `fork()` system calls
 - Does `fork()` duplicate only the calling thread or all threads?
 - Running threads? Threads trapped in system call?
 - Linux `fork()` copies only the calling thread
- Signal handling
 - Which thread to deliver signals to?
 - Segmentation fault kills process or thread?
- When using threads
 - Make sure to use re-entrant functions
 - Only stack variables for per-call data (no globals)
 - If you want globals? Use thread-local storage (`pthread_key_create`), or an array with one entry per-thread