

Synchronization I

COMS W4118

References: Operating Systems Concepts, Linux Kernel Development, previous W4118s

Copyright notice: care has been taken to use only those web images deemed by the instructor to be in the public domain. If you see a copyrighted image on any slide and are the copyright owner, please contact the instructor. It will be removed.

Banking example

```
int balance = 0;
int main()
{
    pthread_t t1, t2;
    pthread_create(&t1, NULL, deposit, (void*)1);
    pthread_create(&t2, NULL, withdraw, (void*)2);
    pthread_join(t1, NULL);
    pthread_join(t2, NULL);
    printf("all done: balance = %d\n", balance);
    return 0;
}
```

```
void* deposit(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i)
        ++ balance;
}
```

```
void* withdraw(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i)
        -- balance;
}
```

Results of the banking example

```
$ gcc -Wall -lpthread -o bank bank.c
```

```
$ bank
```

```
all done: balance = 0
```

```
$ bank
```

```
all done: balance = 140020
```

```
$ bank
```

```
all done: balance = -94304
```

```
$ bank
```

```
all done: balance = -191009
```

Why?

A closer look at the banking example

```
$ objdump -d bank
```

```
...
```

```
08048464 <deposit>:
```

```
...
```

```
// ++ balance
```

```
8048473: a1 80 97 04 08      mov  0x8049780,%eax
```

```
8048478: 83 c0 01            add  $0x1,%eax
```

```
804847b: a3 80 97 04 08      mov  %eax,0x8049780
```

```
...
```

```
0804849b <withdraw>:
```

```
...
```

```
// -- balance
```

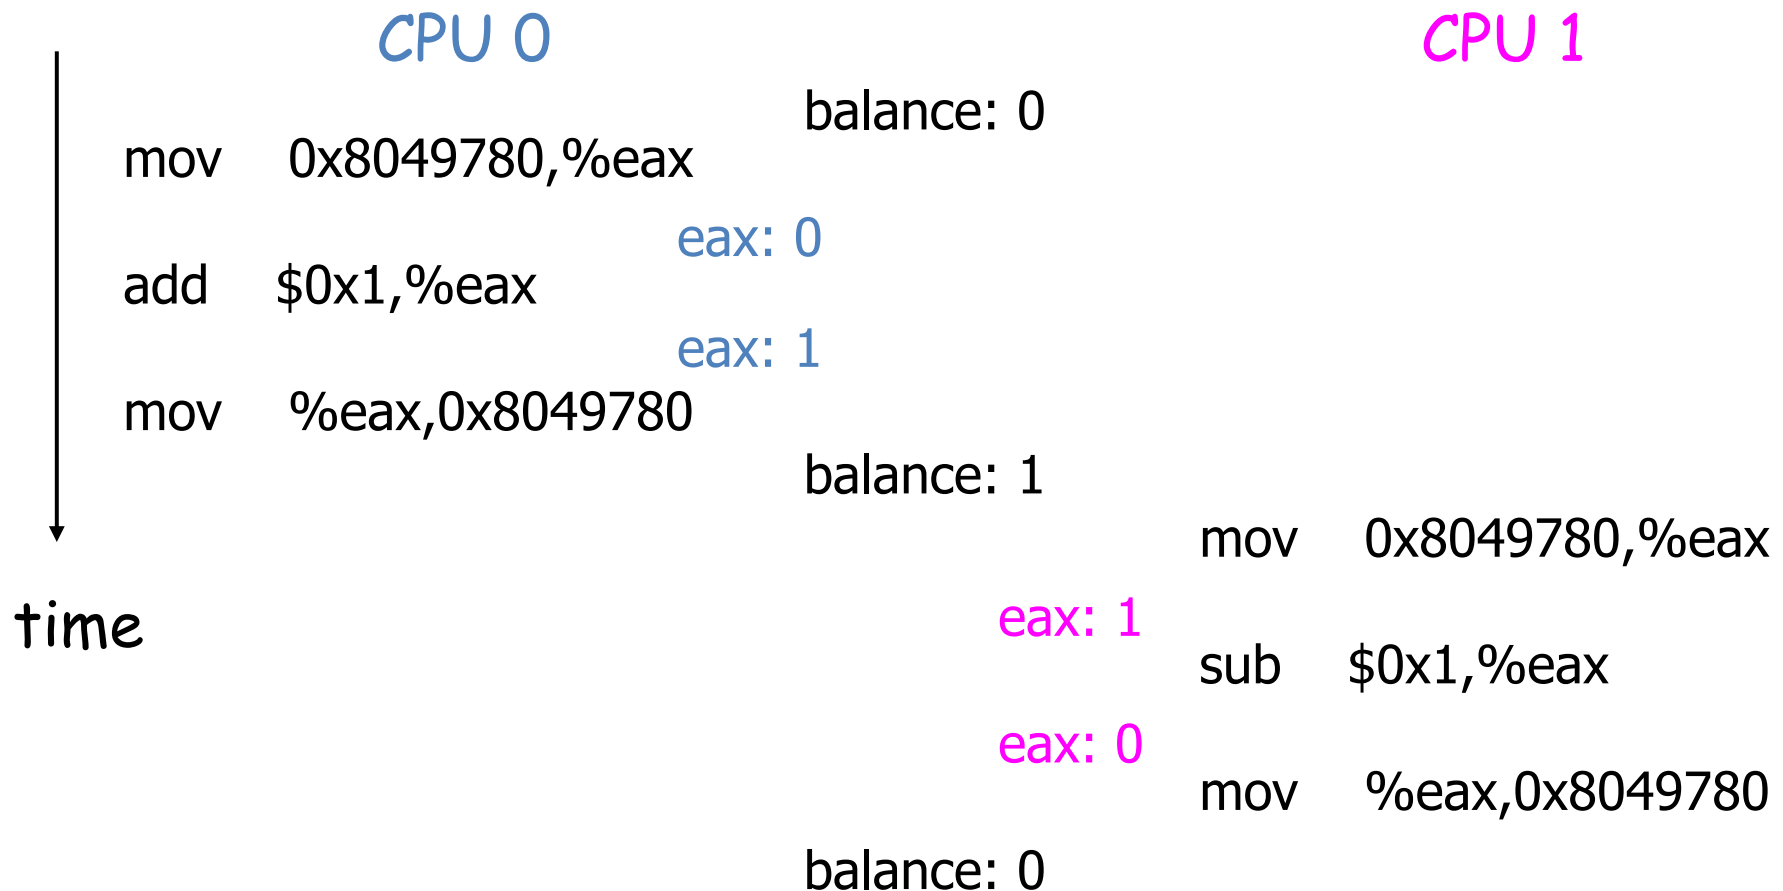
```
80484aa: a1 80 97 04 08      mov  0x8049780,%eax
```

```
80484af: 83 e8 01            sub  $0x1,%eax
```

```
80484b2: a3 80 97 04 08      mov  %eax,0x8049780
```

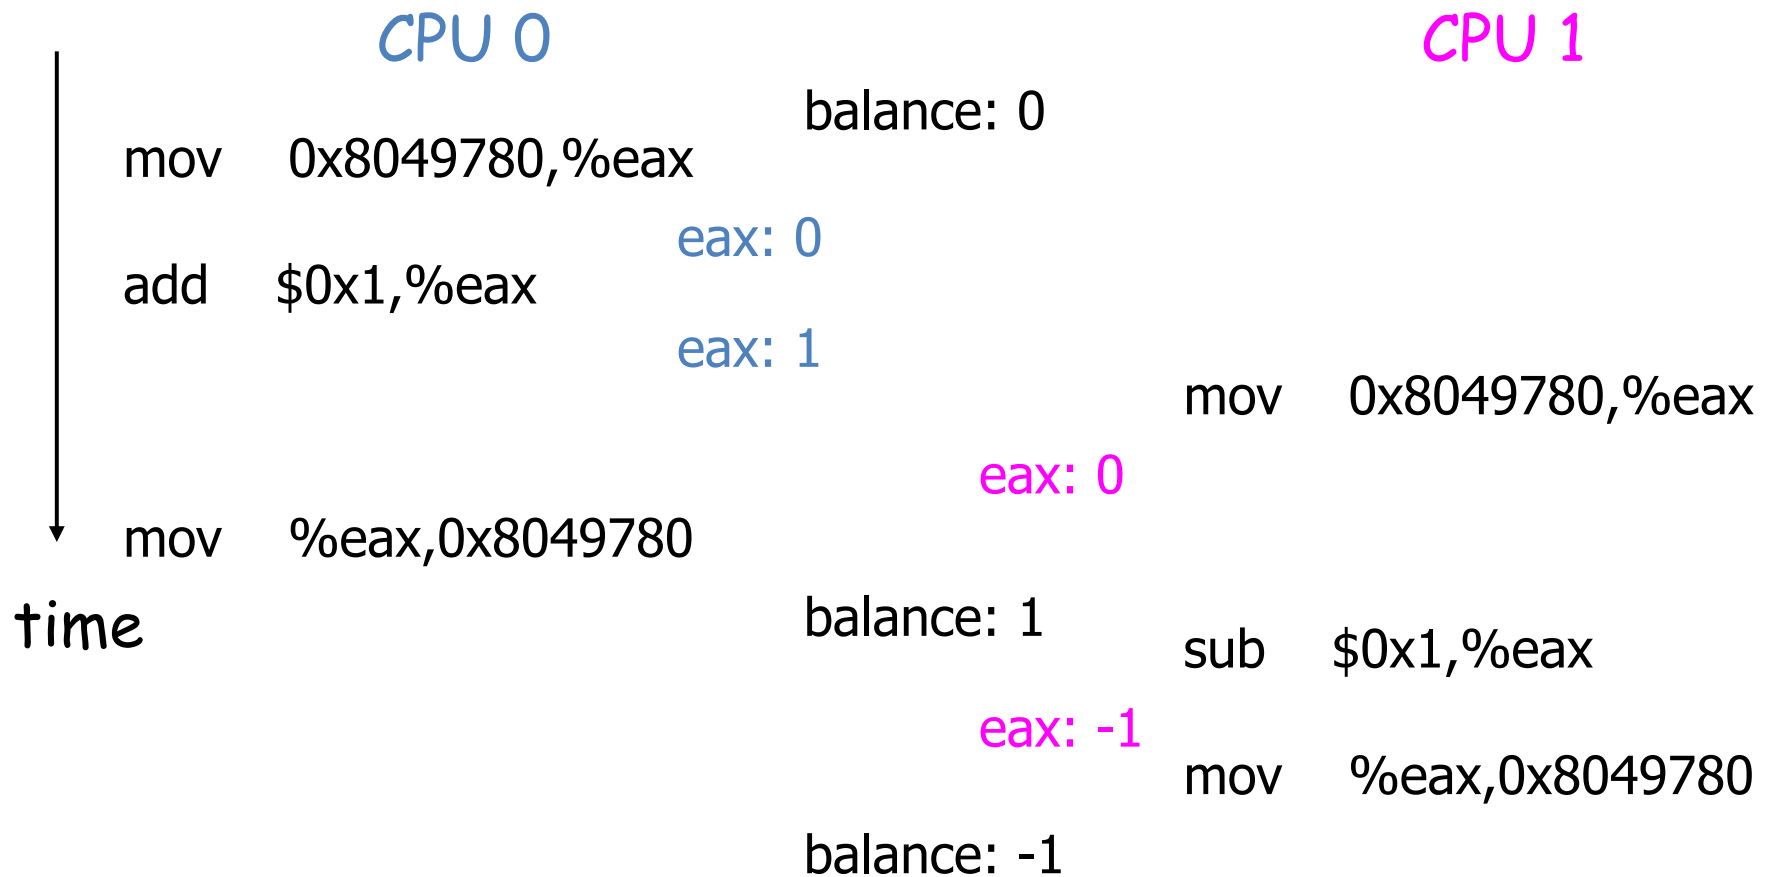
```
...
```

One possible schedule



One deposit and one withdraw,
balance unchanged. Correct

Another possible schedule



One deposit and one withdraw,
balance becomes less. Wrong!

Race condition

- Definition: a timing dependent error involving shared state
- Can be very bad
 - “non-deterministic:” don’t know what the output will be, and it is likely to be different across runs
 - Hard to detect: too many possible schedules
 - Hard to debug: “heisenbug,” debugging changes timing so hides bugs (vs “bohr bug”)

How to avoid race conditions?

- **Atomic operations**: no other instructions can be interleaved, executed “as a unit” “all or none”, guaranteed by hardware
- A possible solution: create a super instruction that does what we want atomically
 - `inc 0x8049780`
- Problem
 - Can’t anticipate **every possible** way we want atomicity
 - Increases hardware complexity, **slows down** other instructions

```
// ++ balance  
mov    0x8049780,%eax  
add    $0x1,%eax  
mov    %eax,0x8049780
```

...

```
// -- balance  
mov    0x8049780,%eax  
sub    $0x1,%eax  
mov    %eax,0x8049780
```

...

Layered approach to synchronization

- Hardware provides simple **low-level atomic operations**, upon which we can build **high-level, synchronization primitives**, upon which we can implement critical sections and build correct multi-threaded/multi-process programs

Properly synchronized application

High-level synchronization
primitives

Hardware-provided low-level
atomic operations

Example synchronization primitives

- Low-level atomic operations
 - On uniprocessor, disable/enable interrupt
 - On x86, aligned load and store of words
 - Special instructions
- High-level synchronization primitives
 - Lock
 - Semaphore
 - Monitor

Outline

- Critical section requirements
- Implementing locks
- Readers-writer lock
- RCU

Avoid race conditions

- **Critical section:** a segment of code that accesses a shared variable (or resource)
- No more than one thread in critical section at a time.

```
// ++ balance  
mov    0x8049780,%eax  
add    $0x1,%eax  
mov    %eax,0x8049780  
...
```

```
// -- balance  
mov    0x8049780,%eax  
sub    $0x1,%eax  
mov    %eax,0x8049780  
...
```

Critical section requirements

- **Safety (aka mutual exclusion):** no more than one thread in critical section at a time.
- **Liveness (aka progress):**
 - If multiple threads simultaneously request to enter critical section, must allow one to proceed
 - Must not depend on threads outside critical section
- **Bounded waiting (aka starvation-free)**
 - Must eventually allow waiting thread to proceed
- Makes no assumptions about the speed and number of CPU
 - However, assumes each thread makes progress

Critical section desirable properties

- **Efficient**: don't consume too much resource while waiting
 - Don't busy wait (spin wait) for a long time. Better to relinquish CPU and let other thread run
- **Fair**: don't make one thread wait longer than others. Hard to do efficiently
- **Simple**: should be easy to use

Implementing critical section using locks

- **lock(l)**: acquire lock exclusively; wait if not available
- **unlock(l)**: release exclusive access to lock

`pthread_mutex_t l = PTHREAD_MUTEX_INITIALIZER`

```
void* deposit(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i) {
        pthread_mutex_lock(&l);
        ++ balance;
        pthread_mutex_unlock(&l);
    }
}
```

```
void* withdraw(void *arg)
{
    int i;
    for(i=0; i<1e7; ++i) {
        pthread_mutex_lock(&l);
        -- balance;
        pthread_mutex_unlock(&l);
    }
}
```

Outline

- Critical section requirements
- Implementing locks
- Readers-writer lock
- RCU

Version 1: Disable interrupts

- **Can cheat on uniprocessor**: implement locks by disabling and enabling interrupts

```
lock()
{
    disable_interrupt();
}
```

```
unlock()
{
    enable_interrupt();
}
```

- **Good**: simple!
- **Bad**:
 - Both operations are **privileged**, can't let user program use
 - Doesn't work on **multiprocessors**
 - Can't use for long critical sections

Version 2: Software Locks

- **Peterson's algorithm**: software-based lock implementation (2 page paper with proof)
- **Good**: doesn't require much from hardware
- Only assumptions:
 - Loads and stores are **atomic**
 - They execute **in order**
 - **Does not require** special hardware instructions

Reference: G. L. Peterson: "Myths About the Mutual Exclusion Problem", *Information Processing Letters* 12(3) 1981, 115–116

Software-based lock: 1st attempt

// 0: lock is available, 1: lock is held by a thread

```
int flag = 0;
```

```
lock()
```

```
{
```

```
    while (flag == 1)
```

```
        ; // spin wait
```

```
    flag = 1;
```

```
}
```

```
unlock()
```

```
{
```

```
    flag = 0;
```

```
}
```

- Idea: use one flag, test then set; if unavailable, **spin-wait**
- Problem?
 - **Not safe**: both threads can be in critical section
 - **Not efficient**: busy wait, particularly bad on uniprocessor (will solve this later)

Unsafe software lock, 1st attempt

```
lock()
{
    1: while (flag == 1)
        ; // spin wait
    2: flag = 1;
}

flag=0;
```

Thread 0:
call lock()
1: while (flag ==1) // it is 0, so
continue

2: flag = 1;

```
unlock()
{
    3: flag = 0;
}
```

Thread 1:

call lock()
1: while(flag == 1) // it is 0, so
continue

2: flag = 1; // ! Thread 0 is already
in critical section

In general, adversarial scheduler model useful to think about concurrency problems

Software-based locks: 2nd attempt

// 1: a thread wants to enter critical section, 0: it doesn't

```
int flag[2] = {0, 0};
```

```
lock()
```

```
{
```

```
    flag[self] = 1; // I need lock
```

```
    while (flag[1-self] == 1)
```

```
        ; // spin wait
```

```
}
```

```
unlock()
```

```
{
```

```
    // not any more
```

```
    flag[self] = 0;
```

```
}
```

- Idea: use per thread flags, set then test, to achieve mutual exclusion
- Why doesn't work?
 - **Not live:** can deadlock

Deadlock: 2nd attempt

// 1: a thread wants to enter critical section, 0: it doesn't

int flag[2] = {0, 0};

lock()

{

flag[self] = 1; // I need lock
while (flag[1-self] == 1)
; // spin wait

}

unlock()

{

// not any more
flag[self] = 0;

}

Thread 0

call lock()

flag[0] = 1;

while (flag[1] == 1) ;
// spins forever too!

Thread1

flag[1] = 1;
while (flag[0] == 1) ;
//spins forever!

...

Software-based locks: 3rd attempt

// whose turn is it?

int turn = 0;

lock()

{

 // wait for my turn

 while (turn == 1 - self)

 ; // spin wait

}

unlock()

{

 // I'm done. your turn

 turn = 1 - self;

}

- Idea: strict alternation to achieve mutual exclusion
- Why doesn't work?
 - **Not live**: depends on threads outside critical section
 - Can't handle repeated calls to lock by same thread

Software-based locks: final attempt (Peterson's algorithm)

```
// whose turn is it?
int turn = 0;
// 1: a thread wants to enter critical section, 0: it doesn't
int flag[2] = {0, 0};

lock()
{
    flag[self] = 1; // I need lock
    turn = 1 - self;
    // wait for my turn
    while (flag[1-self] == 1
    && turn == 1 - self)
        ; // spin wait while the
        // other thread has intent
        // AND it is the other
        // thread's turn
}

unlock()
{
    // not any more
    flag[self] = 0;
}
```

- Why works?
 - Safe?
 - Live?
 - Bounded wait?

Software-based lock

- Problem
 - It's hard!
 - $N > 2$ threads? (Lamport's Bakery algorithm)
 - Modern out of order processors?

Multiprocessor Challenges

- Modern processors are out-of-order/speculative
 - Reorder instructions to keep execution units full
 - Try very hard to avoid inconsistency
 - Guarantees valid only within single execution stream
- Memory access guarantees on x86
 - x86 is relatively conservative with reordering
 - Loads not reordered with other loads
 - Stores not reordered with other stores
 - Stores not reordered with older loads
 - All loads and stores to same location are not reordered
 - Load can reorder with older store to different addr
- Breaks Peterson's algorithm!

Reference: <http://www.linuxjournal.com/article/8211>

<http://www.intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-software-developer-manual-325462.pdf>

Instruction Reordering affects Locking

Thread 0

```
Lock: flag[0] = 1; // I need lock
    turn = 1;
    while (flag[1]==1 && turn==1) ;
}
```

Thread 1

```
Lock: flag[1] = 1; // I need lock
    turn = 0;
    while (flag[0]==1 && turn==0) ;
}
```

- Possible for mutual exclusion to be violated?

— Yes!

Reorder

```
Lock: r1 = Load(flag[1])

turn = 1;
flag[0] = 1; // I need lock
while (r1==1 && turn==1);
// flag[1]==0
}
```

```
Lock: flag[1] = 1; // I need lock
    turn = 0;
    while (flag[0]==1 && turn==0);
    // flag[0]==0
}
```

Memory Barriers

- A memory barrier or fence
 - Ensures that all memory operations up to the barrier are executed before proceeding
- x86 provides several memory fence instructions
 - Relatively expensive (100s of cycles)
 - mfence: all prior memory accesses completed
 - lfence: all prior loads completed
 - sfence: all prior stores flushed

```
lock() {  
    flag[self] = 1; // I need lock  
    turn = 1 - self;  
    sfence; // Store barrier  
    while (flag[1-self] == 1 && turn == 1 - self);  
}
```

Lamport's Bakery Algorithm

- Support more than 2 processes
 - Integer tokens (increasing numbers)
 - Each customer gets next largest token
 - Same token? Smaller thread_id gets priority
 - Smallest token enters critical region

```
bool flag[1..NUM_THREADS] = {0}; // Want to enter
```

```
int token[1..NUM_THREADS] = {0}; // My token
```

```
lock(i) { // Lock by thread i
```

```
    flag[i] = 1;
```

```
    token[i] = 1 + max(token[0..NUM_THREADS-1]);
```

```
    flag[i] = 0;
```

```
    for (j = 1; j <= NUM_THREADS; j++) {
```

```
        while (flag[j]); // Is j getting token?
```

```
        while ((token[j] && ((token[j], j) < (token[i], i))); // j has smaller token?
```

```
    }
```

```
unlock(integer i) {  
    token[i] = 0;
```

```
}
```

Reference: A New Solution of Dijkstra's Concurrent Programming Problem. L. Lamport. Communications of the ACM, 1974. <http://research.microsoft.com/en-us/um/people/lamport/pubs/bakery.pdf>

Version 3: Hardware Instructions

// 0: lock is available, 1: lock is held by a thread

int flag = 0;

lock()

{

while(test_and_set(&flag))

;

}

unlock()

{

flag = 0;

}

- Problem with the test-then-set approach: **test and set are not atomic**
- Fix: **special atomic operation**
 - int test_and_set (int *lock) {
 int old = *lock;
 *lock = 1;
 return old;
}
 - Atomically returns *lock and sets *lock to 1

Implementing test_and_set on x86

```
long test_and_set(volatile long* lock)
{
    int old;
    asm("xchgl %0, %1"
        : "=r"(old), "+m"(*lock) // output
        : "0"(1)                  // input
        : "memory"                // can clobber anything in memory
    );
    return old;
}
```

- `xchg reg, addr`: atomically swaps `*addr` and `reg`
- Spin locks on x86 are implemented using this instruction
- x86 also provides a `lock` prefix that allows bus to be locked for inst
- In Linux:
 - Arch independent: [kernel/spinlock.c](#)
 - Arch dependent: [arch/x86/include/asm/spinlock.h](#)

Spin-wait or block?

- Problem of spin-wait: waste CPU cycles
 - Worst case: thread holding a busy-wait lock gets preempted, other threads try to acquire the same lock
- On uniprocessor: should **not** use spin-lock
 - Yield CPU when lock not available (need OS support)
- On multi-processor
 - Thread holding lock gets preempted → ???
 - Correct action depends on how long before lock release
 - Lock released “quickly” → ?
 - Lock released “slowly” → ?

Problem with simple yield

```
lock()
{
    while(test_and_set(&flag))
        yield();
}
```

- Problem:
 - Still a lot of context switches: **thundering herd**
 - Starvation possible
- Why? **No control** over who gets the lock next
- **Need explicit control over who gets the lock**

Version 4: Sleep Locks

```
lock() {  
  while  
  (test_and_set(&flag))  
    add myself to wait queue  
    yield  
  ...  
}
```

```
unlock() {  
  flag = 0  
  if(any thread in wait queue)  
    wake up one wait thread  
  ...  
}
```

← Lock from another thread?

- The idea: **add thread to queue when lock unavailable;** in **unlock()**, wake up one thread in queue
- Problem I: **lost wakeup**
- Problem II: **wrong thread gets lock**

Lost wakeup

```
lock() {  
  1: while (test_and_set(&flag))  
  2: add myself to wait queue  
  3: yield  
  ...  
}
```

Thread 0:

```
call lock()  
while (test_and_set(&flag)) {  
  
  add myself to wait queue  
  yield  
} // wait forever (or until next unlock)!
```

```
unlock() {  
  4: flag = 0  
  5: if (any thread in wait queue)  
  6: wake up one wait thread  
  ...  
}
```

Thread 1

```
call unlock()  
flag = 0  
if (any thread in wait queue) // No!  
  wake_up_one_wait_thread
```

Wrong thread gets lock

```
lock() {  
  1: while (test_and_set(&flag))  
  2: add myself to wait queue  
  3: yield  
  ...  
}
```

```
unlock() {  
  4: flag = 0  
  5: if (any thread in wait queue)  
  6: wake up one wait thread  
  ...  
}
```

Thread 0:

```
call lock()  
while (test_set(&flag))  
  add myself to wait queue  
  yield
```

Thread 1

```
call unlock()  
flag = 0  
if (thread in wait queue)  
  wake_up_thread
```

Thread 2

```
call lock()  
while (test_set(&flag))
```

- Fix: `unlock()` directly transfers lock to waiting thread

Implementing locks: version 4, the code

```
typedef struct __mutex_t {  
    int flag;      // 0: mutex is available, 1: mutex is not available  
    int guard;     // guard lock to avoid losing wakeups  
    queue_t *q;    // queue of waiting threads  
} mutex_t;
```

```
void lock(mutex_t *m) {  
    while (test_and_set(m->guard))  
        ; //acquire guard lock by spinning  
    if (m->flag == 0) {  
        m->flag = 1; // acquire mutex  
        m->guard = 0;  
    } else {  
        enqueue(m->q, self);  
        m->guard = 0;  
        yield();  
    }  
}
```

```
void unlock(mutex_t *m) {  
    while (test_and_set(m->guard))  
        ;  
    if (queue_empty(m->q))  
        // release mutex; no one wants mutex  
        m->flag = 0;  
    else  
        // direct transfer mutex to next thread  
        wakeup(dequeue(m->q));  
    m->guard = 0;  
}
```

Adaptive Mutexes

- Cons of Spinlocks
 - Inefficient if lock is held for long duration
- Cons of Sleeplocks
 - Higher overhead, state maintenance
- Solaris, OS X, FreeBSD
 - Idea: use spinlock if holder is currently running, sleeplock otherwise
 - Best of both worlds

Outline

- Critical section requirements
- Implementing locks
- Readers-writer lock

Readers-Writers problem

- A **reader** is a thread that needs to look at the shared data but won't change it
- A **writer** is a thread that modifies the shared data
- Example: making an airline reservation
- Courtois et al 1971

Readers-writer lock

```
rwlock_t lock;
```

Writer

```
write_lock (&lock);  
...  
// write shared data  
...  
write_unlock (&lock);
```

Reader

```
read_lock (&lock);  
...  
// read shared data  
...  
read_unlock (&lock);
```

- **read_lock**: acquires lock in read (shared) mode
 - Lock is not acquired or is acquired in read mode → success
 - Otherwise (lock is in write mode) → wait
- **write_lock**: acquires lock in write (exclusive) mode
 - Lock is not acquired → success
 - Otherwise → wait

Implementing readers-writer lock

```
struct rwlock_t {  
    int nreader;    // init to 0  
    lock_t guard;  // init to unlocked  
    lock_t lock;    // init to unlocked  
};
```

```
write_lock(rwlock_t *l)  
{  
    lock(&l->lock);  
}
```

```
write_unlock(rwlock_t *l)  
{  
    unlock(&l->lock);  
}
```

```
read_lock(rwlock_t *l)  
{  
    lock(&l->guard);  
    ++ nreader;  
    if(nreader == 1) // first reader  
        lock(&l->lock);  
    unlock(&l->guard);  
}
```

```
read_unlock(rwlock_t *l)  
{  
    lock(&l->guard);  
    -- nreader;  
    if(nreader == 0) // last reader  
        unlock(&l->lock);  
    unlock(&l->guard);  
}
```

Problem: may starve writer!

Driving out readers in a RW-Lock

```
struct rwlock_t {  
    int nreader;    // init to 0  
    lock_t guard;   // init to unlocked  
    lock_t lock;    // init to unlocked  
    lock_t writer;  // init to unlocked  
};
```

```
write_lock(rwlock_t *l)  
{  
    lock(&l->writer);  
    lock(&l->lock);  
    unlock(&l->writer);  
}
```

```
write_unlock(rwlock_t *l)  
{  
    unlock(&l->lock);  
}
```

```
read_lock(rwlock_t *l)  
{  
    lock(&l->writer);  
    lock(&l->guard);  
    ++ nreader;  
    if(nreader == 1) // first reader  
        lock(&l->lock);  
    unlock(&l->guard);  
    unlock(&l->writer);  
}
```

```
read_unlock(rwlock_t *l)  
{  
    lock(&l->guard);  
    -- nreader;  
    if(nreader == 0) // last reader  
        unlock(&l->lock);  
    unlock(&l->guard);  
}
```

Q: In write_lock, can we just use guard instead of writer lock?