

Linux System Overview: From Boot To Panic

Boot process

Under BIOS-based systems:

1. Power-on self-test (POST) and peripheral initializations
2. Jump to the boot code in the first 440 bytes of Master Boot Record (MBR)
3. MBR boot code locates and launches boot loader – ex) GRUB, Syslinux
4. Boot loader reads its configuration and possibly presents menu
5. Boot loader loads the kernel and launches it
6. Kernel unpacks the `initramfs` (initial RAM filesystem)
 - initial temporary root filesystem
 - contains device drivers needed to mount real root filesystem
7. Kernel switches to real root filesystem
8. Kernel launches the first user-level process `/sbin/init` (pid == 1)
 - `/sbin/init` is the mother of all processes
 - traditional SysV-style init reads `/etc/inittab`
 - in Arch, `/sbin/init` is a symlink to `systemd`

User session

1. `init` calls `getty`
 - `getty` is called for each virtual terminal (normally 6 of them)
 - `getty` checks user name and password against `/etc/passwd`
2. `getty` calls `login`, which then runs the user's shell – ex) `bash`
 - `login` sets the user's environment variables
 - the user's shell is specified in `/etc/passwd`

Process control

Displaying process hierarchy

```
ps axf      # display process tree

ps axfj     # with more info

ps axfjww   # even if lines wrap around terminal
```

Creating processes

Simple shell program from APUE3, 1.6:

```
#include "apue.h"
#include <sys/wait.h>

int
main(void)
{
    char    buf[MAXLINE];    /* from apue.h */
    pid_t   pid;
    int     status;

    printf("%% "); /* print prompt (printf requires %% to print %) */
    while (fgets(buf, MAXLINE, stdin) != NULL) {
        if (buf[strlen(buf) - 1] == '\n')
            buf[strlen(buf) - 1] = 0; /* replace newline with null */

        if ((pid = fork()) < 0) {
            err_sys("fork error");
        } else if (pid == 0) { /* child */
            execlp(buf, buf, (char *)0);
            err_ret("couldn't execute: %s", buf);
            exit(127);
        }

        /* parent */
        if ((pid = waitpid(pid, &status, 0)) < 0)
            err_sys("waitpid error");
        printf("%% ");
    }
    exit(0);
}
```

Questions:

- What happens if the parent process terminates before its children?
 - create this condition and see it in `ps` output
- What happens if a child process has terminated, but the parent never calls `waitpid()` ?
 - create this condition and see it in `ps` output

Signals

How do you terminate the following program?

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>

static void sig_int(int signo)
{
    printf("stop pressing ctrl-c!\n");
}

int main()
{
    if (signal(SIGINT, &sig_int) == SIG_ERR) {
        perror("signal() failed");
        exit(1);
    }

    int i = 0;
    for (;;) {
        printf("%d\n", i++);
        sleep(1);
    }
}
```

How about this one?

```
#include <stdio.h>
#include <stdlib.h>
#include <signal.h>
#include <unistd.h>

static void sig_int(int signo)
{
    printf("stop pressing ctrl-c!\n");
}

static void sig_term(int signo)
{
    printf("stop trying to kill me!\n");
}

int main()
{
    if (signal(SIGINT, &sig_int) == SIG_ERR) {
        perror("signal() failed");
        exit(1);
    }

    if (signal(SIGTERM, &sig_term) == SIG_ERR) {
        perror("signal() failed");
        exit(1);
    }

    int i = 0;
    for (;;) {
        printf("%d\n", i++);
        sleep(1);
    }
}
```

Signals and system calls

The following program, from APUE3 section 1.9, adds signal handling to the simple shell program we looked at before:

```

#include "apue.h"
#include <sys/wait.h>

static void sig_int(int);      /* our signal-catching function */

int
main(void)
{
    char    buf[MAXLINE];     /* from apue.h */
    pid_t   pid;
    int     status;

    if (signal(SIGINT, sig_int) == SIG_ERR)
        err_sys("signal error");

    printf("%% "); /* print prompt (printf requires %% to print %) */
    while (fgets(buf, MAXLINE, stdin) != NULL) {
        if (buf[strlen(buf) - 1] == '\n')
            buf[strlen(buf) - 1] = 0; /* replace newline with null */

        if ((pid = fork()) < 0) {
            err_sys("fork error");
        } else if (pid == 0) { /* child */
            execlp(buf, buf, (char *)0);
            err_ret("couldn't execute: %s", buf);
            exit(127);
        }

        /* parent */
        if ((pid = waitpid(pid, &status, 0)) < 0)
            err_sys("waitpid error");
        printf("%% ");
    }
    exit(0);
}

void
sig_int(int signo)
{
    printf("interrupt\n%% ");
}

```

Does it work as expected in Linux? What about other UNIX platforms like Mac OS X? What's going on here?

X session

Let's take a look at the processes in your X session:

```

145 ?      Ss      0:00 login -- jae
406 tty1   Ss      0:00  \_ -bash
421 tty1   S+      0:00      \_ xinit /etc/xdg/xfce4/xinitrc -- /etc/X11/xinit/xs
422 ?      Ss      0:03          \_ /usr/bin/X -nolisten tcp :0 vt1
425 tty1   S       0:00          \_ sh /etc/xdg/xfce4/xinitrc
430 tty1   Sl      0:00              \_ xfce4-session
448 tty1   S       0:00                  \_ xfwm4 --display :0.0 --sm-client-id 2
450 tty1   Sl      0:00                  \_ Thunar --sm-client-id 2ed5a75a5-2f67-
452 tty1   Sl      0:01                  \_ xfce4-panel --display :0.0 --sm-clien
494 tty1   S       0:00                      | \_ /usr/lib/xfce4/panel/wrapper /usr
498 tty1   S       0:00                      | \_ /usr/lib/xfce4/panel/wrapper /usr
454 tty1   Sl      0:00                  \_ xfdesktop --display :0.0 --sm-client-
456 tty1   Sl      0:01                  \_ xfce4-terminal --geometry=100x36 --di
505 tty1   S       0:00                      \_ gnome-pty-helper
507 pts/0   Ss+     0:00                      \_ bash
518 pts/1   Ss      0:00                      \_ bash
2330 pts/1   R+      0:00                      | \_ ps afx
1218 pts/2   Ss+     0:00                      \_ bash

```

Explorations:

- Identify the function of each process
- Understand the network client-server architecture of X window system
- Try running a simpler X session – `twm` or even just `xterm`

Kernel module

Here is a skeleton kernel module from OSCE2, Chapter 2, Programming Project:

```
#include <linux/init.h>
#include <linux/module.h>
#include <linux/kernel.h>

/* This function is called when the module is loaded. */
int simple_init(void)
{
    printk(KERN_INFO "Loading Module\n");

    return 0;
}

/* This function is called when the module is removed. */
void simple_exit(void) {
    printk(KERN_INFO "Removing Module\n");
}

/* Macros for registering module entry and exit points. */
module_init( simple_init );
module_exit( simple_exit );

MODULE_LICENSE("GPL");
MODULE_DESCRIPTION("Simple Module");
MODULE_AUTHOR("SGG");
```

Here is the Makefile:

```
obj-m += simple.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
```

Can you modify the code to cause the kernel panic?

References

- [Source code for this lecture](#)
- [Arch Boot Process](#)

~~~~~

*Last updated: 2014-01-27*