

Nonblocking I/O and I/O multiplexing

Nonblocking I/O

- Two ways to make “slow” systems calls nonblocking:
 1. call `open()` with `O_NONBLOCK`
 2. call `fcntl()` to turn on `O_NONBLOCK` file status flag (recall that file status flag is part of file table entry – the middle layer)
- Nonblocking slow system call returns `-1` with `errno` set to `EAGAIN` if it would have blocked
- Nonblocking write example:

```
#include "apue.h"
#include <errno.h>
#include <fcntl.h>

char    buf[500000];

int
main(void)
{
    int    ntowrite, nwrite;
    char    *ptr;

    ntowrite = read(STDIN_FILENO, buf, sizeof(buf));
    fprintf(stderr, "read %d bytes\n", ntowrite);

    set_fl(STDOUT_FILENO, O_NONBLOCK); /* set nonblocking */

    ptr = buf;
    while (ntowrite > 0) {
        errno = 0;
        nwrite = write(STDOUT_FILENO, ptr, ntowrite);
        fprintf(stderr, "nwrite = %d, errno = %d\n", nwrite, errno);

        if (nwrite > 0) {
            ptr += nwrite;
            ntowrite -= nwrite;
        }
    }

    clr_fl(STDOUT_FILENO, O_NONBLOCK); /* clear nonblocking */

    exit(0);
}
```

- `set_fl()` and `clr_fl()`:

```
#include "apue.h"
#include <fcntl.h>

void set_fl(int fd, int flags) /* flags are file status flags to turn on */
{
    int    val;

    if ((val = fcntl(fd, F_GETFL, 0)) < 0)
        err_sys("fcntl F_GETFL error");

    val |= flags;      /* turn on flags */

    if (fcntl(fd, F_SETFL, val) < 0)
        err_sys("fcntl F_SETFL error");
}

void clr_fl(int fd, int flags) /* flags are file status flags to turn off */
{
    int    val;

    if ((val = fcntl(fd, F_GETFL, 0)) < 0)
        err_sys("fcntl F_GETFL error");

    val &= ~flags;     /* turn flags off */

    if (fcntl(fd, F_SETFL, val) < 0)
        err_sys("fcntl F_SETFL error");
}
```

I/O multiplexing

- `cat` :

```
while ((n = read(STDIN_FILENO, buf, BUFSIZ)) > 0)
    if (write(STDOUT_FILENO, buf, n) != n)
        err_sys("write error");
```

- What about netcat (`nc`)?
- `select()` :

```
#include <sys/select.h>
```

```
int select(int maxfdp1, fd_set *restrict readfds,  
           fd_set *restrict writefds, fd_set *restrict exceptfds,  
           struct timeval *restrict tvptr);
```

Returns: count of ready descriptors, 0 on timeout, -1 on error

```
int FD_ISSET(int fd, fd_set *fdset);
```

Returns: nonzero if fd is in set, 0 otherwise

```
void FD_CLR(int fd, fd_set *fdset);
```

```
void FD_SET(int fd, fd_set *fdset);
```

```
void FD_ZERO(fd_set *fdset);
```

- `select()` will always get interrupted on signals, even when the `SA_RESTART` flag was used
- Nonblocking write example, augmented with `select()`:

```
#include "apue.h"
#include <errno.h>
#include <fcntl.h>
#include <sys/select.h>

char    buf[5000000];

int
main(void)
{
    int    ntowrite, nwrite;
    char    *ptr;

    ntowrite = read(STDIN_FILENO, buf, sizeof(buf));
    fprintf(stderr, "read %d bytes\n", ntowrite);

    set_fl(STDOUT_FILENO, O_NONBLOCK); /* set nonblocking */

    ptr = buf;
    while (ntowrite > 0) {
        errno = 0;

        // Wait until stdout is ready before we call write()
        fd_set write_fds;
        FD_ZERO(&write_fds);
        FD_SET(STDOUT_FILENO, &write_fds);
        select(STDOUT_FILENO + 1, NULL, &write_fds, NULL, NULL);

        nwrite = write(STDOUT_FILENO, ptr, ntowrite);
        fprintf(stderr, "nwrite = %d, errno = %d\n", nwrite, errno);

        if (nwrite > 0) {
            ptr += nwrite;
            ntowrite -= nwrite;
        }
    }

    clr_fl(STDOUT_FILENO, O_NONBLOCK); /* clear nonblocking */

    exit(0);
}
```

~~~~~  
*Last updated: 2014-02-20*