

UNIX File I/O

File I/O system calls

open

Example (taken from `man open` in Linux):

```
#include <fcntl.h>
...
int fd;
mode_t mode = S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH;
char *pathname = "/tmp/file";
...
fd = open(pathname, O_WRONLY | O_CREAT | O_TRUNC, mode);
...
```

Another example – creating a lock file:

```
fd = open("/var/run/myprog.pid", O_WRONLY | O_CREAT | O_EXCL, 0644);
```

creat

Redundant: `creat(path, mode)` is equivalent to
`open(path, O_WRONLY|O_CREAT|O_TRUNC, mode)`

And it should have been called `create`, says Ken Thompson

close

```
int close(int fildes);
```

lseek

```
off_t lseek(int fildes, off_t offset, int whence);
```

- If `whence` is `SEEK_SET`, the file offset shall be set to `offset` bytes.
- If `whence` is `SEEK_CUR`, the file offset shall be set to its current location plus `offset`.
- If `whence` is `SEEK_END`, the file offset shall be set to the size of the file plus `offset`.

read

```
ssize_t read(int fildes, void *buf, size_t nbyte);
```

- returns number of bytes read, 0 if end of file, -1 on error

Note:

- Number of bytes read may be less than the requested nbyte
- `read()` may block forever on a “slow” read from pipes, FIFOs (aka named pipes), sockets, or keyboard
- For sockets, `read(socket, buf, nbyte)` is equivalent to `recv(socket, buf, nbyte, 0)`

```
recv(int socket, void *buffer, size_t length, int flags)
```

- normally, `recv()` blocks until it has received at least 1 byte
- returns num bytes received, 0 if connection closed, -1 if error

write

```
ssize_t write(int fildes, const void *buf, size_t nbyte);
```

- returns number of bytes written, -1 on error

Note:

- Number of bytes written may be less than the requested nbyte – ex) filling up a disk
- `write()` may block forever on a “slow” write into pipes, FIFOs, or sockets
- For sockets, `write(socket, buf, nbyte)` is equivalent to `send(socket, buf, nbyte, 0)`

```
send(int socket, const void *buffer, size_t length, int flags)
```

- normally, `send()` blocks until it sends all bytes requested
- returns num bytes sent or -1 for error
- If the file was opened with `O_APPEND` flag, the file offset gets set to the end of the file prior to each write
 - setting of the offset and writing happen in an atomic operation

Signals and system calls (revisited)

Recall `shell2.c` from previous lecture:

```

#include "apue.h"
#include <sys/wait.h>

static void sig_int(int);      /* our signal-catching function */

int
main(void)
{
    char    buf[MAXLINE];     /* from apue.h */
    pid_t   pid;
    int     status;

    if (signal(SIGINT, sig_int) == SIG_ERR)
        err_sys("signal error");

    printf("%% "); /* print prompt (printf requires %% to print %) */
    while (fgets(buf, MAXLINE, stdin) != NULL) {
        if (buf[strlen(buf) - 1] == '\n')
            buf[strlen(buf) - 1] = 0; /* replace newline with null */

        if ((pid = fork()) < 0) {
            err_sys("fork error");
        } else if (pid == 0) { /* child */
            execlp(buf, buf, (char *)0);
            err_ret("couldn't execute: %s", buf);
            exit(127);
        }

        /* parent */
        if ((pid = waitpid(pid, &status, 0)) < 0)
            err_sys("waitpid error");
        printf("%% ");
    }
    exit(0);
}

void
sig_int(int signo)
{
    printf("interrupt\n%% ");
}

```

Three problems:

1. "Slow" system calls may get interrupted on signals
 - `errno` set to `EINTR`
2. Signals get lost
 - signal handling resets after each signal

3. Signal handler calls non-reentrant functions

- `malloc()` , `free()` , standard I/O functions

Solutions:

- Don't use `signal()` ; use `sigaction()` instead
- From a signal handler, call only *async-signal-safe* functions (see `man 7 signal`)

File sharing

1. Kernel data structures for open files

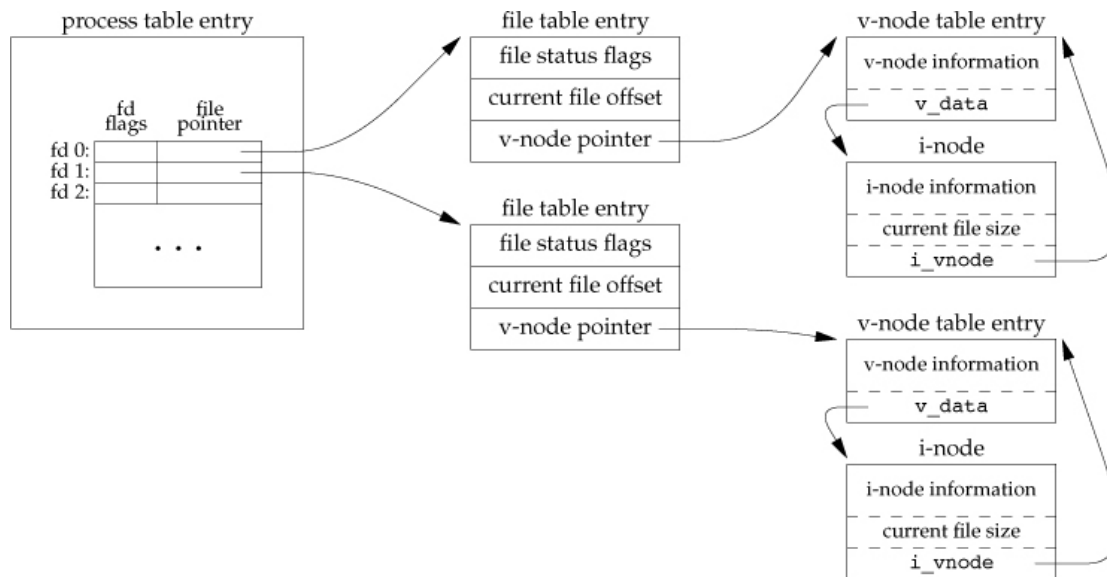


Figure 3.7, APUE

2. Two independent processes with the same file open

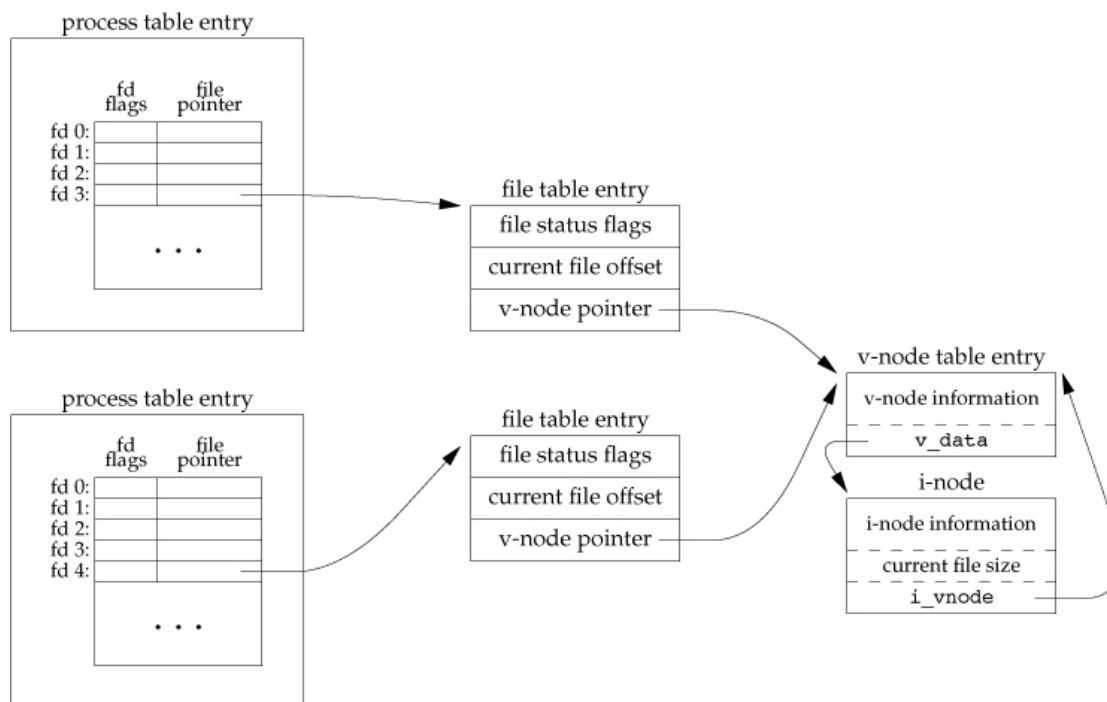


Figure 3.8, APUE


3. Kernel data structures after `dup(1)`

Figure 3.9, APUE

4. Sharing of open files between parent and child after `fork`

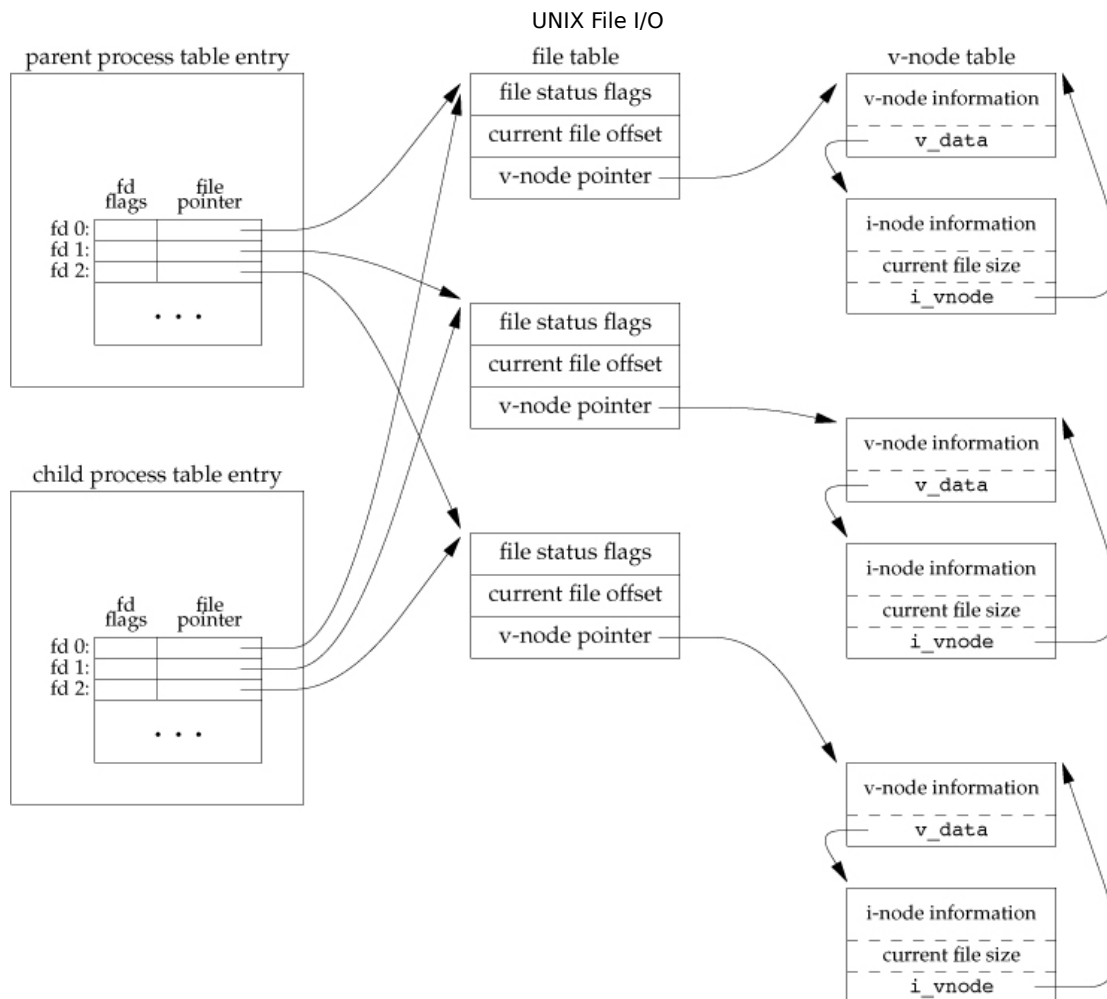


Figure 8.2, APUE

~~~~~

*Last updated: 2015-01-30*