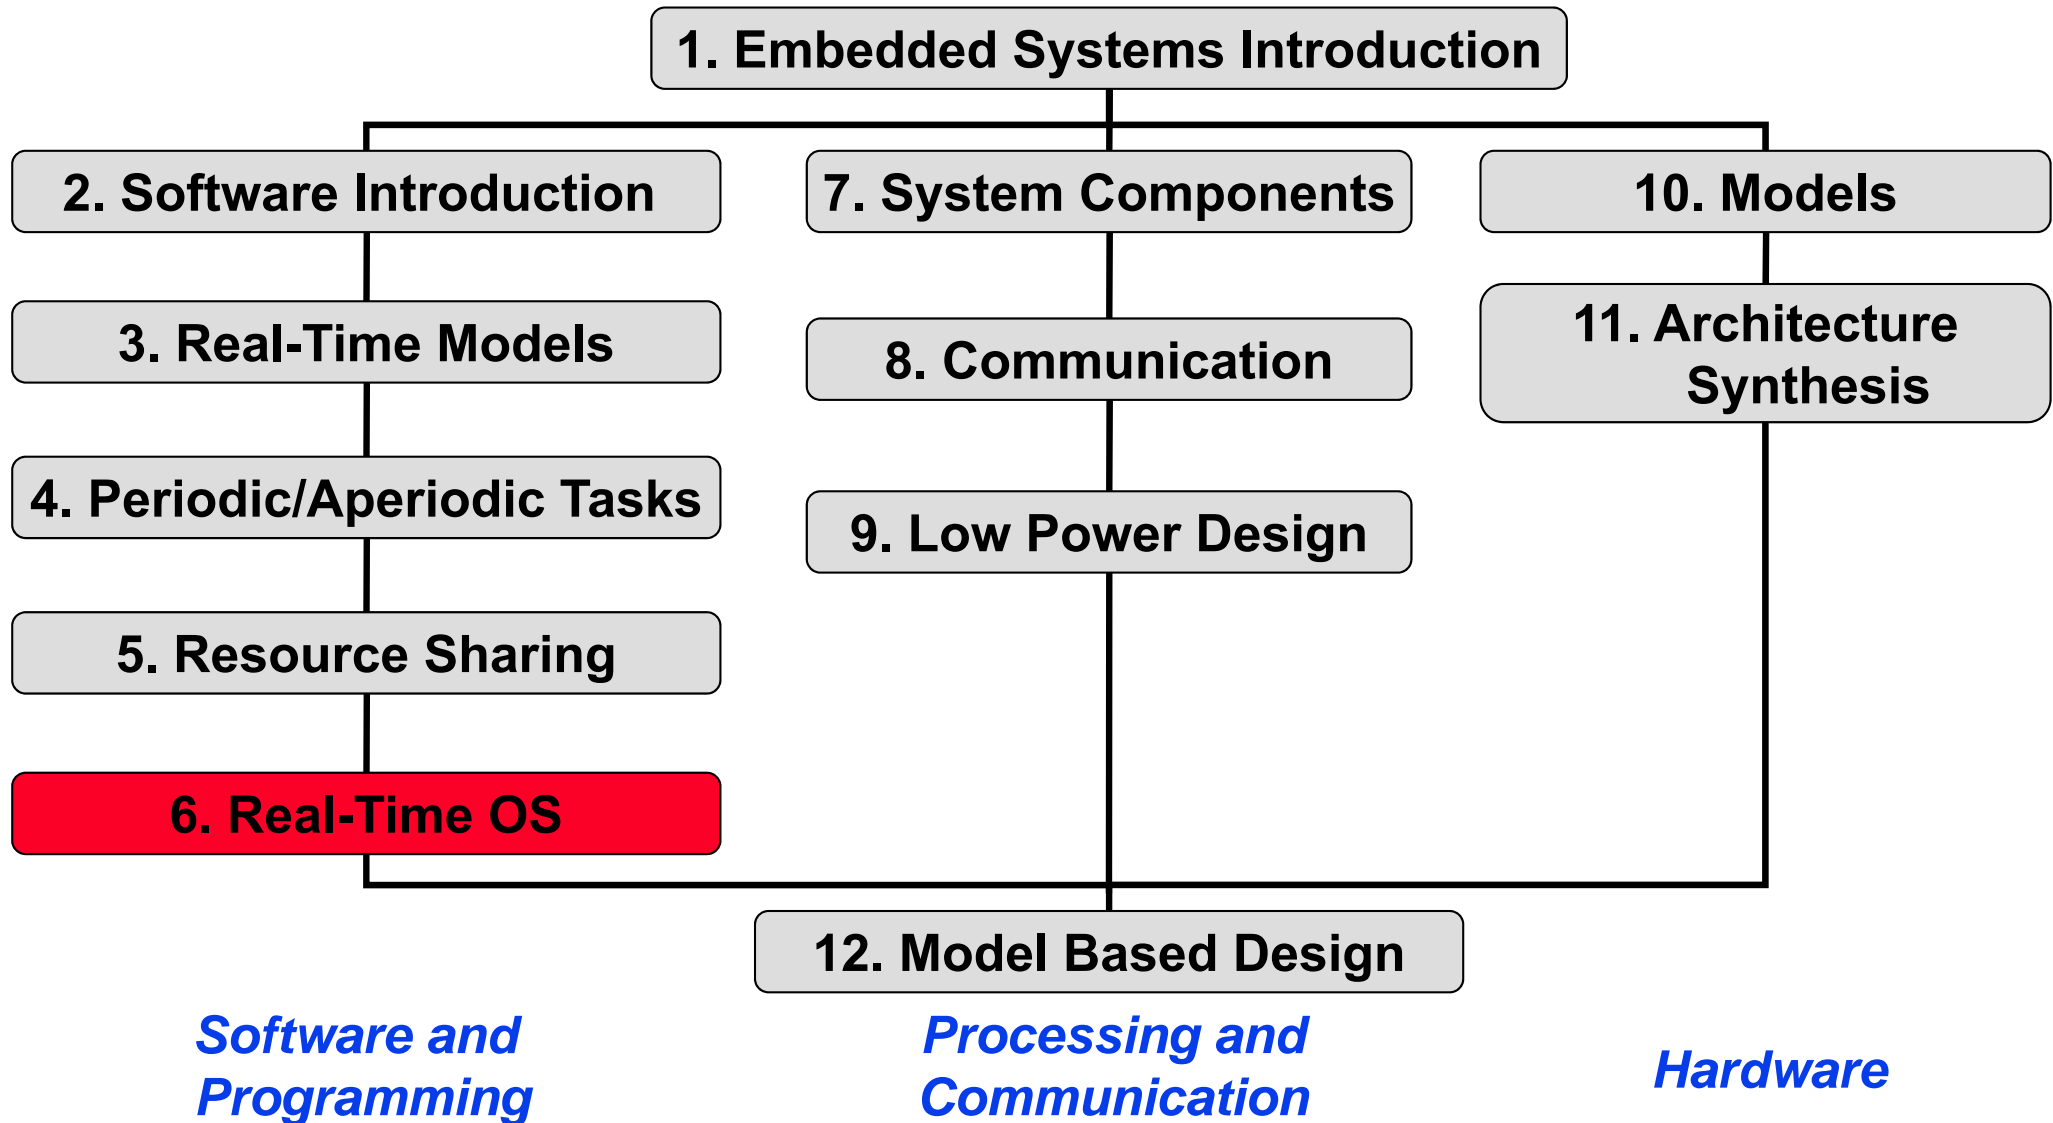


# Embedded Systems

## 6. Real-Time Operating Systems

Lothar Thiele

# Contents of Course



# Embedded OS

---

## ► *Why an OS at all?*

- Same reasons why we need one for a traditional computer.
- Not all services are needed for any device.

## ► Large variety of *requirements* and environments:

- Critical applications with high functionality (medical applications, space shuttle, process automation, ...).
- Critical applications with small functionality (ABS, pace maker, ...)
- Not very critical applications with varying functionality (smart phone, smart card, microwave oven, ...)

# Embedded OS

---

- ▶ Why is a *desktop OS not suited*?
  - Monolithic kernel is too feature reach.
  - Monolithic kernel is not modular, fault-tolerant, configurable, modifiable, ... .
  - Takes too much memory space.
  - It is often too ressource hungry in terms of computation time.
  - Not designed for mission-critical applications.
  - Timing uncertainty too large.

# Embedded Operating Systems

---

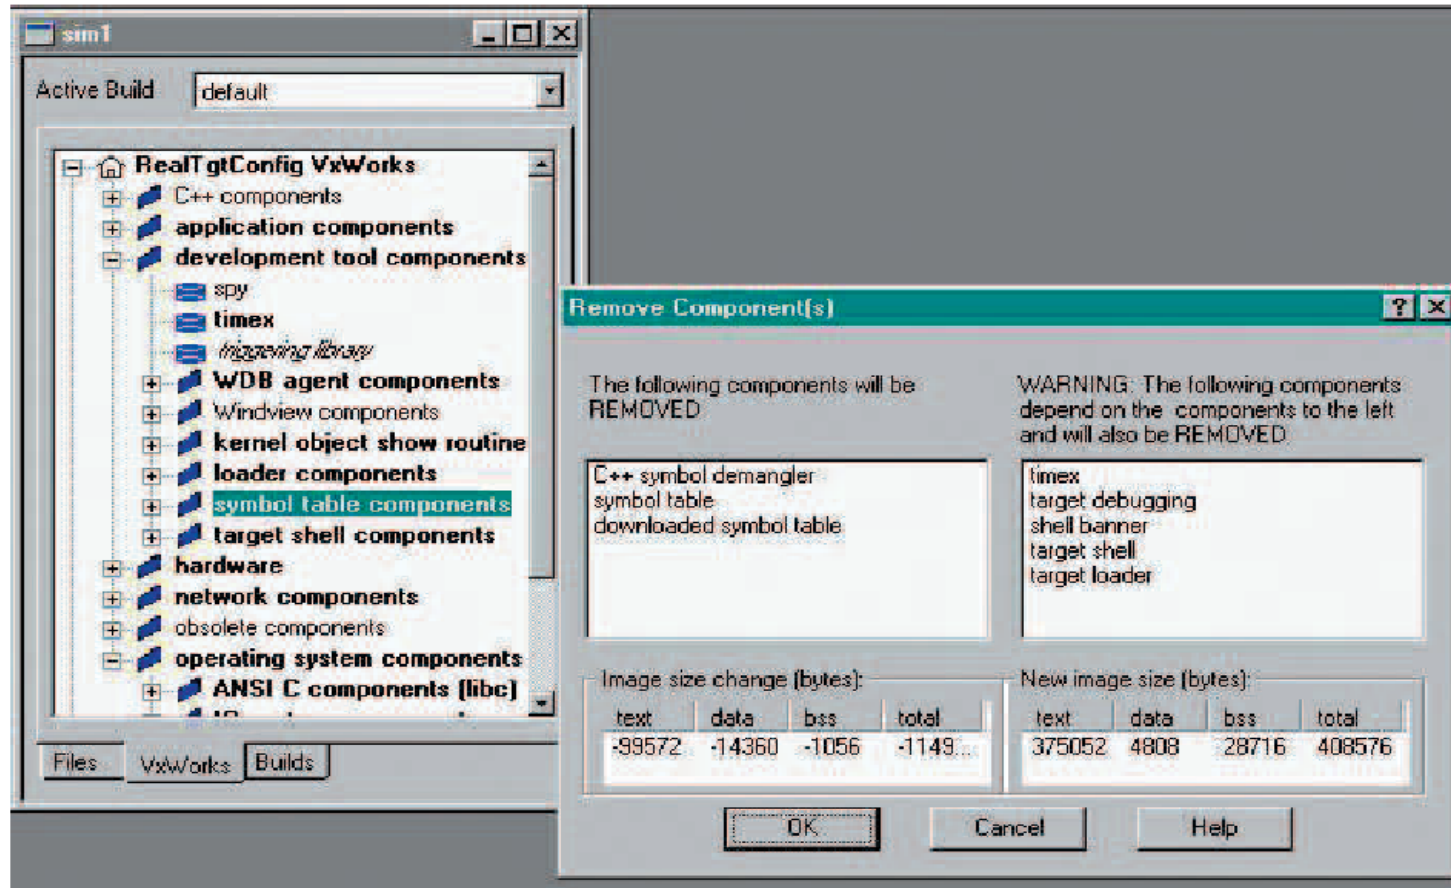
## *Configurability*

- No single RTOS will fit all needs, no overhead for unused functions/data tolerate: configurability is needed.
- For example, there are many embedded systems without external memory, a keyboard, a screen or a mouse.

## *Configurability examples:*

- Simplest form: remove unused functions (by linker for example).
- Conditional compilation (using `#if` and `#ifdef` commands).
- Validation is a potential problem of systems with a large number of derived operating systems:
  - each derived operating system must be tested thoroughly;
  - for example, eCos (open source RTOS from Red Hat) includes 100 to 200 configuration points.

# Example: Configuration of VxWorks



Automatic dependency analysis and size calculations allow users to quickly custom-tailor the VxWORKS operating system.

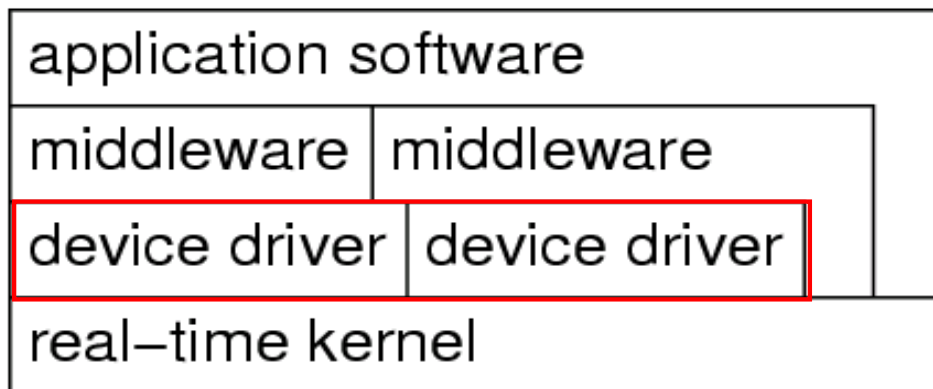
© Windriver

# Embedded operating systems

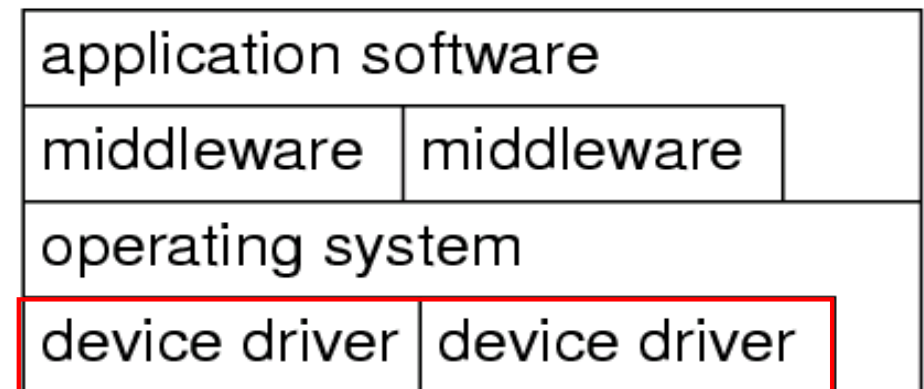
*Device drivers handled by tasks* instead of hidden integrated drivers:

- Improve predictability; everything goes through scheduler
- Effectively no device that needs to be supported by all versions of the OS, except maybe the system timer.

## RTOS



## Standard OS



# Embedded Operating Systems

---

## *Interrupts can be employed by any process*

- For standard OS: this would be serious source of unreliability.
- But embedded programs can be considered to be tested ...
- It is possible to let interrupts directly start or stop tasks (by storing the tasks start address in the interrupt table). More efficient and predictable than going through OS interfaces and services.
- However, composability suffers: if a specific task is connected to some interrupt, it may be difficult to add another task which also needs to be started by the same event.
- If real-time processing is of concern, time to handle interrupts need to be considered. For example, interrupts may be handled by the scheduler.



# Embedded Operating Systems

---

## *Protection mechanisms are not always necessary:*

- Embedded systems are typically designed for a single purpose, untested programs rarely loaded, software considered reliable.
- *Privileged I/O* instructions not necessary and tasks can do their own I/O.

Example: Let `switch` be the address of some switch. Simply use

`load register, switch`

instead of a call to the underlying operating system.

- However, protection mechanisms may be needed for safety and security reasons.

# Real-time Operating Systems

---

- ▶ *A real-time operating system is an operating system that supports the construction of real-time systems.*

- ▶ **Three key requirements:**

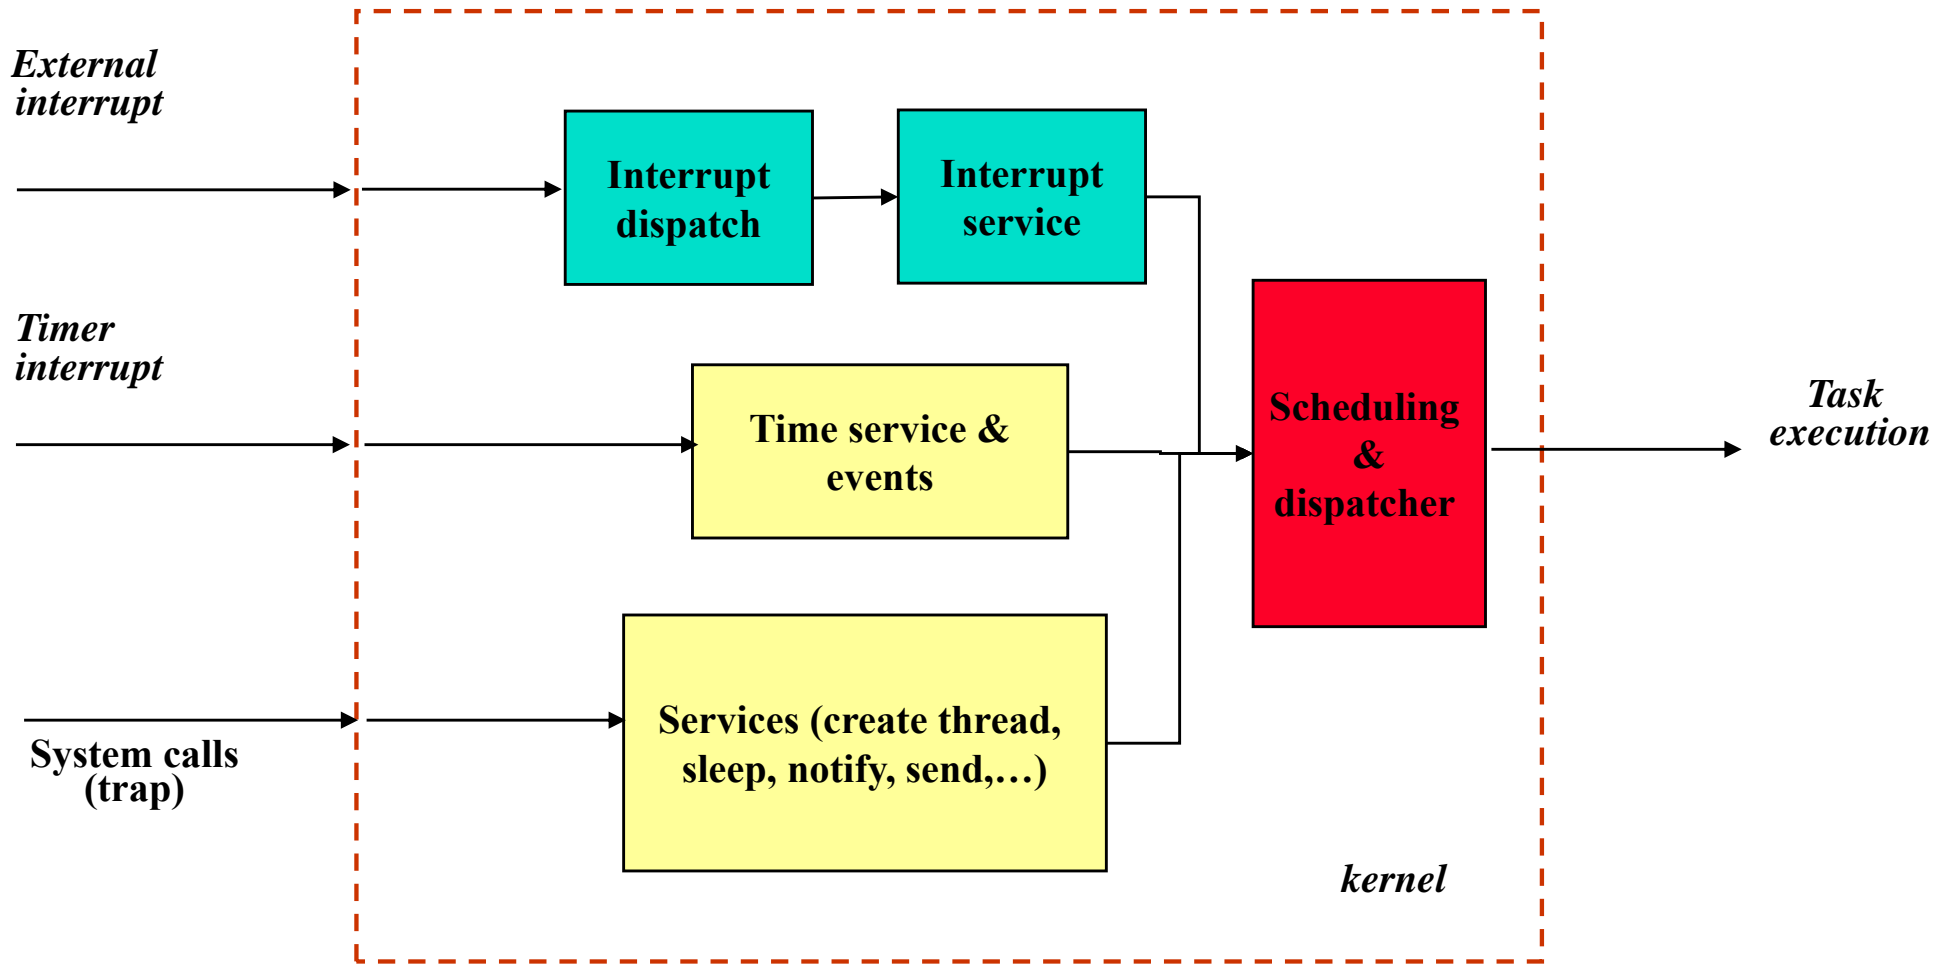
1. The timing behavior of the OS must be *predictable*.

$\forall$  services of the OS: Upper bound on the execution time!

RTOSs must be deterministic (unlike standard Java for example):

- upper bounds on blocking times need to be available, i.e. during which interrupts are disabled,
- almost all activities are controlled by a real-time scheduler.

# Task Management Services



# Real-time Operating Systems

---

## 2. OS must *manage the timing and scheduling*

- OS possibly has to be aware of deadlines; (unless scheduling is done off-line).
- OS must provide precise time services with high resolution.

## 3. The OS must be *fast*

- Practically important.

# Main Functionality of RTOS-Kernels

---

## ► *Task management:*

- Execution of **quasi-parallel tasks** on a processor using processes or threads (lightweight process) by
  - maintaining process states, process queuing,
  - allowing for preemptive tasks (fast context switching) and quick interrupt handling
- CPU **scheduling** (guaranteeing deadlines, minimizing process waiting times, fairness in granting resources such as computing power)
- Process **synchronization** (critical sections, semaphores, monitors, mutual exclusion)
- Inter-process **communication** (buffering)
- Support of a **real-time clock** as an internal time reference

# Task Management

---

## ► *Task synchronization:*

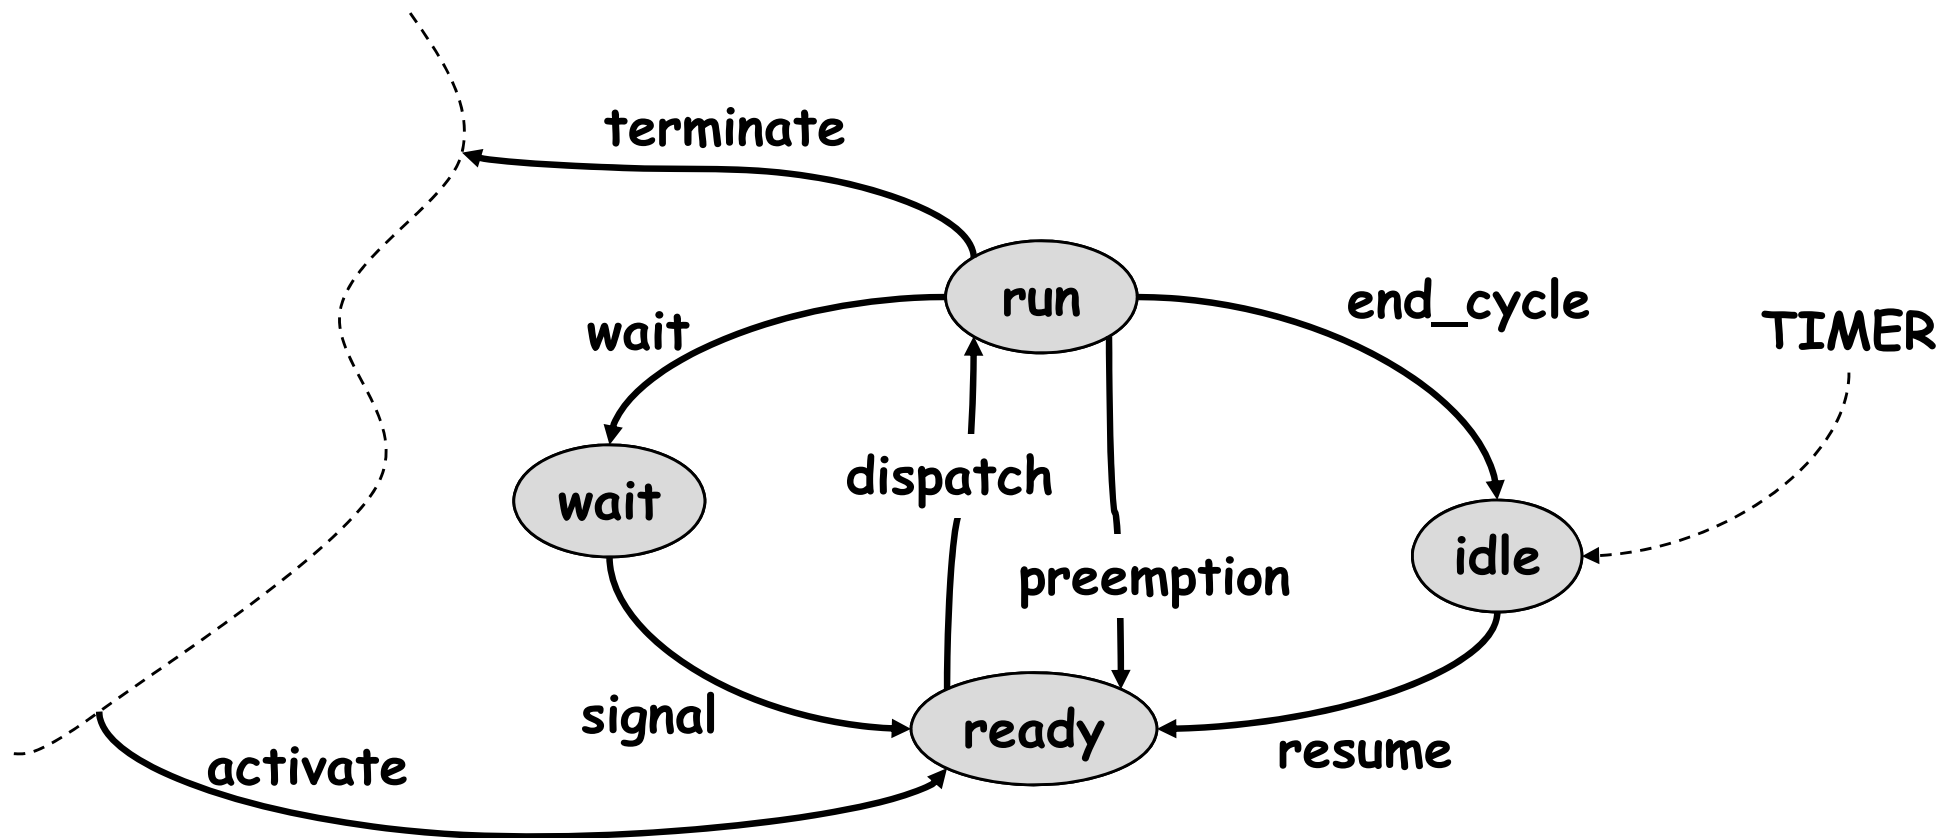
- In classical operating systems, synchronization and mutual exclusion is performed via semaphores and monitors.
- In real-time OS, special semaphores and a deep integration into scheduling is necessary (priority inheritance protocols, ....).

## ► *Further responsibilities:*

- Initializations of internal data structures (tables, queues, task description blocks, semaphores, ...)

# Task States

## ► Minimal Set of Task States:



# Task states

---

- ▶ **Run:**

- A task enters this state as it starts executing on the processor

- ▶ **Ready:**

- State of those tasks that are ready to execute but cannot be executed because the processor is assigned to another task.

- ▶ **Wait:**

- A task enters this state when it executes a synchronization primitive to wait for an event, e.g. a wait primitive on a semaphore. In this case, the task is inserted in a queue associated with the semaphore. The task at the head is resumed when the semaphore is unlocked by a signal primitive.

- ▶ **Idle:**

- A periodic job enters this state when it completes its execution and has to wait for the beginning of the next period.

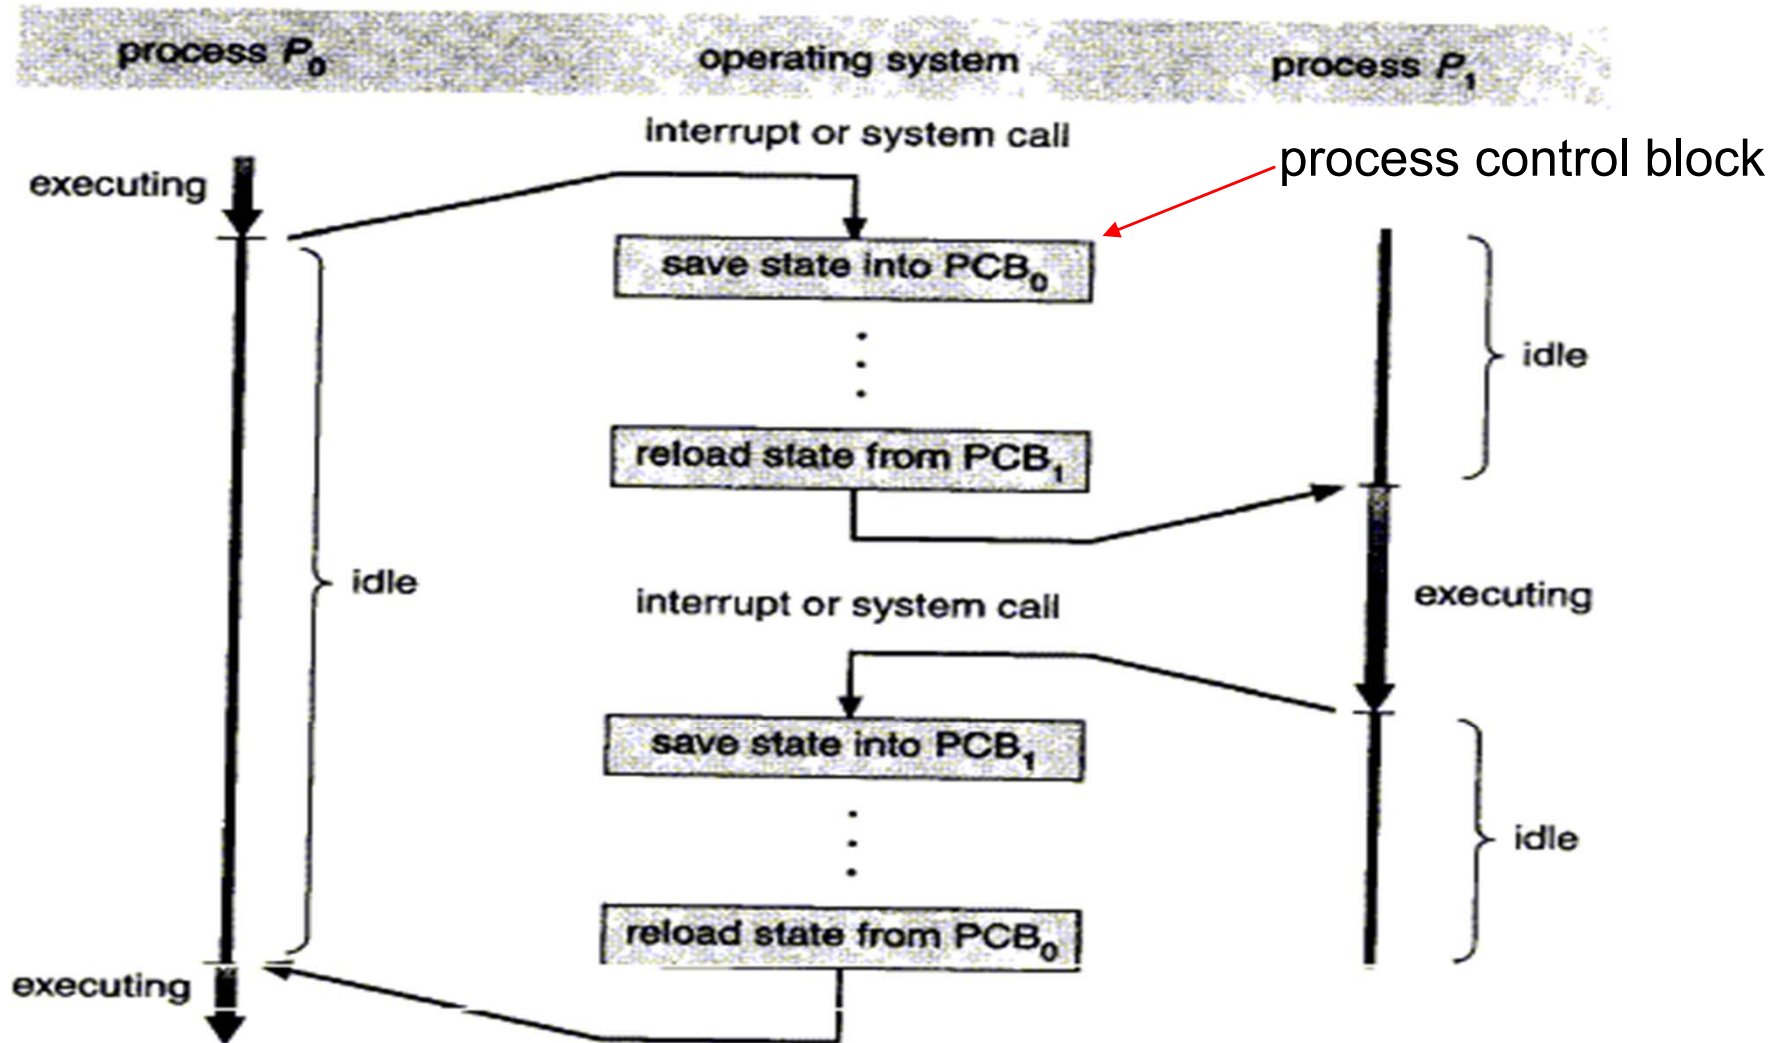


# Threads


---

- ▶ A **thread** is the smallest sequence of programmed instructions that can be managed independently by a scheduler; e.g., a thread is a basic unit of CPU utilization.
- ▶ Multiple threads can exist within the same process and share resources such as memory, while different processes do not share these resources:
  - Typically **shared by threads**: memory.
  - Typically **owned by threads**: registers, stack.
- ▶ **Thread** advantages and characteristics:
  - Faster to switch between threads; switching between user-level threads requires no major intervention by the operating system.
  - Typically, an application will have a separate thread for each distinct activity.
  - Thread Control Block (TCB) stores information needed to manage and schedule a thread

# Context Switching

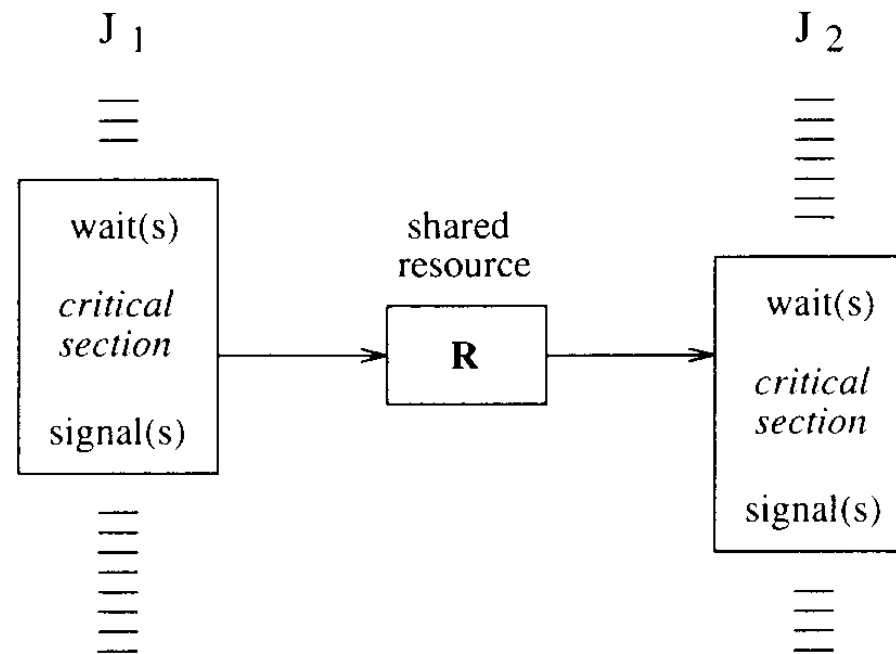


# Multiple Threads within a Process



# Communication Mechanisms

- **Problem:** the use of shared resources for implementing message passing schemes may cause priority inversion and blocking.



# Communication mechanisms

---


## ► *Synchronous communication:*

- Whenever two tasks want to communicate they must be synchronized for a message transfer to take place (*rendez-vous*)
- They have to wait for each other.
- *Problem* in case of dynamic real-time systems: Estimating the maximum blocking time for a process rendez-vous.
- In a *static* real-time environment, the problem can be solved off-line by transforming all synchronous interactions into precedence constraints.

# Communication mechanisms

## ► *Asynchronous communication:*

- Tasks do not have to wait for each other
- The sender just deposits its message into a channel and continues its execution; similarly the receiver can directly access the message if at least a message has been deposited into the channel.
- More suited for real-time systems than synchronous comm.
- **Mailbox:** Shared memory buffer, FIFO-queue, basic operations are send and receive, usually has fixed capacity.
- **Problem:** Blocking behavior if channel is full or empty; alternative approach is provided by cyclical asynchronous buffers.



# Class 1: Fast Proprietary Kernels

---

## *Fast proprietary kernels*

*For hard real-time systems, these kernels are questionable, because they are designed to be fast, rather than to be predictable in every respect*

Examples include

FreeRTOS, QNX, eCOS, RT-LINUX, VxWORKS, LynxOS.



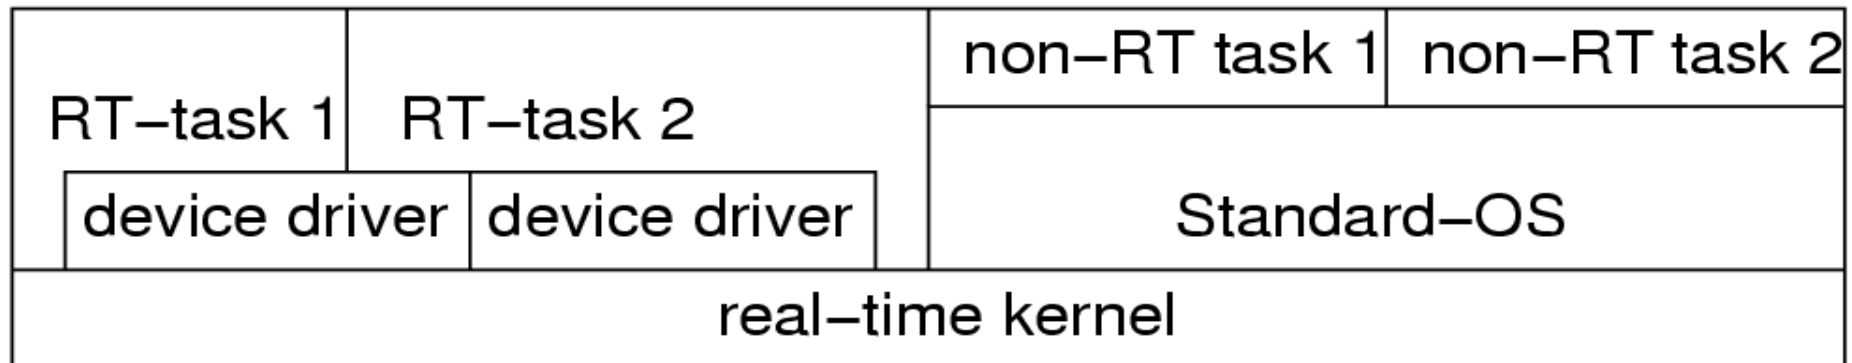
# Class 2: Extensions to Standard OSs

## *Real-time extensions to standard OS:*

Attempt to exploit comfortable main stream OS.

RT-kernel running all RT-tasks.

Standard-OS executed as one task.



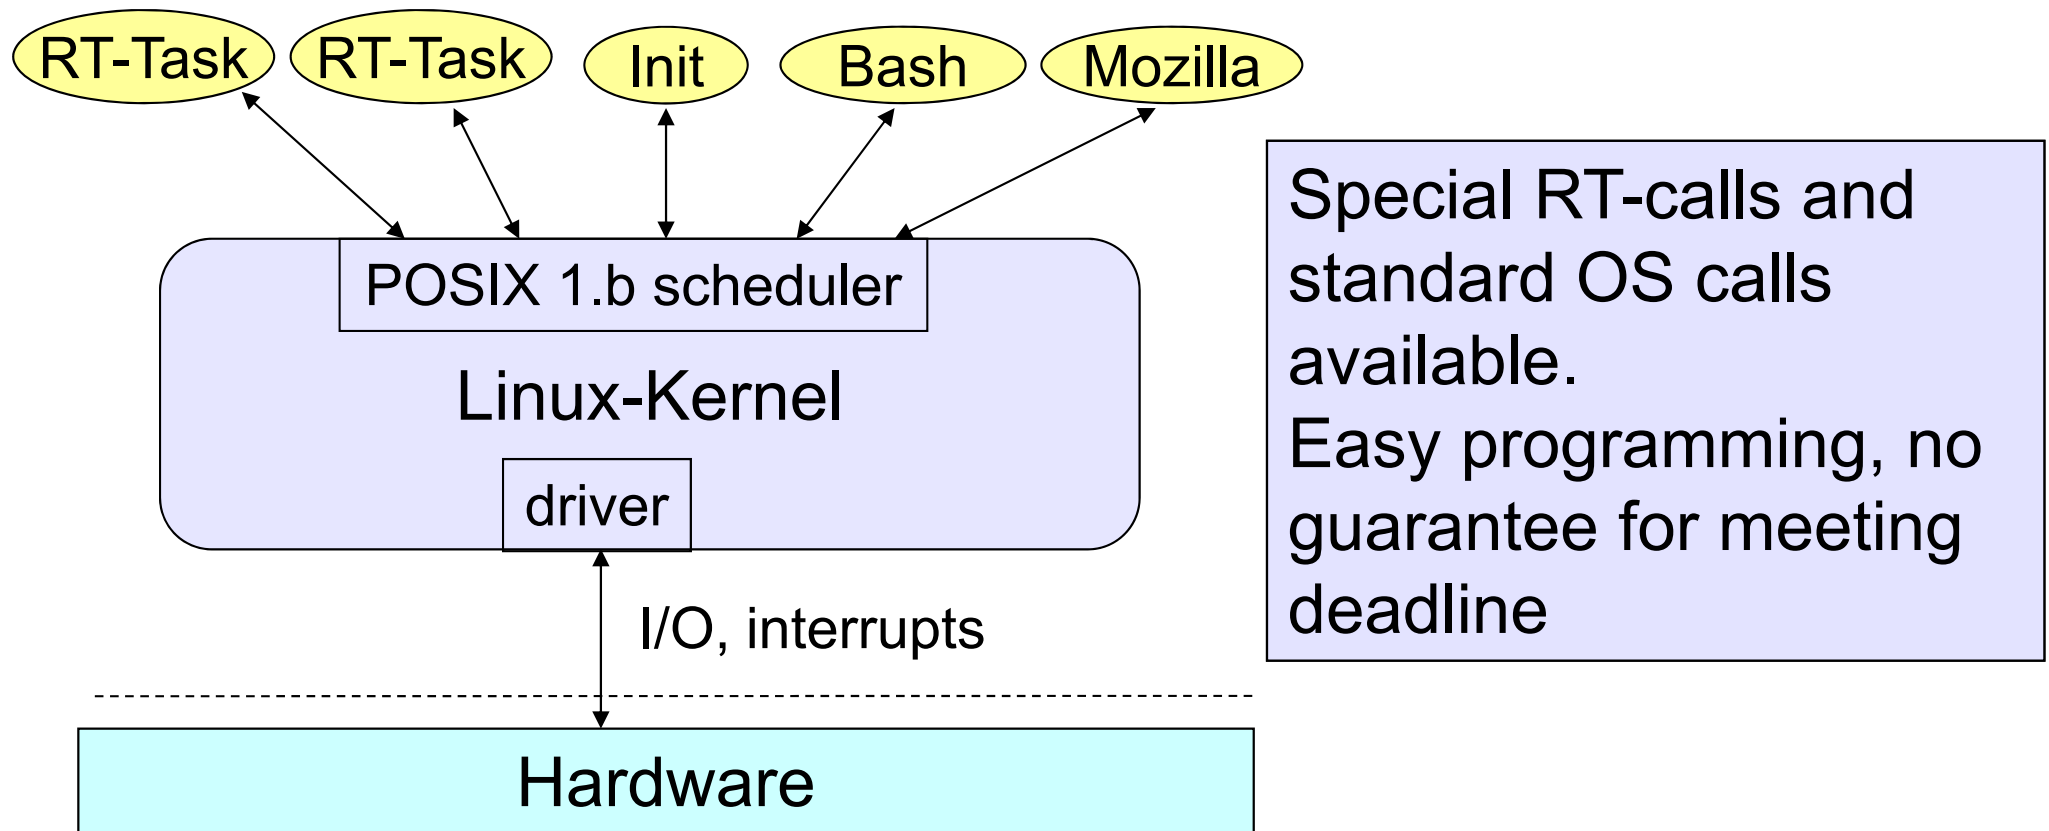
- + Crash of standard-OS does not affect RT-tasks;
- RT-tasks cannot use Standard-OS services;  
less comfortable than expected

revival of the concept:  
hypervisor

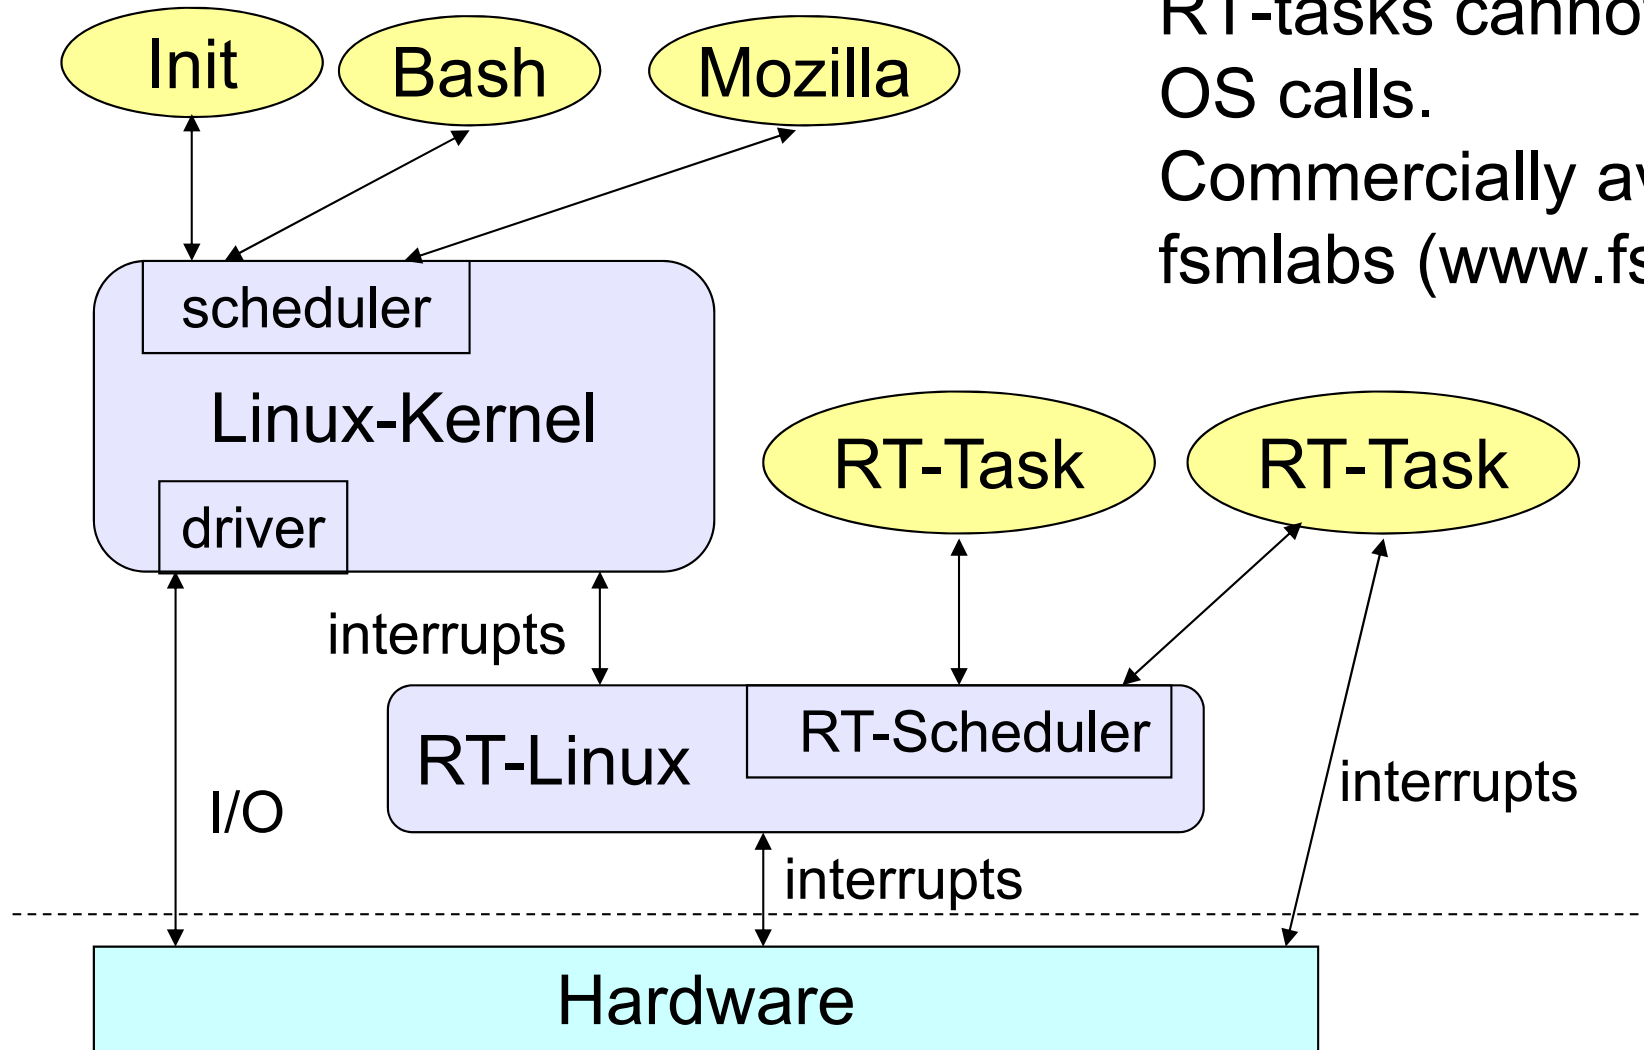


# Example: Posix 1.b RT-extensions to Linux

Standard scheduler can be replaced by POSIX scheduler implementing priorities for RT tasks



# Example: RT Linux



RT-tasks cannot use standard OS calls.  
Commercially available from fsm labs ([www.fsmlabs.com](http://www.fsmlabs.com))

# Class 3: Research Systems

---

## *Research systems trying to avoid limitations:*

- Include L4, seL4, NICTA, ERIKA, SHARK

## *Research issues:*

- low overhead memory protection,
- temporal protection of computing resources
- RTOSes for on-chip multiprocessors
- quality of service (QoS) control (besides real-time constraints)
- formally verified kernel properties

List of current real-time operating systems:

[http://en.wikipedia.org/wiki/Comparison\\_of\\_real-time\\_operating\\_systems](http://en.wikipedia.org/wiki/Comparison_of_real-time_operating_systems)