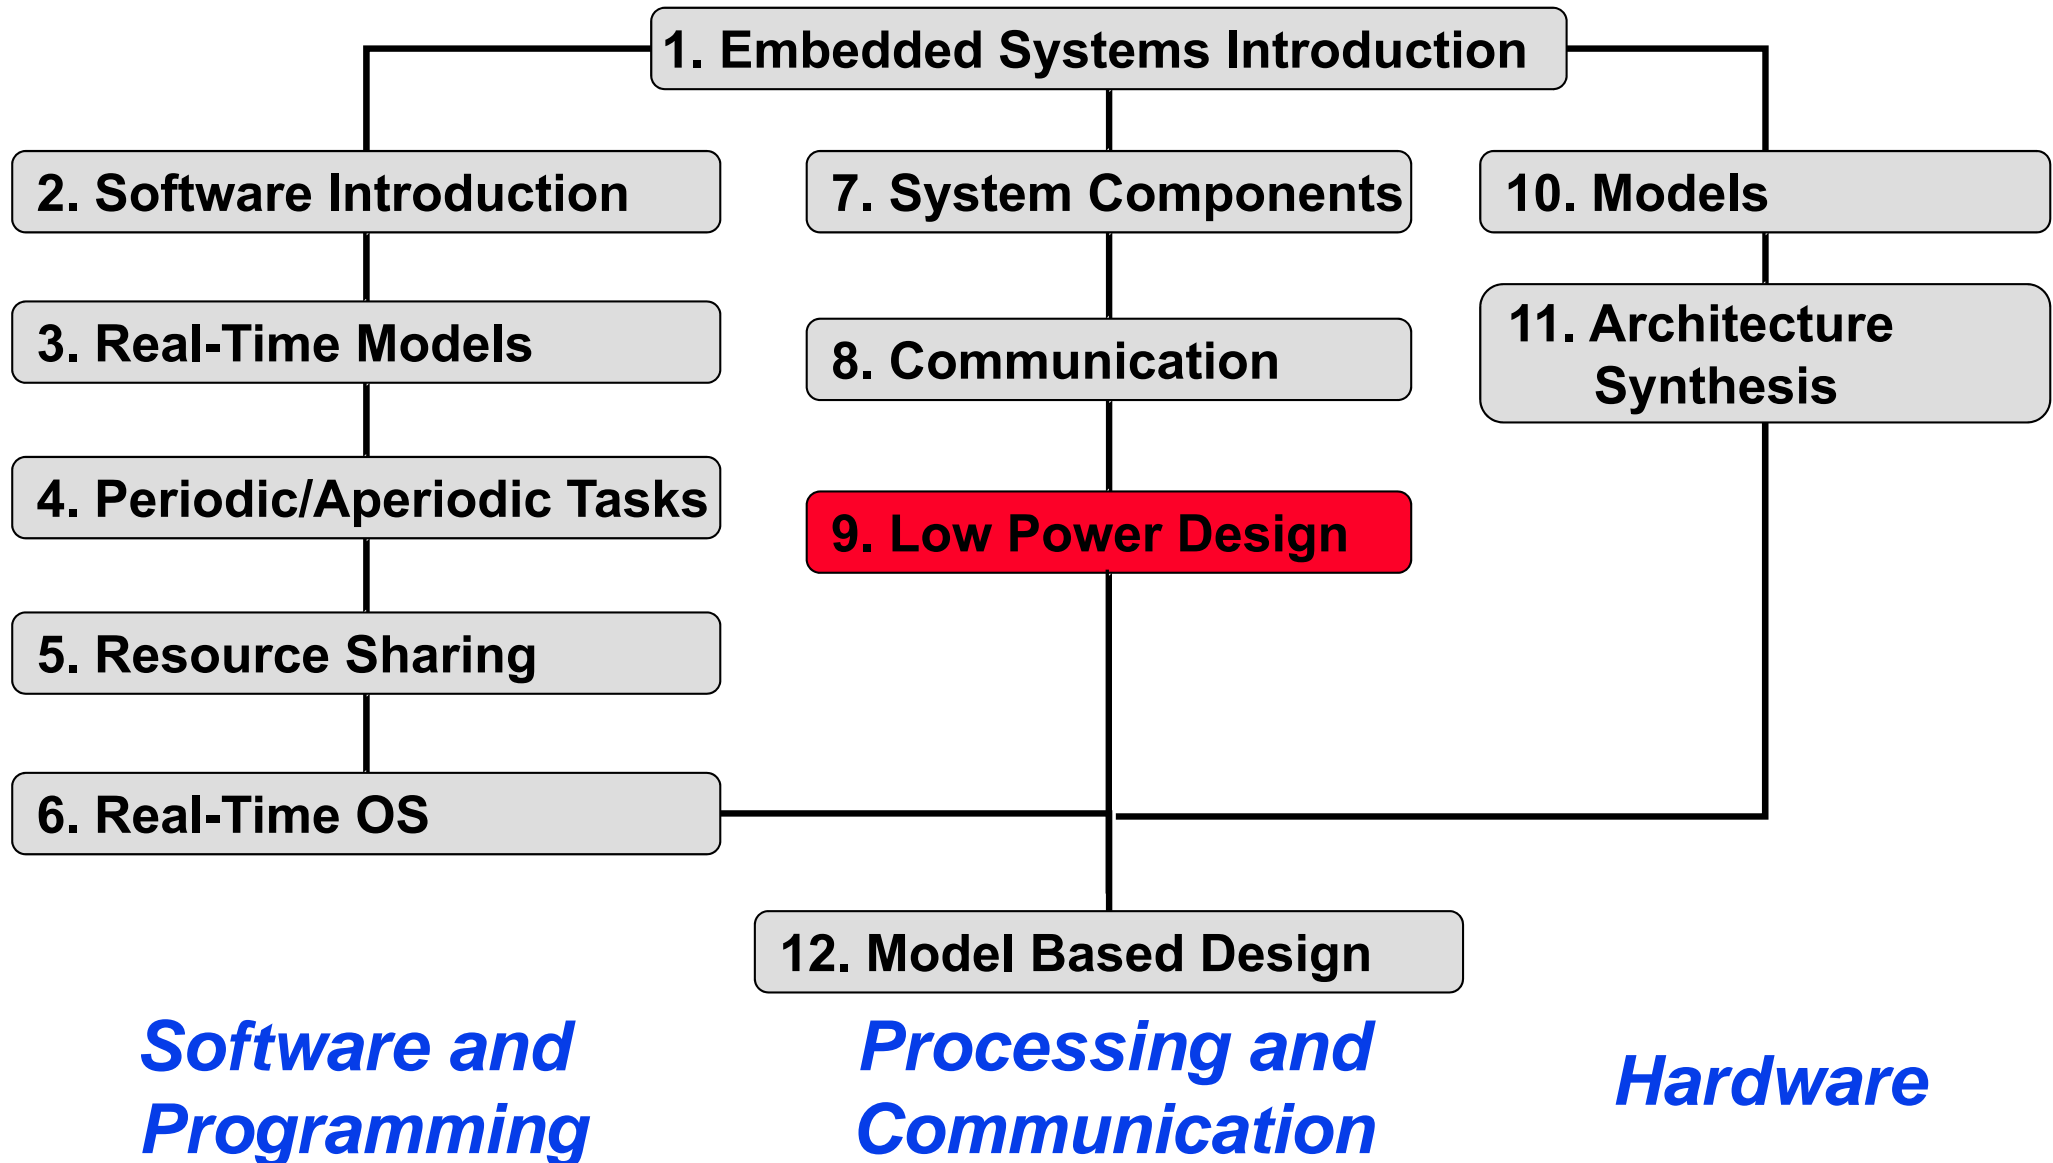


# Embedded Systems

## 9. Low Power Design

Lothar Thiele

# Contents of Course



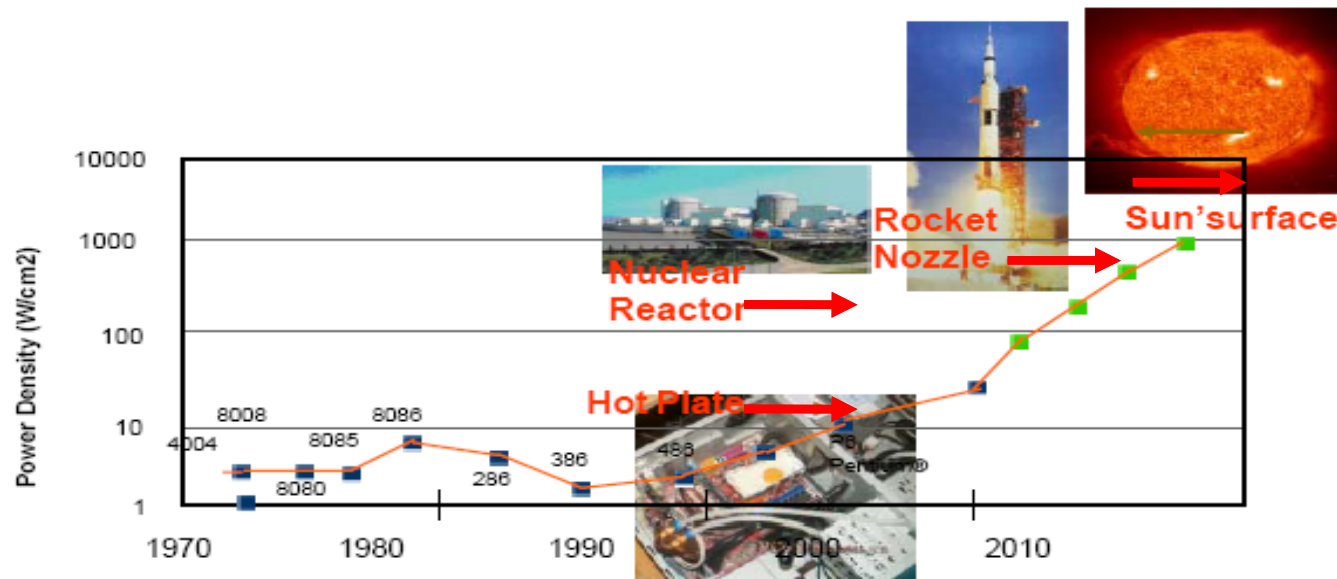
# Topics

---

- ▶ ***General Remarks***
- ▶ Power and Energy
- ▶ Basic Techniques
  - Parallelism
  - VLIW (parallelism and reduced overhead)
  - Dynamic Voltage Scaling
  - Dynamic Power Management

# Power and Energy Consumption

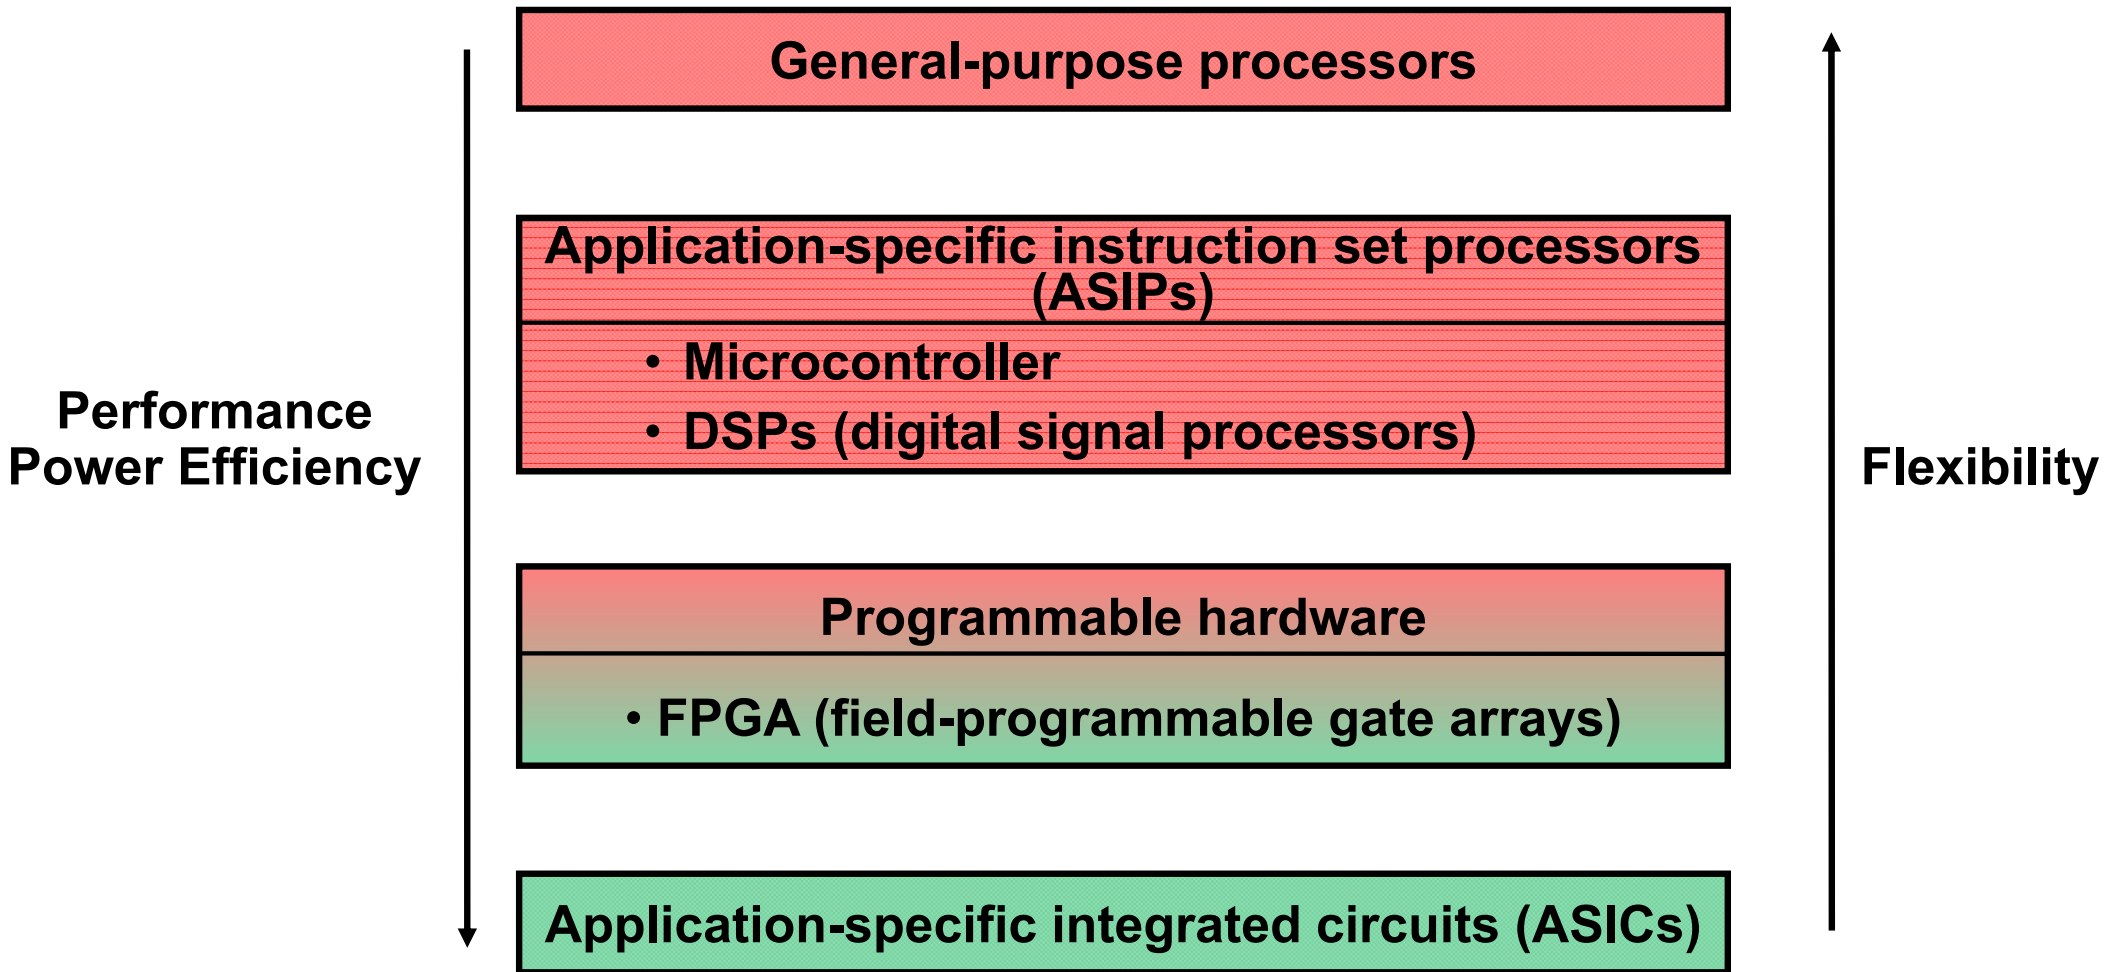
Need for efficiency (power and energy):




***„Power is considered as the most important constraint in embedded systems.“*** [in: L. Eggermont (ed): Embedded Systems Roadmap 2002, STW]

***“Power demands are increasing rapidly, yet battery capacity cannot keep up.”*** [in Diztel et al.: Power-Aware Architecting for data-dominated applications, 2007, Springer]

# Implementation Alternatives




# The Power/Flexibility Conflict



Necessary to **optimize HW and SW**.  
Use **heterogeneous architectures**.  
Apply **specialization techniques**.

# Energy Efficiency



© Hugo De Man,  
IMEC, Philips, 2007

# Topics


---

- ▶ General Remarks
- ▶ *Power and Energy*
- ▶ Basic Techniques
  - Parallelism
  - VLIW (parallelism and reduced overhead)
  - Dynamic Voltage Scaling
  - Dynamic Power Management



# Power and Energy are Related

$$E = \int P(t) dt$$



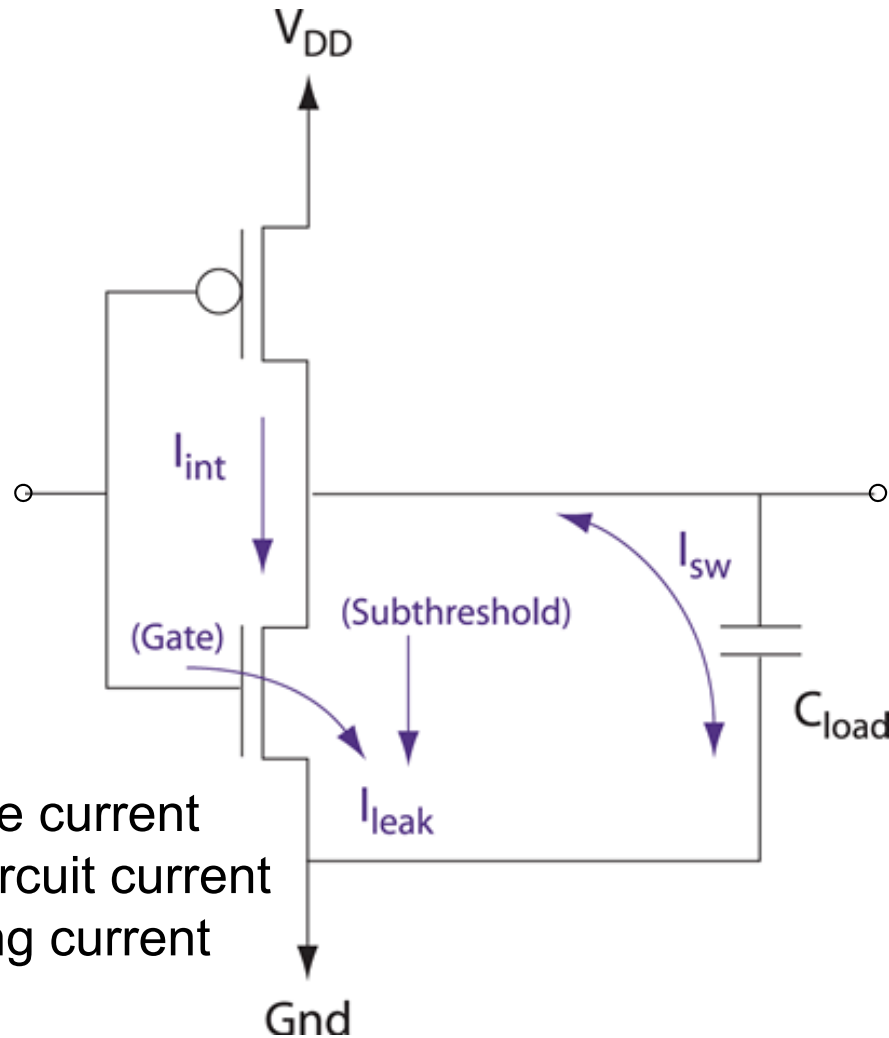
In many cases, faster execution also means less energy, but the opposite may be true if power has to be increased to allow faster execution.

# Low Power vs. Low Energy

---


- ▶ Minimizing the *power consumption* is important for
  - the design of the power supply
  - the design of voltage regulators
  - the dimensioning of interconnect
  - cooling (short term cooling)
    - high cost (estimated to be rising at \$1 to \$3 per Watt for heat dissipation [Skadron et al. ISCA 2003])
    - limited space
- ▶ Minimizing the *energy consumption* is important due to
  - restricted availability of energy (mobile systems)
  - limited battery capacities (only slowly improving)
  - very high costs of energy (solar panels, in space)
  - long lifetimes, low temperatures

# Power Consumption of a CMOS Gate



$I_{leak}$  : leakage current  
 $I_{int}$  : short circuit current  
 $I_{sw}$  : switching current


subthreshold and gate-oxide leakage



# Power Consumption of CMOS Processors

## ► *Main sources:*

- Dynamic power consumption
  - charging and discharging capacitors
- Short circuit power consumption
  - short circuit path between supply rails during switching
- Leakage
  - leaking diodes and translators
  - becomes one of the major factors due to shrinking feature sizes in semiconductor technology



(Micro32 Keynotes by Fred Pollack)

# Dynamic Voltage Scaling (DVS)

**Power consumption of CMOS circuits (ignoring leakage):**

$$P \sim \alpha C_L V_{dd}^2 f$$

$V_{dd}$  : supply voltage

$\alpha$  : switching activity

$C_L$  : load capacity

$f$  : clock frequency

**Delay for CMOS circuits:**

$$\tau \sim C_L \frac{V_{dd}}{(V_{dd} - V_T)^2}$$

$V_{dd}$  : supply voltage

$V_T$  : threshold voltage

$$V_T \ll V_{dd}$$

Decreasing  $V_{dd}$  reduces  $P$  quadratically ( $f$  constant).

The gate delay increases reciprocally with decreasing  $V_{dd}$ .

Maximal frequency  $f_{\max}$  decreases linearly with decreasing  $V_{dd}$ .

# Potential for Energy Optimization: DVS


$$P \sim \alpha C_L V_{dd}^2 f$$

$$E \sim \alpha C_L V_{dd}^2 f t = \alpha C_L V_{dd}^2 (\#cycles)$$

Saving energy for a given task:

- Reduce the supply voltage  $V_{dd}$
- Reduce switching activity  $\alpha$
- Reduce the load capacitance  $C_L$
- Reduce the number of cycles  $\#cycles$

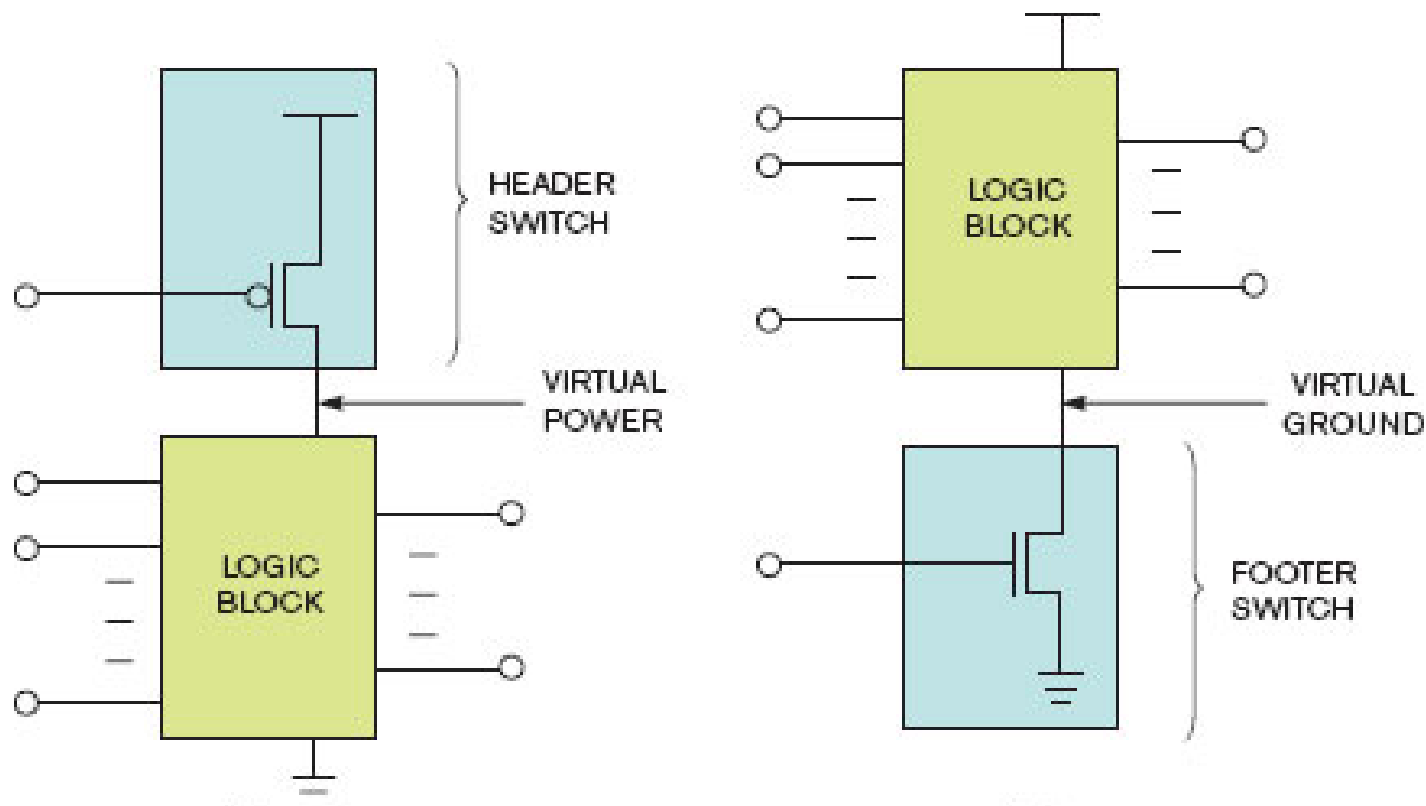
# Example: Voltage Scaling



[Courtesy, Yasuura, 2000]

# Power Supply Gating

- ▶ Power gating is one of the most effective ways of minimizing static power consumption (leakage)
  - Cut-off power supply to inactive units/components
  - Reduces leakage






# Topics

---

- ▶ General Remarks
- ▶ Power and Energy
- ▶ Basic Techniques
  - *Parallelism*
  - VLIW (parallelism and reduced overhead)
  - Dynamic Voltage Scaling
  - Dynamic Power Management


# Use of Parallelism



$$E \sim V_{dd}^2 (\text{\#cycles})$$


$$E_2 = \frac{1}{4} E_1$$

# Use of Pipelining



$$E \sim V_{dd}^2 (\text{\#cycles})$$

$$E_2 = \frac{1}{4} E_1$$



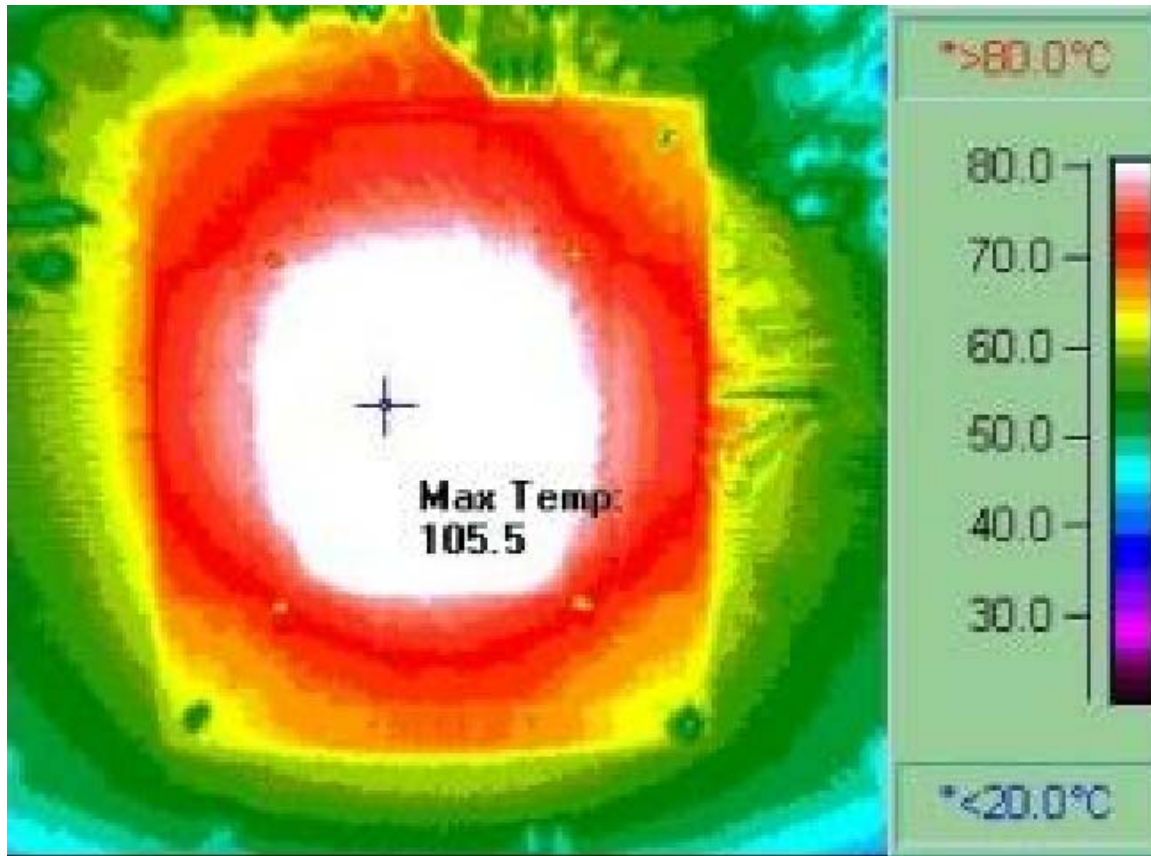
# Topics

---


- ▶ General Remarks
- ▶ Power and Energy
- ▶ Basic Techniques
  - Parallelism
  - *VLIW (parallelism and reduced overhead)*
  - Dynamic Voltage Scaling
  - Dynamic Power Management

# New ideas help ...

Pentium



Crusoe

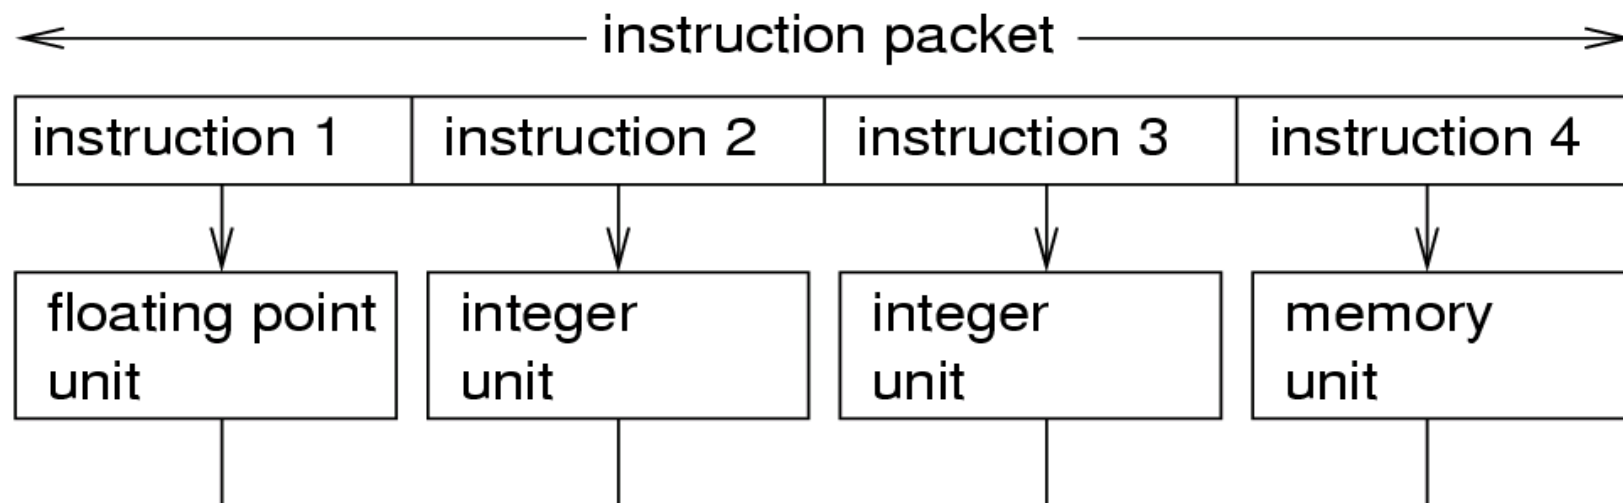


Running the same multimedia application.

As published by Transmeta [www.transmeta.com]

# VLIW Architectures

- ▶ *Large degree of parallelism*
  - many computational units, (deeply) pipelined
- ▶ *Simple hardware architecture*
  - explicit parallelism (parallel instruction set)
  - parallelization is done offline (compiler)




# Transmeta was a typical VLIW Architecture


---

- 128-bit instructions (bundles):
  - 4 operations per instruction
  - 2 combinations of instructions allowed
- Register files
  - 64 integer, 32 floating point
- Some interesting features
  - 6 stage pipeline (2x fetch, decode, register read, execute, write)
  - x86 ISA execution using software techniques
    - Skip the binary compatibility problem!!
    - Interpretation and just-in-time binary translation
  - Speculation support


# Transmeta



e.g. Sun, MIPS, ARM



e.g. Intel, AMD



e.g. Transmeta  
(VLIW)




# Topics


---

- ▶ General Remarks
- ▶ Power and Energy
- ▶ Basic Techniques
  - Parallelism
  - VLIW (parallelism and reduced overhead)
  - ***Dynamic Voltage Scaling***
  - Dynamic Power Management

# Spatial vs. Dynamic Voltage Management



Not all components require same performance.



Required performance may change over time

# Potential for Energy Optimization: DVS

$$P \sim \alpha C_L V_{dd}^2 f$$


$$E \sim \alpha C_L V_{dd}^2 f t = \alpha C_L V_{dd}^2 (\#cycles)$$

Saving energy for a given task:

- Reduce the supply voltage  $V_{dd}$
- Reduce switching activity  $\alpha$
- Reduce the load capacitance  $C_L$
- Reduce the number of cycles  $\#cycles$

# Example: INTEL Xscale


## POWER-PERFORMANCE COMPARISON



OS should schedule distribution of the energy budget.

From Intel's Web Site

# Example: Voltage Scaling




[Courtesy, Yasuura, 2000]

# DVS Example: a) Complete task ASAP

$V_{dd}$ [V]	5.0	4.0	2.5
Energy per cycle [nJ]	40	25	10
$f_{max}$ [MHz]	50	40	25
cycle time [ns]	20	25	40


Task that needs to execute  $10^9$  cycles within 25 seconds.



$$E_a = 10^9 \times 40 \times 10^{-9} \\ = 40 \text{ [J]}$$


# DVS Example: b) Two voltages

$V_{dd}$ [V]	5.0	4.0	2.5
Energy per cycle [nJ]	40	25	10
$f_{max}$ [MHz]	50	40	25
cycle time [ns]	20	25	40



# DVS Example: c) Optimal Voltage


$V_{dd}$ [V]	5.0	4.0	2.5
Energy per cycle [nJ]	40	25	10
$f_{max}$ [MHz]	50	40	25
cycle time [ns]	20	25	40



$$E_c = 10^9 \times 25 \times 10^{-9} = 25 \text{ [J]}$$



# DVS: Optimal Strategy



$$z = a \cdot x + (1-a) \cdot y$$

Execute task in fixed time  $T$  with variable voltage  $V_{dd}(t)$ :


gate delay:  $\tau \sim \frac{1}{V_{dd}}$

execution rate:  $f(t) \sim V_{dd}(t)$

invariant:  $\int V_{dd}(t) dt = \text{const.}$

- **case A**: execute at voltage  $x$  for  $T \cdot a$  time units and at voltage  $y$  for  $(1-a) \cdot T$  time units;  
energy consumption  $T \cdot (P(x) \cdot a + P(y) \cdot (1-a))$
- **case B**: execute at voltage  $z = a \cdot x + (1-a) \cdot y$  for  $T$  time units;  
energy consumption  $T \cdot P(z)$

# DVS: Optimal Strategy




- ▶ If possible, running at a constant frequency (voltage) minimizes the energy consumption for dynamic voltage scaling:
  - **case A** is always worse if the power consumption is a convex function of the supply voltage

# DVS: Offline Scheduling on One Processor

- ▶ Let us model a set of independent tasks as follows:
  - We suppose that a task  $v_i \in V$ 
    - requires  $c_i$  computation time at normalized processor frequency 1
    - arrives at time  $a_i$
    - has (absolute) deadline constraint  $d_i$
- ▶ How do we schedule these tasks such that all these tasks can be finished **no later than their deadlines** and the energy consumption is **minimized**?
  - YDS Algorithm from “A Scheduling Model for Reduce CPU Energy”, Frances Yao, Alan Demers, and Scott Shenker, FOCS 1995.”

If possible, running at a constant frequency (voltage) minimizes the energy consumption for dynamic voltage scaling.

# YDS Algorithm for Offline Scheduling




- Define **intensity**  $G([z, z'])$  in some time interval  $[z, z']$ :
  - average accumulated execution time of all tasks that have arrival and deadline in  $[z, z']$  relative to the length of the interval  $z' - z$

$$V'([z, z']) = \{v_i \in V : z \leq a_i < d_i \leq z'\}$$

$$G([z, z']) = \sum_{v_i \in V'([z, z'])} c_i / (z' - z)$$

# YDS Algorithm for Offline Scheduling

- **Step 1:** Execute jobs in the interval with the highest intensity by using the earliest-deadline first schedule and running at the intensity as the frequency.



3,6,5

2,6,3

0,8,2

6,14,6

10,14,6

11,17,2

12,17,2

$a_i, d_i, c_i$

$$G([0,6]) = (5+3)/6=8/6, G([0,8]) = (5+3+2)/(8-0) = 10/8,$$

$$G([0,14]) = (5+3+2+6+6)/14=11/7, G([0,17]) = (5+3+2+6+6+2+2)/17=26/17$$

$$G([2,6]) = (5+3)/(6-2)=2, G([2,14]) = (5+3+6+6)/(14-2) = 5/3,$$

$$G([2,17]) = (5+3+6+6+2+2)/15=24/15$$


$$G([3,6]) = 5/3, G([3,14]) = (5+6+6)/(14-3) = 17/11, G([3,17]) = (5+6+6+2+2)/14=21/14$$

$$G([6,14]) = 12/(14-6)=12/8, G([6,17]) = (6+6+2+2)/(17-6)=16/11$$

$$G([10,14]) = 6/4, G([10,17]) = 10/7, G([11,17]) = 4/6, G([12,17]) = 2/5$$




# YDS Algorithm for Offline Scheduling

- **Step 1:** Execute jobs in the interval with the highest intensity by using the earliest-deadline first schedule and running at the intensity as the frequency.



# YDS Algorithm for Offline Scheduling

- **Step 2:** Adjust the arrival times and deadlines by excluding the possibility to execute at the previous critical intervals.




0,8,2	0,4,2
6,14,6	2,10,6
10,14,6	6,10,6
11,17,2	7,13,2
12,17,2	8,13,2

$a_i, d_i, c_i$

# YDS Algorithm for Offline Scheduling



- **Step 3:** Run the algorithm for the revised input again



$$G([0,4])=2/4, G([0,10]) = 14/10, G([0,13])=18/13$$

$$G([2,10])=12/8, G([2,13]) = 16/11, G([6,10])=6/4$$

$$G([6,13])=10/7, G([7,13])=4/6, G([8,13])=4/5$$



0,4,2

2,10,6

6,10,6

7,13,2



8,13,2

$a_i, d_i, c_i$




# YDS Algorithm for Offline Scheduling

- **Step 3:** Run the algorithm for the revised input again
- **Step 4:** Put pieces together



	$V_1$	$V_2$	$V_3$	$V_4$	$V_5$	$V_6$	$V_7$
frequency	2	2	1	1.5	1.5	4/3	4/3

# DVS: Online Scheduling on One Processor



3,6,5
2,6,3
0,8,2
6,14,6
10,14,6
11,17,2
12,17,2
$a_i, d_i, c_i$

## ► Continuously update to the best schedule for all arrived tasks

- Time 0: task  $v_3$  is executed at  $2/8$
- Time 2: task  $v_2$  arrives
  - $G([2,6]) = 3/4$ ,  $G([2,8]) = 4.5/6 = 3/4 \Rightarrow$  execute  $v_2$  at  $3/4$
- Time 3: task  $v_1$  arrives
  - $G([3,6]) = (5+3-3/4)/3 = 29/12$ ,  $G([3,8]) < G([3,6]) \Rightarrow$  execute  $v_2$  and  $v_1$  at  $29/12$
- Time 6: task  $v_4$  arrives
  - $G([6,8]) = 1.5/2$ ,  $G([6,14]) = 7.5/8 \Rightarrow$  execute  $v_3$  and  $v_4$  at  $15/16$
- Time 10: task  $v_5$  arrives
  - $G([10,14]) = 39/16 \Rightarrow$  execute  $v_4$  and  $v_5$  at  $39/16$
- Time 11 and Time 12
  - The arrival of  $v_6$  and  $v_7$  does not change the critical interval
- Time 14:
  - $G([14,17]) = 4/3 \Rightarrow$  execute  $v_6$  and  $v_7$  at  $4/3$

# Remarks on YDS Algorithm

---

## ► Offline

- The algorithm guarantees the minimal energy consumption while satisfying the timing constraints
- The time complexity is  $O(N^3)$ , where  $N$  is the number of tasks in  $V$ 
  - Finding the critical interval can be done in  $O(N^2)$
  - The number of iterations is at most  $N$
- Exercise:
  - For periodic real-time tasks with deadline=period, running at **constant speed with 100% utilization** under EDF has minimum energy consumption while satisfying the timing constraints.

## ► Online

- Compared to the optimal offline solution, the on-line schedule uses at most 27 times of the minimal energy consumption.


# Topics

---

- ▶ General Remarks
- ▶ Power and Energy
- ▶ Basic Techniques
  - Parallelism
  - VLIW (parallelism and reduced overhead)
  - Dynamic Voltage Scaling
  - *Dynamic Power Management*

# Transmeta™ LongRun™ Power Management

In Action



# Dynamic Power Management (DPM)

Dynamic Power management tries to assign optimal **power saving states**

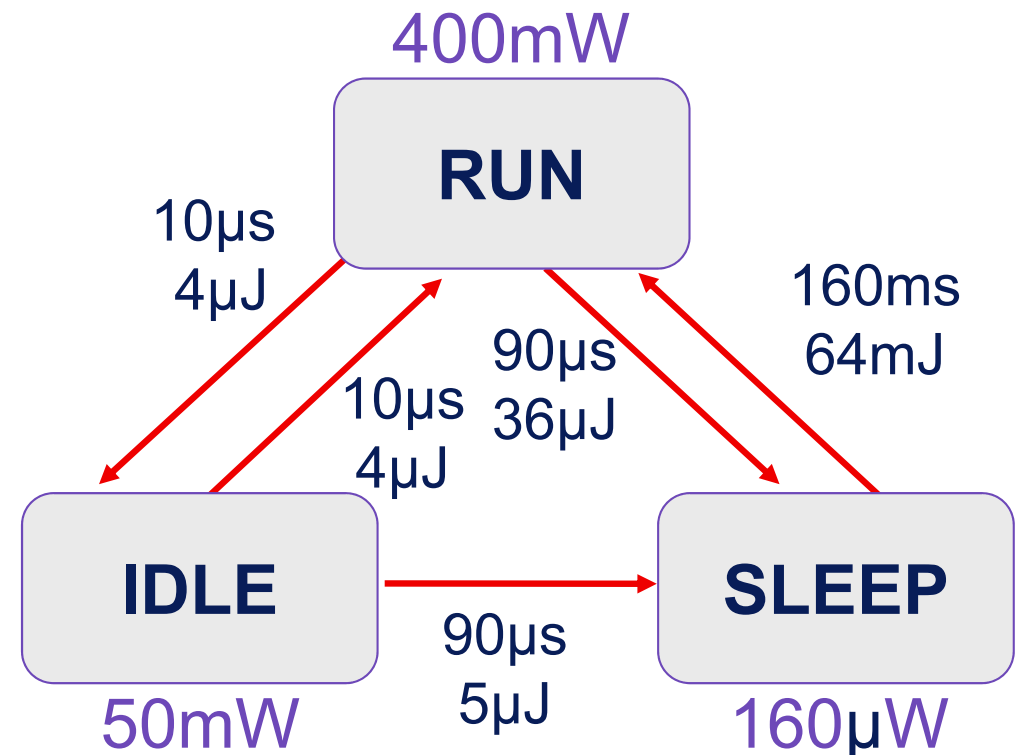
Requires Hardware Support

Example: StrongARM SA1100

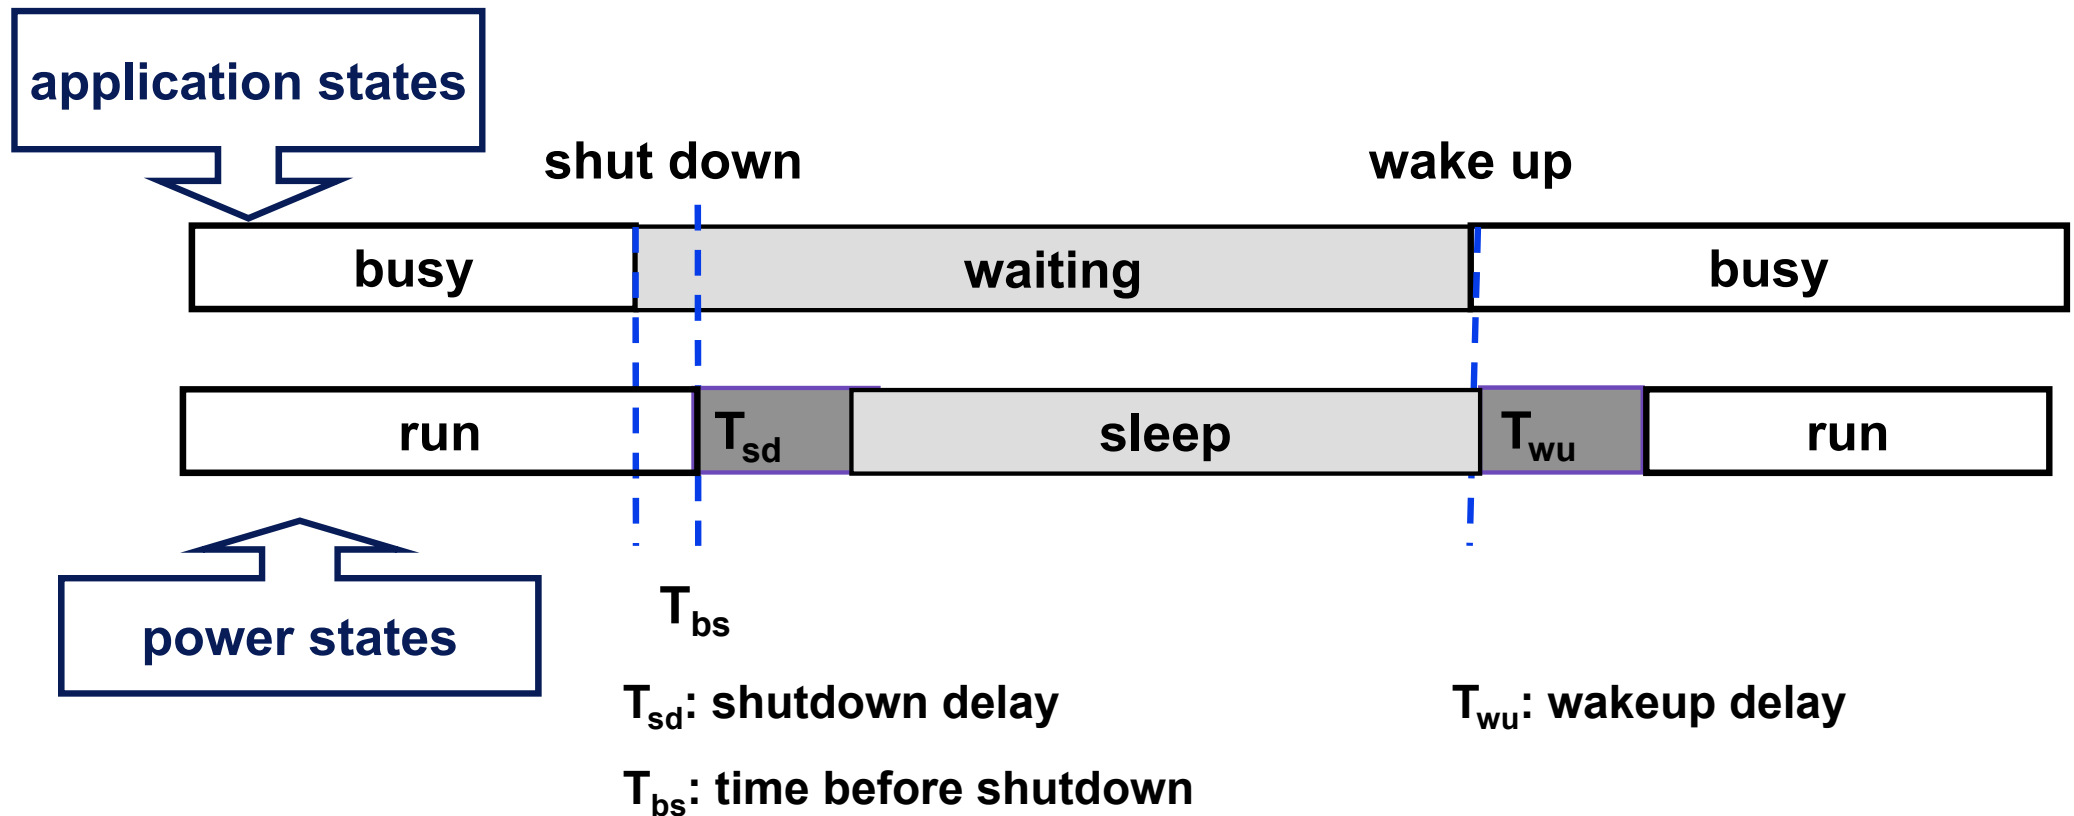
**RUN**: operational

**IDLE**: a SW routine may stop the CPU when not in use, while monitoring interrupts

**SLEEP**: Shutdown of on-chip activity



# Reduce Power According to Workload



Desired: Shutdown only during long idle times  
→ Tradeoff between savings and overhead

# The Challenge

---

Questions:

- When to go to a power-saving state?
- Is an idle period long enough for shutdown?


→ Predicting the future



# Combining DVFS and DPM

## DVS Critical frequency (voltage):


- Running at any frequency/voltage lower than this frequency is not worthwhile for execution.



(Micro32 Keynotes by Fred Pollack)




## power during „run task“ using voltage and frequency scaling




assume constant;  
usually more complex

## energy for executing task



critical voltage

# Procrastination Schedule



- ▶ Execute by using voltages higher or equal to the critical voltage only
  - apply YDS algorithm
  - round up voltages lower than the critical voltage
- ▶ Procrastinate the execution of tasks to aggregate enough time for sleeping
  - Try to reduce the number of times to turn on/off
  - Sleep as long as possible

# Operating System Services

- = Hardware power management
- = Application Software
- = RTOS Power Management Framework

