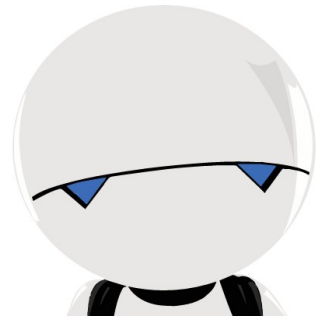


# Deep Learning: CNN

Prof. Marios Savvides

Some slides borrowed from Princeton Vision Group



# 10 BREAKTHROUGH TECHNOLOGIES 2013

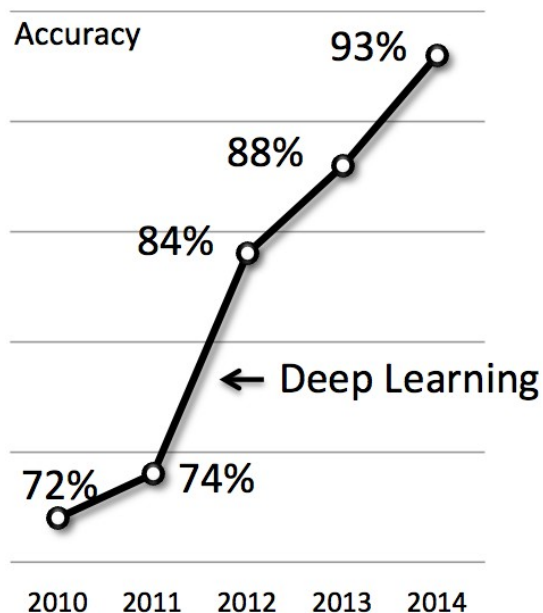
## Deep Learning

With massive amounts of computational power, machines can now recognize objects and translate speech in real time. Artificial intelligence is finally getting smart.



# Deep Learning

IMGENET



*Delving Deep into Rectifiers: Surpassing Human-Level Performance on ImageNet Classification*



Microsoft

95.06%, Feb 06, 2015

*Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariant Shift*



95.18%, Feb 11, 2015

*Deep Image: Scaling up Image Recognition*



95.42%, May 11, 2015

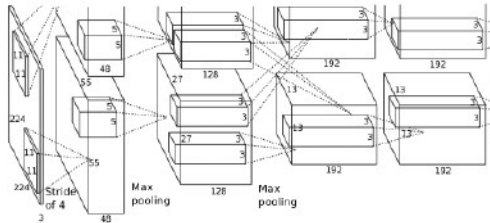
# Deep Learning Recipe

The Deep Learning “Computer Vision Recipe”



Big Data: ImageNet

+



Deep Convolutional Neural Network

+



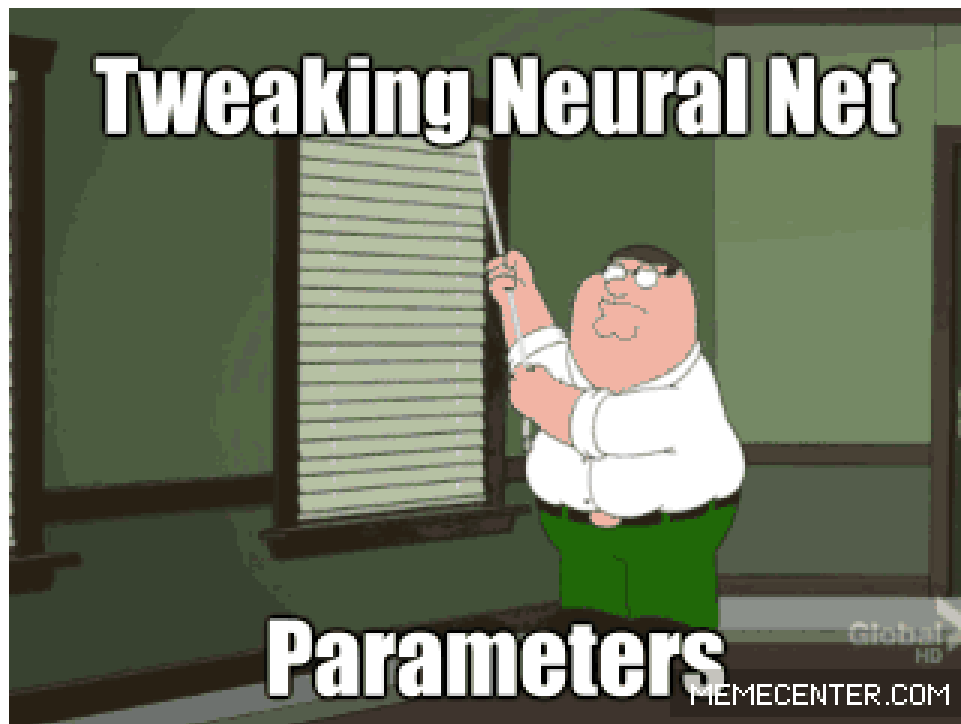
Backprop on GPU

=



Learned Weights

Meanwhile...



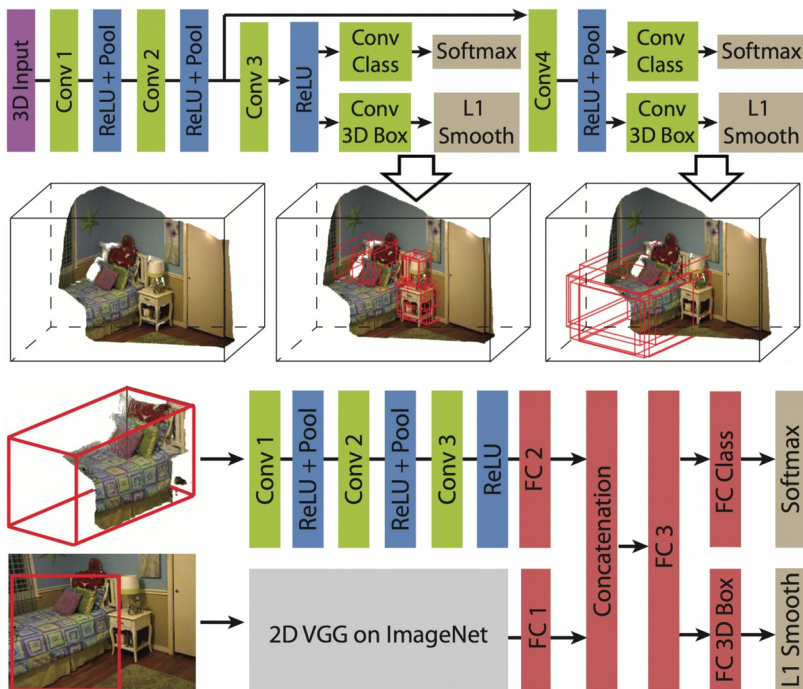
# Big Data + Deep Learning



Big Data is the headache;  
deep learning is the solution.

- Steve Jurvetson

# Successful Case: 3D Object Detector



Improve over 2D deep learning by 13.8 in mAP

Deep Sliding Shapes for Amodal 3D Object Detection in RGB-D Images. Song and Xiao, 2015.

<http://dss.cs.princeton.edu>

# Deep Learning

## **Deep Neural Networks**

### **Convolutional Neural Networks**

Deep Belief Networks

Convolutional Deep Belief Networks

Deep Boltzmann Machines

Stacked (Denoising) Auto-Encoders

Deep Stacking Networks

Tensor Deep Stacking Networks (T-DSN)

Spike-and-Slab RBMs (ssRBMs)

Compound Hierarchical-Deep Models

Deep Coding Networks

Multilayer Kernel Machine

Deep Q-Networks

Memory Networks

Long short-term memory

Semantic Hashing

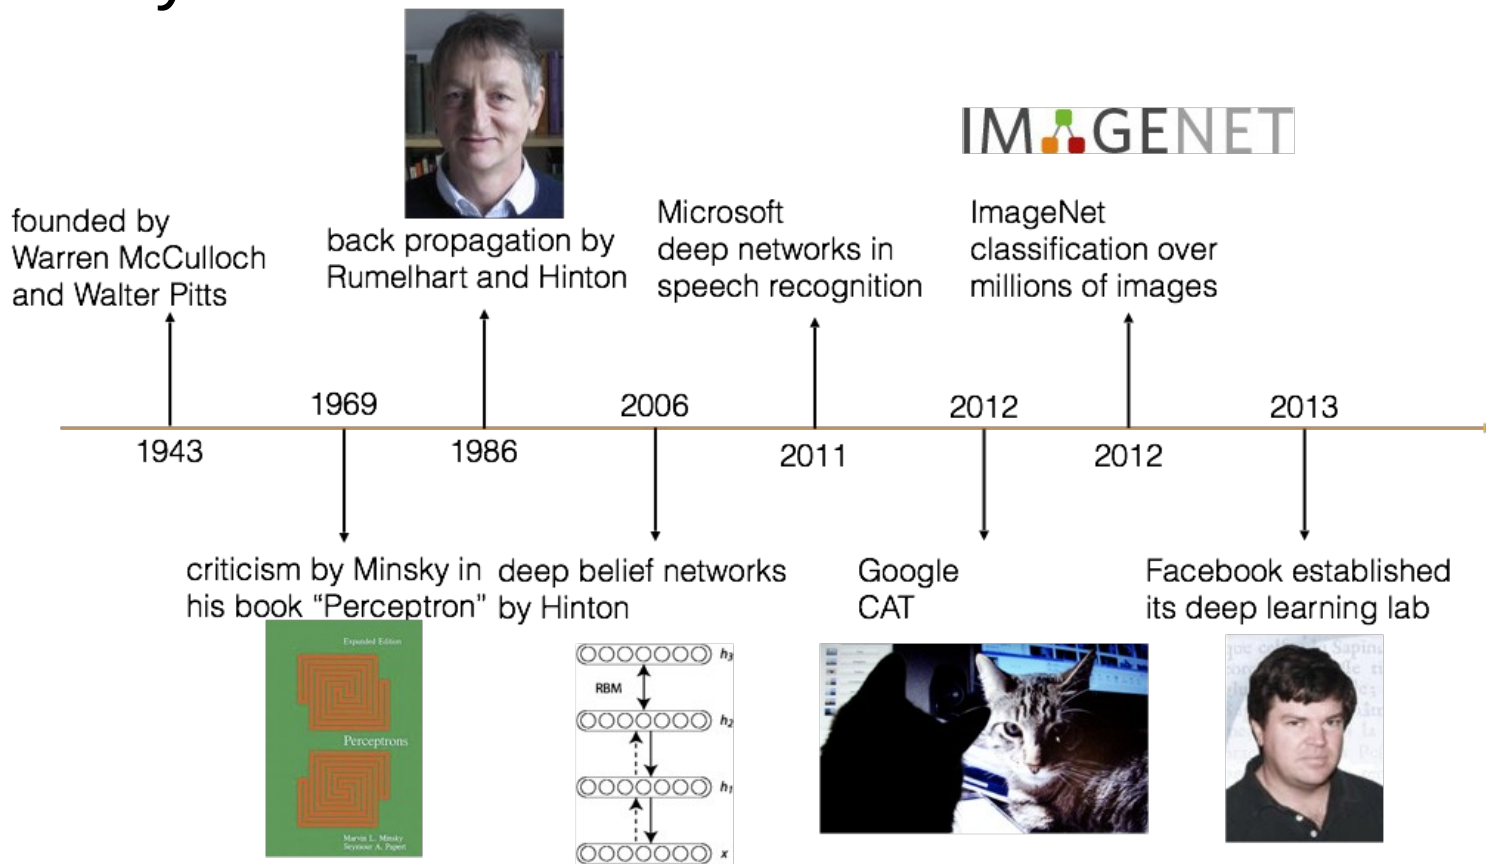
Neural Turing Machines

Memory Networks

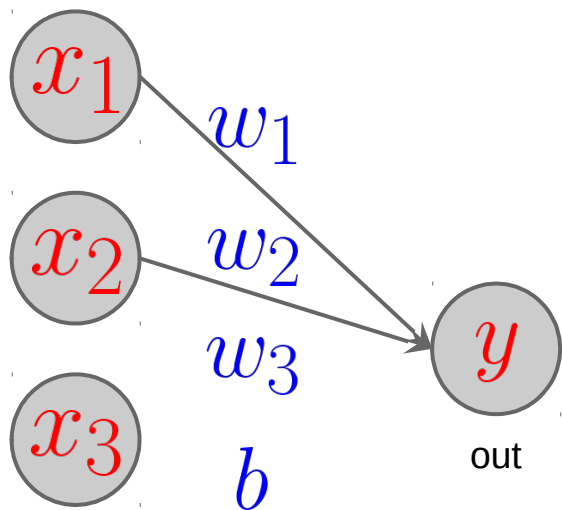
Encoder–Decoder networks



# History



# Neural Network in 60 Seconds



Operations are called layers

Results (in & out) are called responses (aka activation)



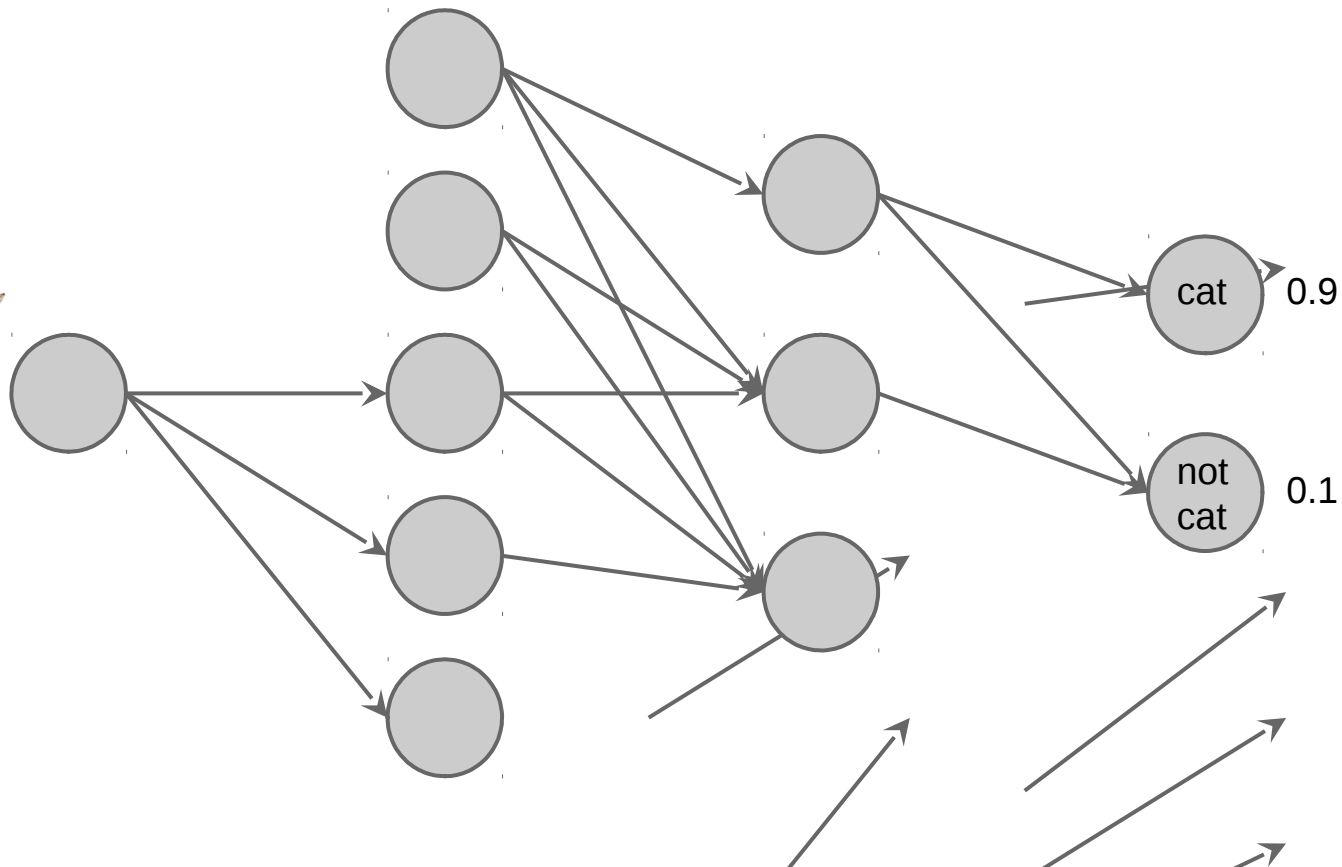
$$y = \sum_i w_i x_i + b$$

$$y = f(x)$$

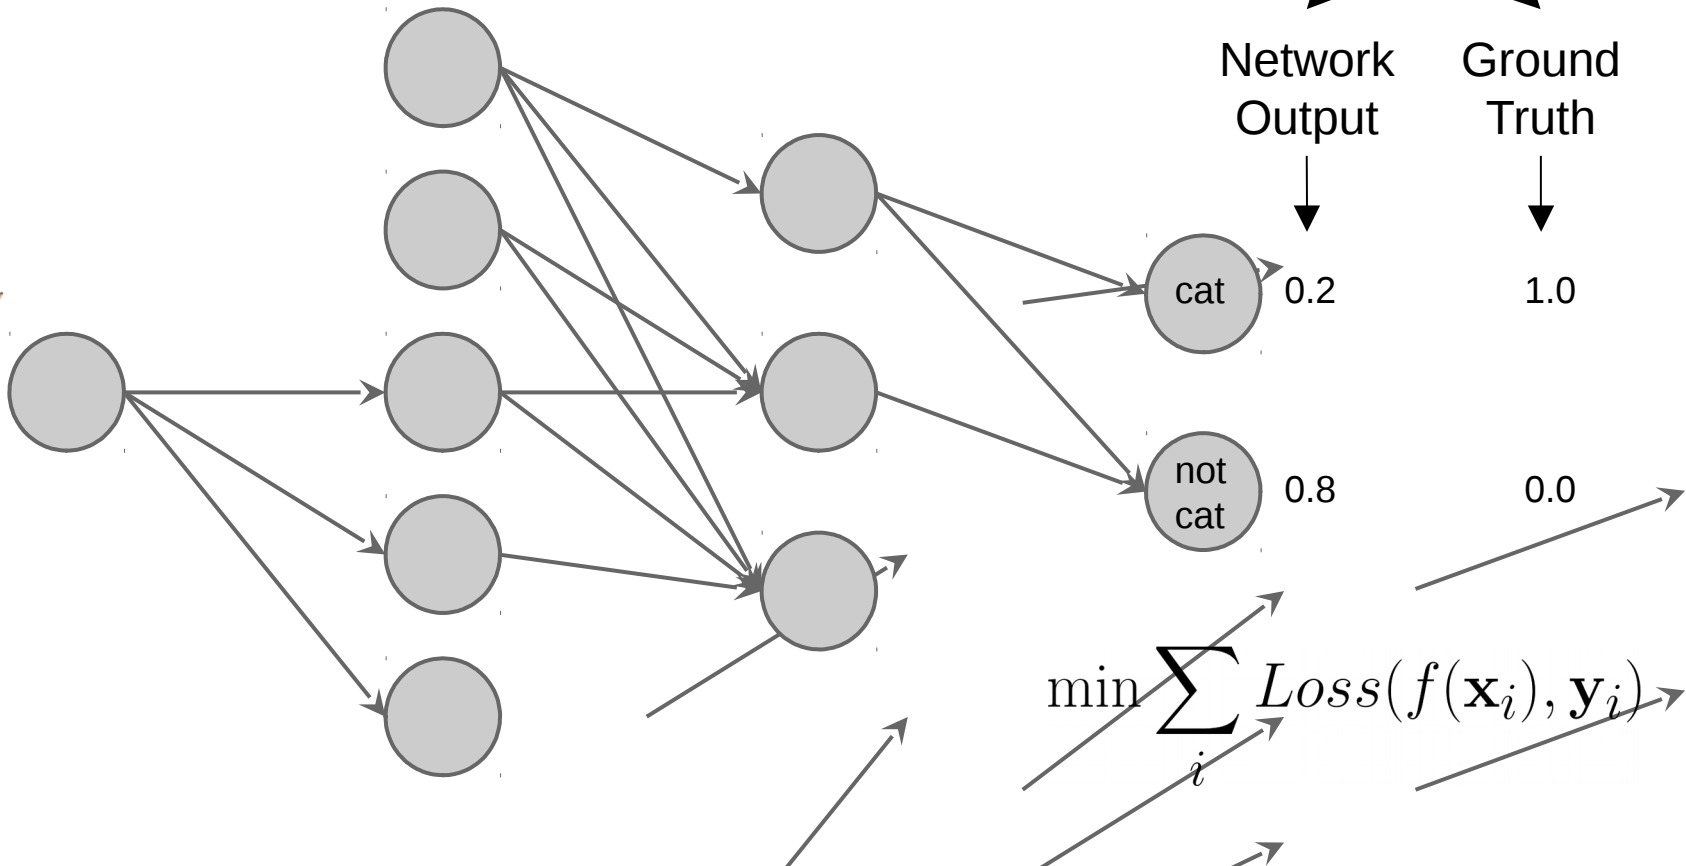
any (almost) differentiable function



# A Good Network



Oops!



# Evolution: Random Search



Evolutionary adversarial training: the discriminator is the predator. This seaweed-looking fish is the product of the generator. The generator is the biochemical machinery of life, fed with randomly-sampled genome. The problem is that the fish doesn't get to see the gradient back-propagated through the predator. The only way to evaluate the gradient is to modify your genes and see if you get eaten.

- Yann LeCun

# Training: Searching for the best weights

$$\min_{\text{weights}} \sum_i \text{Loss}(\text{weights}, \text{data}_i)$$

We need to solve

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} L(\mathbf{w})$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_i L(\mathbf{w}, \mathbf{d}_i)$$

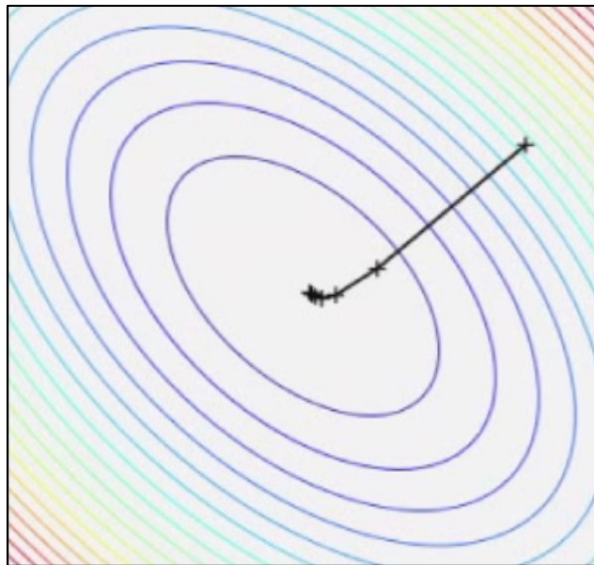
by gradient descent

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \lambda \frac{\partial L}{\partial \mathbf{w}}(\mathbf{w}_t)$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t - \lambda \sum_i \frac{\partial L}{\partial \mathbf{w}}(\mathbf{w}_t, \mathbf{d}_i)$$

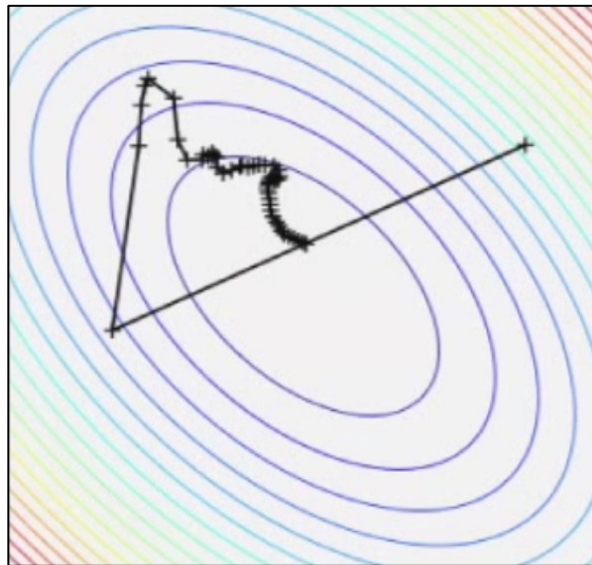
$$\sum_i \frac{\partial L}{\partial \mathbf{w}}(\mathbf{w}_t, \mathbf{d}_i) = \sum_i g(\mathbf{d}_i)$$

Gradient Descent



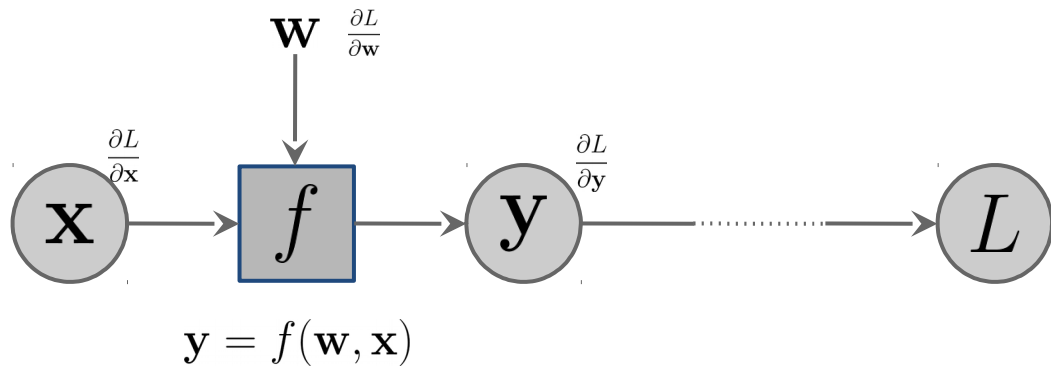
sum over all training data

Stochastic Gradient Descent (SGD)



sum over a random subset of training data  
(called mini-batch). different every iteration.

# Training by Backpropagation (a.k.a. chain rule)



$$\frac{\partial L}{\partial \mathbf{w}} = \sum_i \frac{\partial L}{\partial \mathbf{y}_i} \frac{\partial \mathbf{y}_i}{\partial \mathbf{w}} = \sum_i \frac{\partial L}{\partial \mathbf{y}_i} f'_1(\mathbf{w}_t, \mathbf{x}_i)$$

$$\frac{\partial L}{\partial \mathbf{x}_i} = \frac{\partial L}{\partial \mathbf{y}_i} \frac{\partial \mathbf{y}_i}{\partial \mathbf{x}_i} = \frac{\partial L}{\partial \mathbf{y}_i} f'_2(\mathbf{w}_t, \mathbf{x}_i)$$



# Training

for many iterations:

1. randomly sample a mini-batch
2. forward to compute responses
3. backward to compute gradients (for both weights and responses)
4. update weights with the gradients

# Convolutional Neural Network

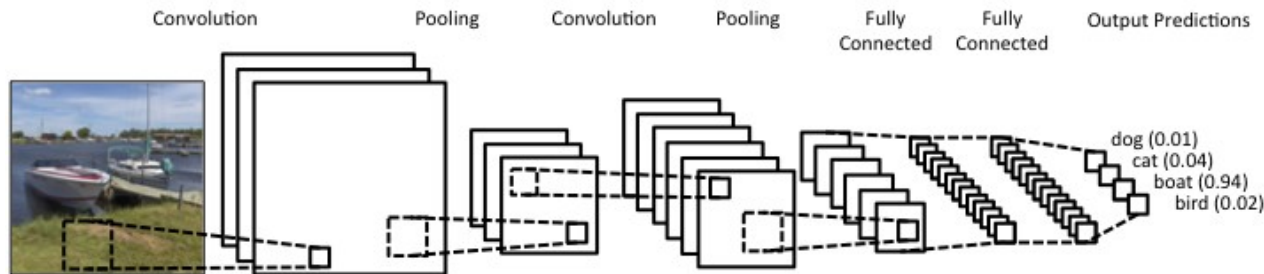
1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

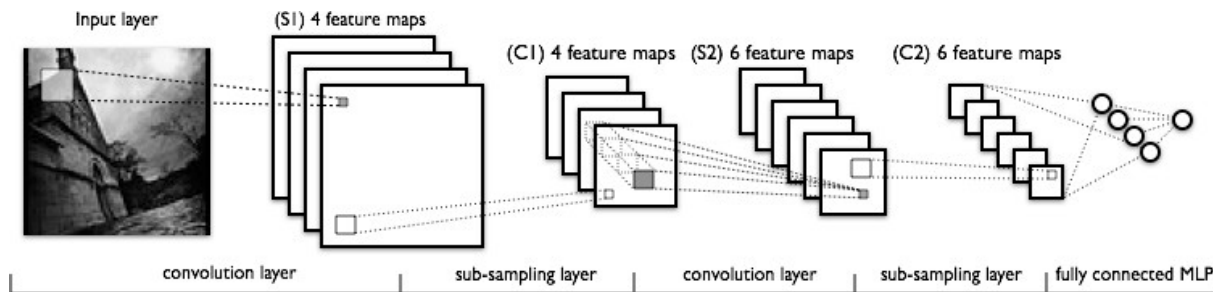
4		

Convolved  
Feature

# Convolutional Neural Network



source: <http://clarifai.com/technology>



source: <http://deeplearning.net/tutorial/lenet.html>

# Parameters for weight update rules

**Learning rate:** parameter that determines how much an updating step influences the current value of the weights.

**Weight decay:** an additional term in the weight update rule that causes the weights to exponentially decay to zero, if no other update is scheduled. Similar to the regularization term in SVM to reduce overfitting.

**Momentum:** To smooth out the gradients over iterations, each update to make is averaged with the weight update we made in the previous iteration. Momentum controls how strong the influence of the previous update.

SGD+L2: 
$$\begin{aligned} \text{update\_history} &\leftarrow \text{momentum} * \text{update\_history} - \text{learning\_rate} * (\text{gradient} + \text{weight\_decay} * \text{weight}) \\ \text{weight} &\leftarrow \text{weight} - \text{update\_history} \end{aligned}$$

A **Layer** is an **operation** on **responses**:  $\mathbf{y} = f(\mathbf{x})$

in's and out's

in-place operation,  $[\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_m] = f(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n)$

allow sharing in's with other backprop layer

Parallel in's and out's  $\mathbf{x} = f(\mathbf{x})$

$$\mathbf{y}_1 = f_1(\mathbf{x}) \quad \mathbf{y}_2 = f_2(\mathbf{x})$$

$$\mathbf{y}_1 = f(\mathbf{x}_1, \mathbf{x}_2) \quad \mathbf{y}_2 = f(\mathbf{x}_3, \mathbf{x}_4) \quad \mathbf{y}_3 = f(\mathbf{x}_5, \mathbf{x}_6)$$

# Layer: Convolution

Common:

**name:**

**phase:**

“Training”, “Testing”, “**TrainingTesting**”

**train\_me:** **true** / false

**Convolution filters:**

**num\_output:** number of filters

**window:** e.g., [5,5,5] for 3D

**padding:** [0,0,...]

**stride:** [1,1,...]

**upscale:** [1,1,...]

**group:** 1

Weight update and initialization:

**weight\_filler:**

“**Xavier**”, “Gaussian”, “Constant”

**weight\_filler\_param:** 0.0

**weight\_lr\_mult:** 1.0

**weight\_decay\_mult:** 1.0

**bias\_filler:**

“Xavier”, “Gaussian”, “**Constant**”

**bias\_filler\_param:** 0.0

**bias\_lr\_mult:** 2.0

**bias\_decay\_mult:** 1.0

# Layer: Convolution

Common:

name:

phase:

“Training”, “Testing”, “TrainingTesting”

train\_me: true / false

Convolution filters:

num\_output: number of filters

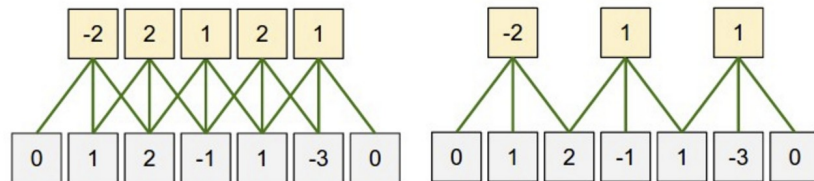
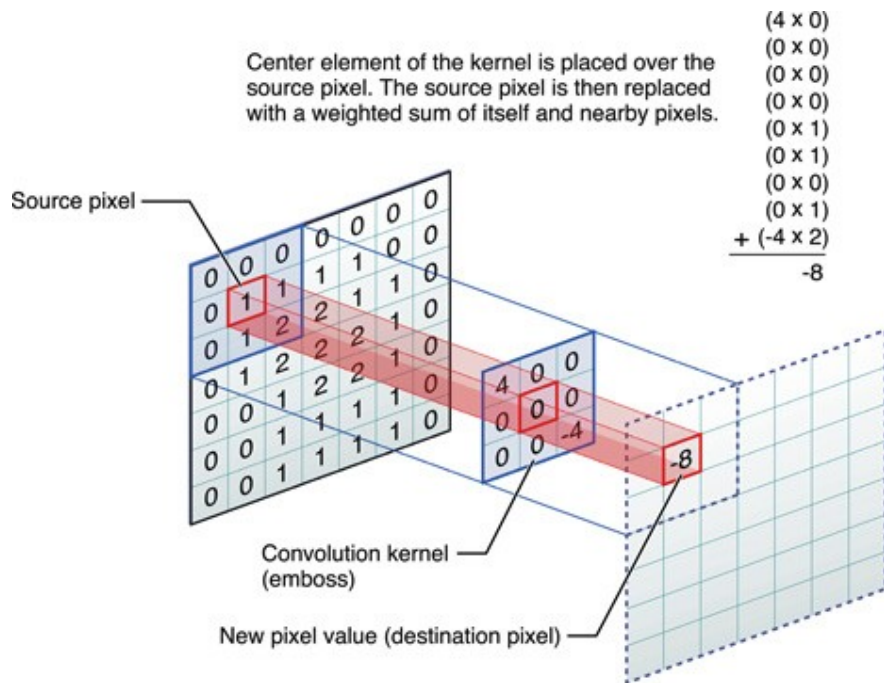
window: e.g., [5,5,5] for 3D

padding: [0,0,...]

stride: [1,1,...]

upscale: [1,1,...]

group: 1



# Layer: Convolution

Common:

name:

phase:

“Training”, “Testing”, “TrainingTesting”

train\_me: true / false

Convolution filters:

num\_output: number of filters

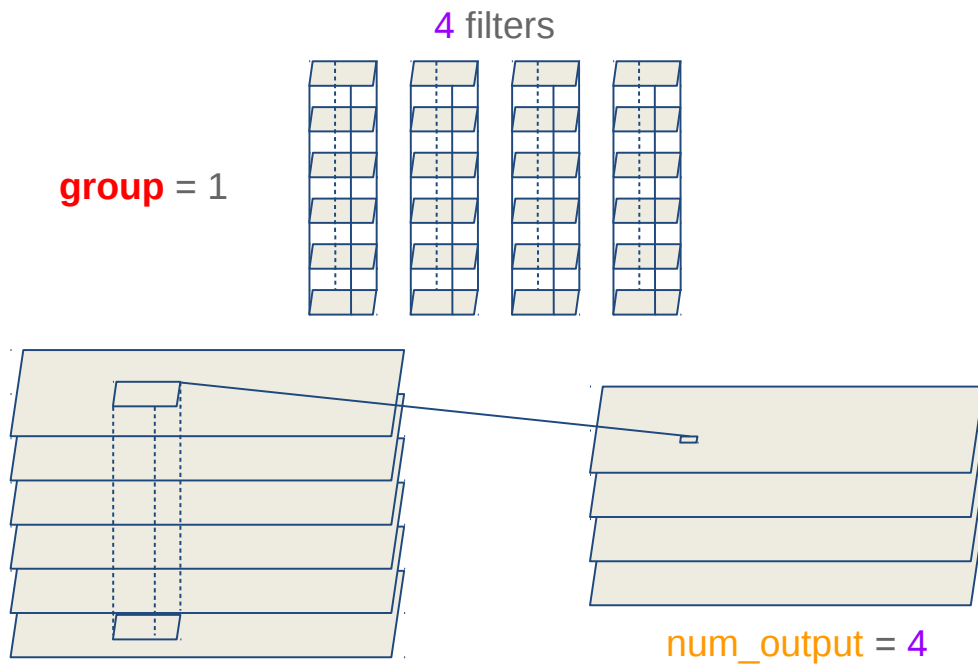
window: e.g., [5,5,5] for 3D

padding: [0,0,...]

stride: [1,1,...]

upscale: [1,1,...]

group: 1





# Layer: Convolution

Common:

name:

phase:

"Training", "Testing", "TrainingTesting"

train\_me: true / false

Convolution filters:

num\_output: number of filters

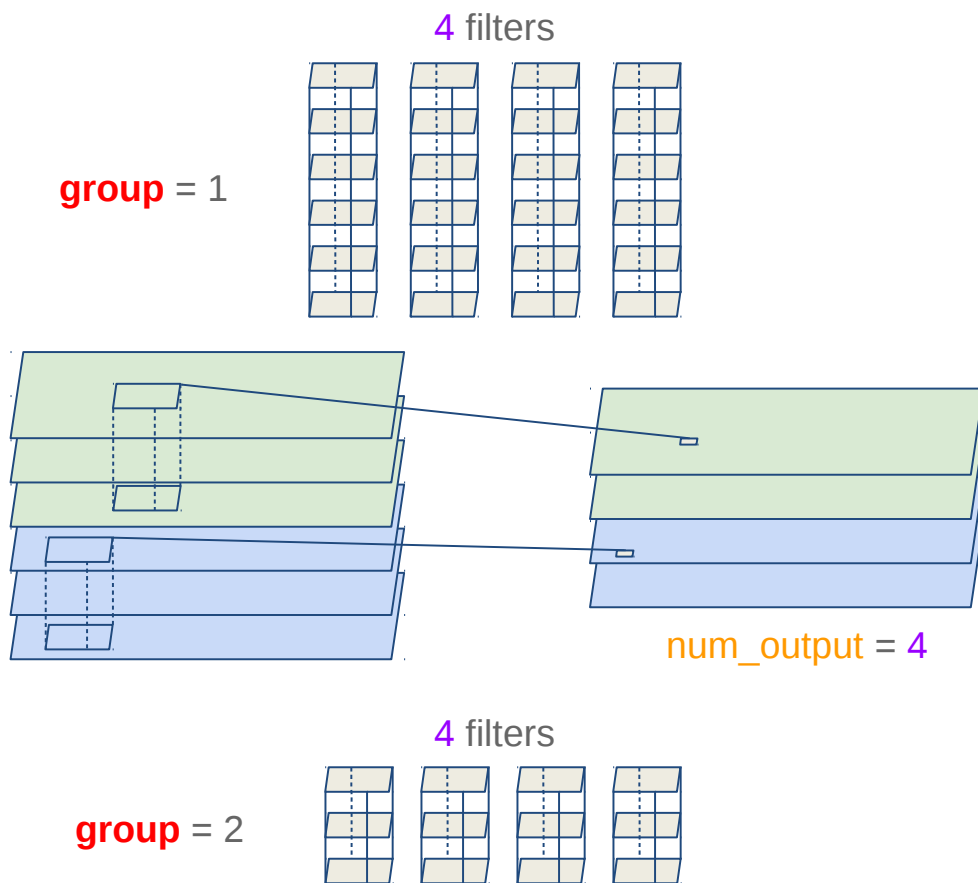
window: e.g., [5,5,5] for 3D

padding: [0,0,...]

stride: [1,1,...]

upscale: [1,1,...]

group: 1



# Layer: InnerProduct (a.k.a. Fully Connected)

Common:

name:

phase:

"Training", "Testing", "TrainingTesting"

train\_me: true / false

Parameters:

num\_output: number of output of the layer

Weight update and initialization:

weight\_filler:

"Xavier", "Gaussian", "Constant"

weight\_filler\_param: 0.0

weight\_lr\_mult: 1.0

weight\_decay\_mult: 1.0

bias\_filler:

"Xavier", "Gaussian", "Constant"

bias\_filler\_param: 0.0

bias\_lr\_mult: 2.0

bias\_decay\_mult: 1.0

$$\mathbf{y} = \mathbf{W}\mathbf{x} + \mathbf{b}$$

size: num\_output × sizeofitem(**X**)

size: num\_output



# Layer: Activation

Common:

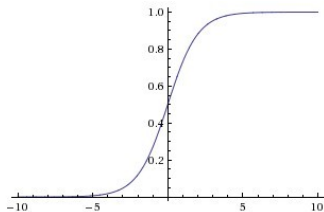
name:

phase:

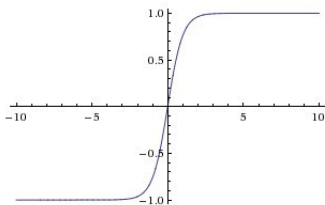
“Training”, “Testing”, “TrainingTesting”

mode:

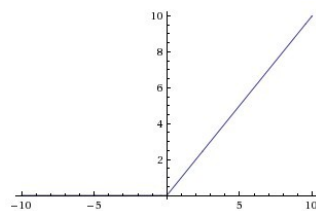
“Sigmoid”, “ReLU”, “TanH”



Sigmoid

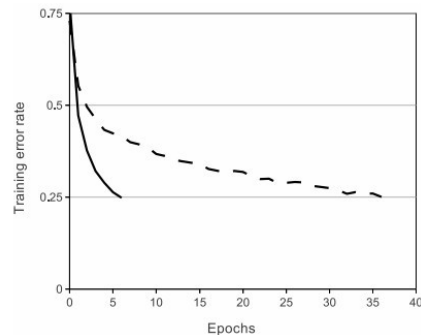


TanH



ReLU

A plot from Krizhevsky *et al.* indicating the 6x improvement in convergence with the ReLU unit compared to the tanh unit.



# Layer: Dropout

Dropout helps to avoid overfitting by randomly drop units (along with their connections) from the network during training to prevent units from co-adapting too much.

During training, dropout layer randomly set some part of an in response to be zeros in the forward based on the `dropout_rate`. In backward, the gradient is multiply by  $1/\text{dropout\_rate}$ .

During testing, the dropout layer does nothing.

**name:**

**phase:**

"Training", "Testing", "TrainingTesting"

**dropout\_rate:** 0.5

# Layer: Softmax

name:

phase:

“Training”, “Testing”, “TrainingTesting”

stable\_gradient: true / false

Strong Recommendation:

For numerical stability, it is highly recommend to set stable\_gradient = true, and use “MultinomialLogistic\_StableSoftmax” for the mode in the loss layer.

For more details about the reasons behind, refer to

<http://freemind.pluskid.org/machine-learning/softmax-vs-softmax-loss-numerical-stability/>

# Layer: Loss

name:

phase:

“Training”, “Testing”, “TrainingTesting”

mode:

“MultinomialLogistic\_StableSoftmax”, “MultinomialLogistic”, “SmoothL1”, “Contrastive”, “EuclideanSSE”,  
“HingeL1”, “HingeL2”, “SigmoidCrossEntropy”, “Infogain”

loss\_weight: 1.0

margin: 1.0

loss\_weights: an array of real numbers for weight each channels

in:

the first in is the network prediction

the second in is the ground truth label

optionally, the third in is a weight array (for the first three mode only)

# Layer: Pooling

Common:

name:

phase:

“Training”, “Testing”, “TrainingTesting”

Parameters:

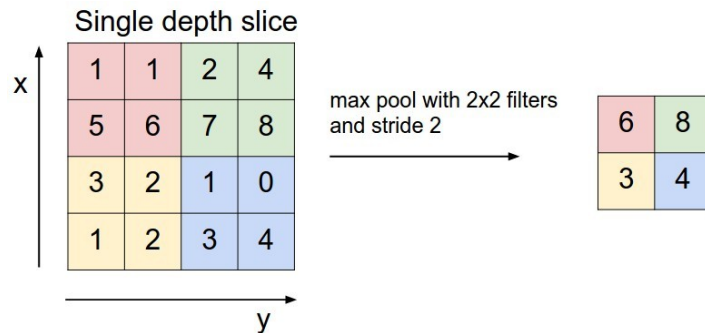
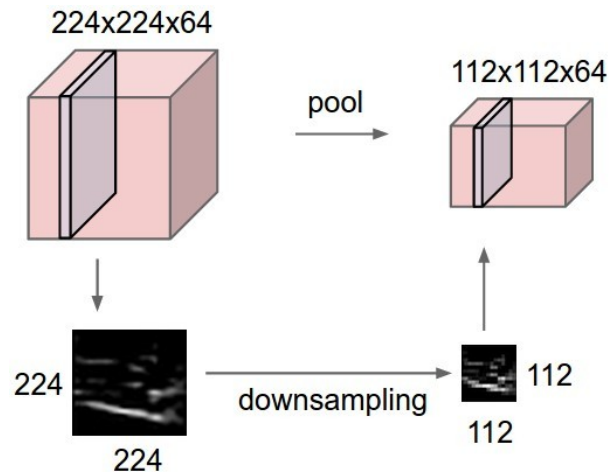
mode:

“max”, “average\_include”, “average\_exclude”

window: e.g., [5,5,5] for 3D

padding: [0,0,...]

stride: [1,1,...]



# Layer: Concat

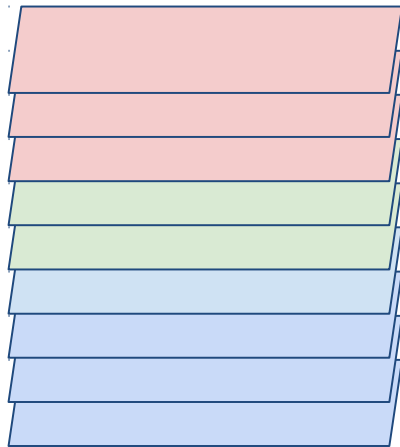
Concatenate inputs in channels:

name:

phase:

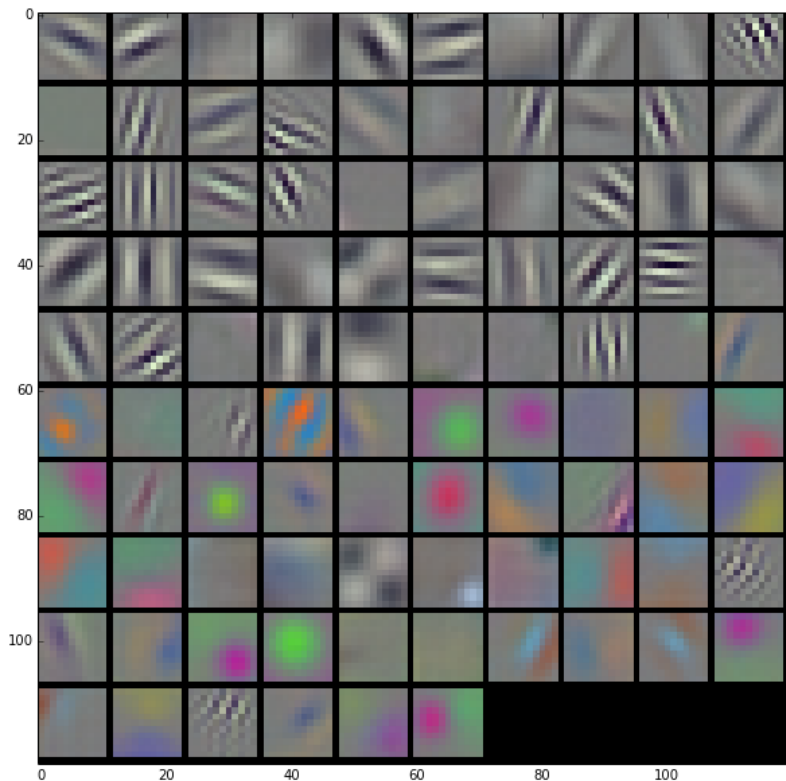
“Training”, “Testing”, “TrainingTesting”

Let's say that there are  $N$  in's and  $M$  out's. It will divide the  $N$  in's into  $M$  groups. Each of the  $M$  groups of ins will be concatenated to generate  $M$  outs in total.

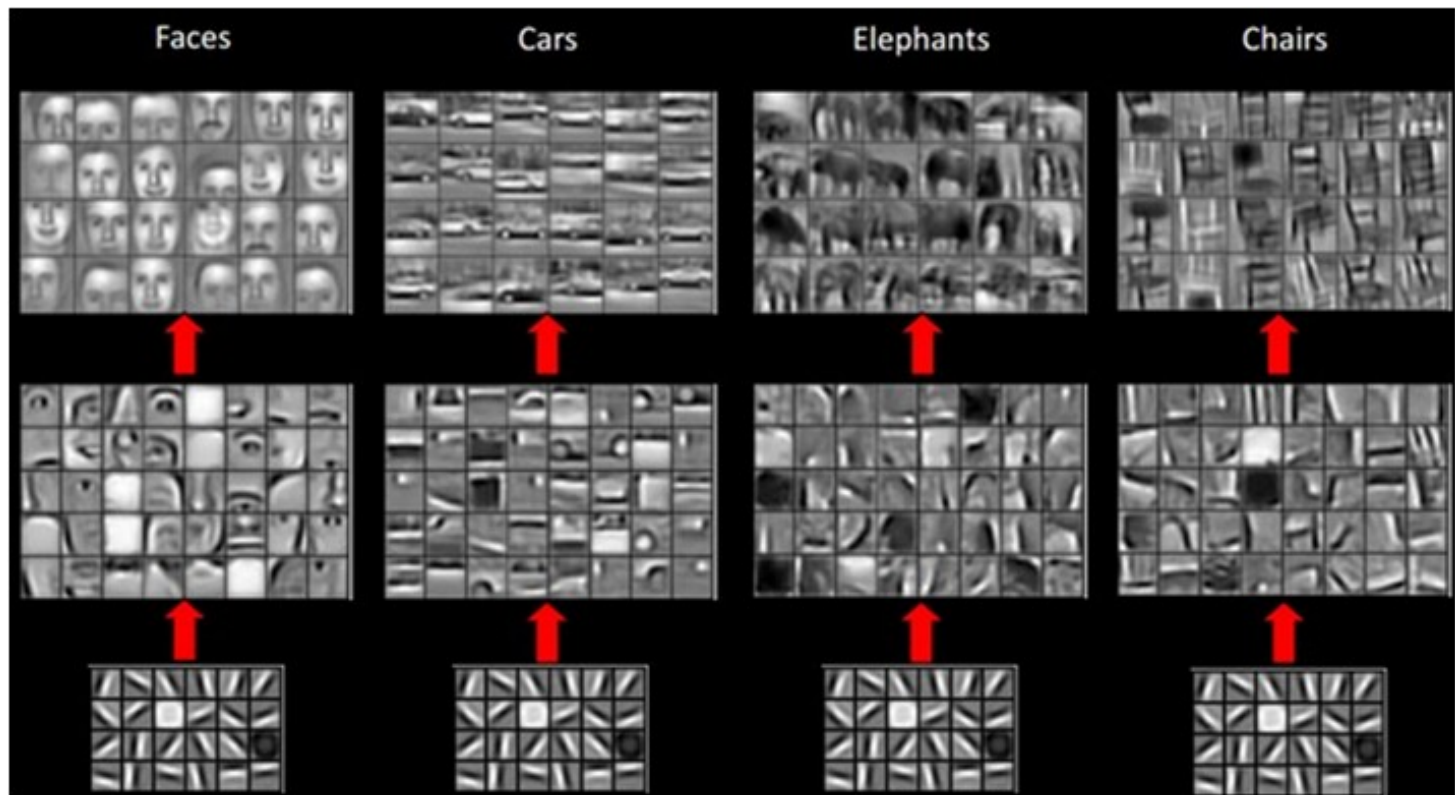




# Deeper into deep learning: Hierarchy is key



# Deeper into deep learning: Hierarchy is key



# Deeper into deep learning: Why CNNs ?

Neural Networks were introduced during the 1960s.

Through the 80s, they contributed to the AI winter

(a lot of expectations and claims, little delivery)

Mainly because data and computing power were not available.

Today, we have both

(but we still need better learning algorithms)



# Deeper into deep learning: Why CNNs ?

Neural Networks vastly overfitted in the 90s with little data.

Yann LeCun ('89) proposed weight sharing within each layer.

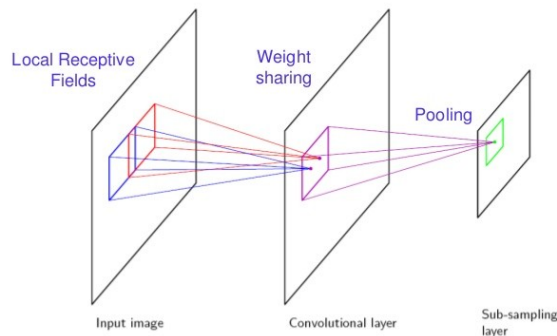
Reduced the number of learnable parameters by a lot.

Helps reduce overfitting, why ?

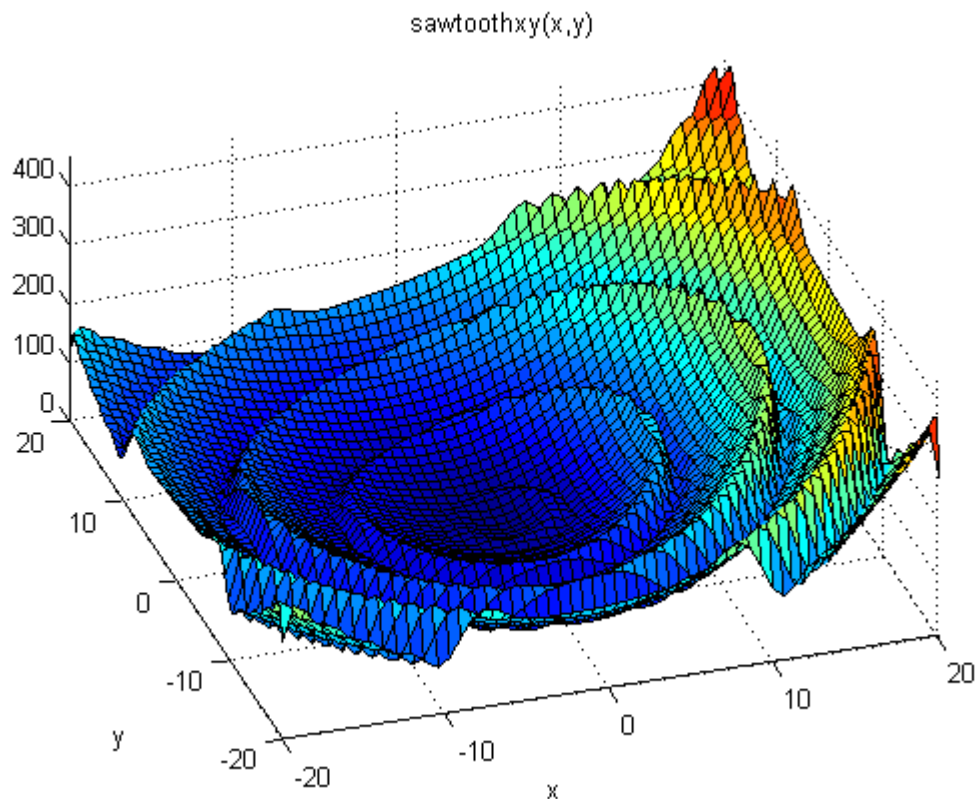
Weight sharing effectively is convolution.

## Convolutional Neural Networks

(LeCun et al., 1989)



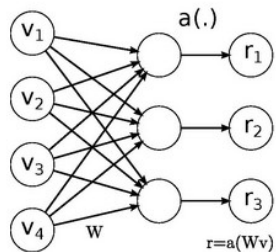
# Deeper into deep learning: Local Minima



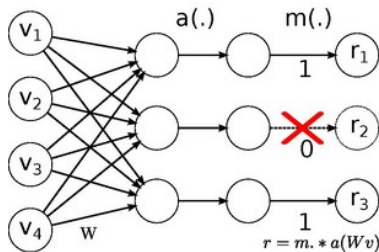
# Deeper into deep learning: Dropout

Drop out helps the network generalize better. One of the most successful recent techniques since scaling up.

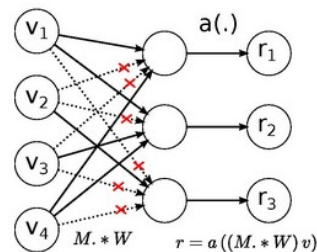
Recall, during training it randomly puts activations to zero (proportional to drop out rate). At test time, it does nothing.



No-Drop Network



DropOut Network

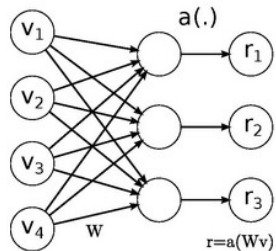


DropConnect Network

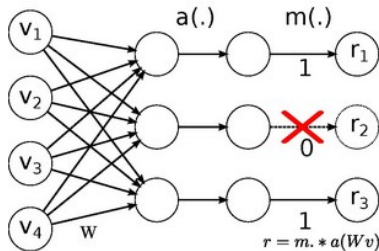
# Deeper into deep learning: Dropout

Randomly “inhibiting” the network keeps it from falling into a premature local minima.

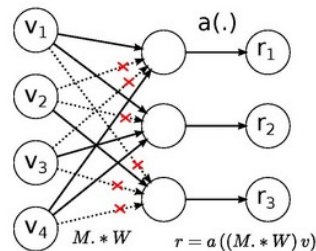
Randomization might help in other things ! (Dropout/DropConnect)



No-Drop Network



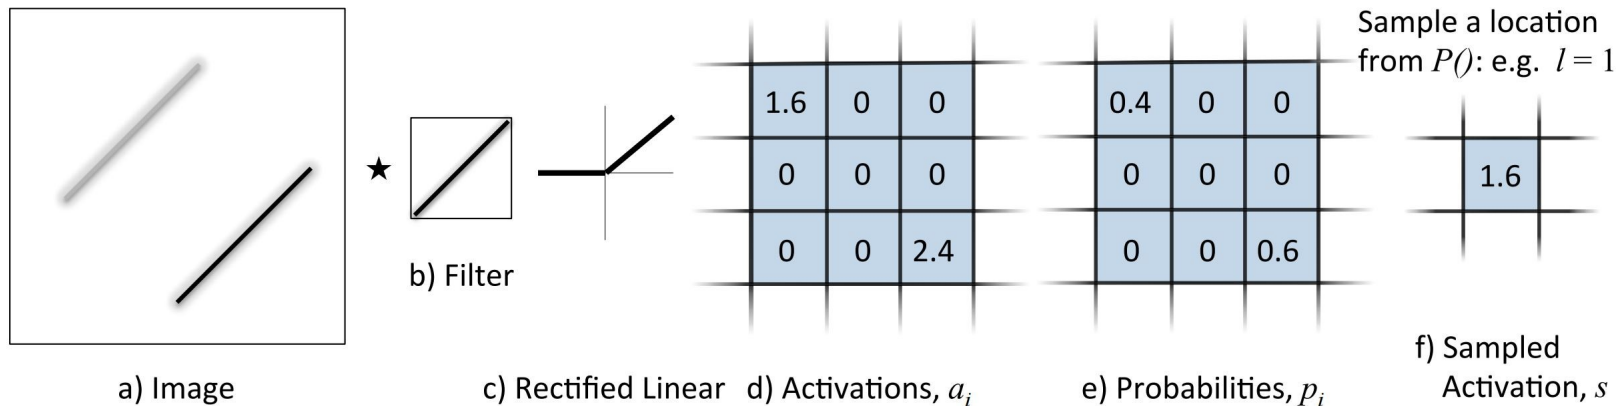
DropOut Network



DropConnect Network

# Deeper into deep learning: Stochastic Pooling

Matt Zeiler and Rob Fergus proposed stochastic pooling in 2013.





# Deeper into deep learning: Stochastic Pooling

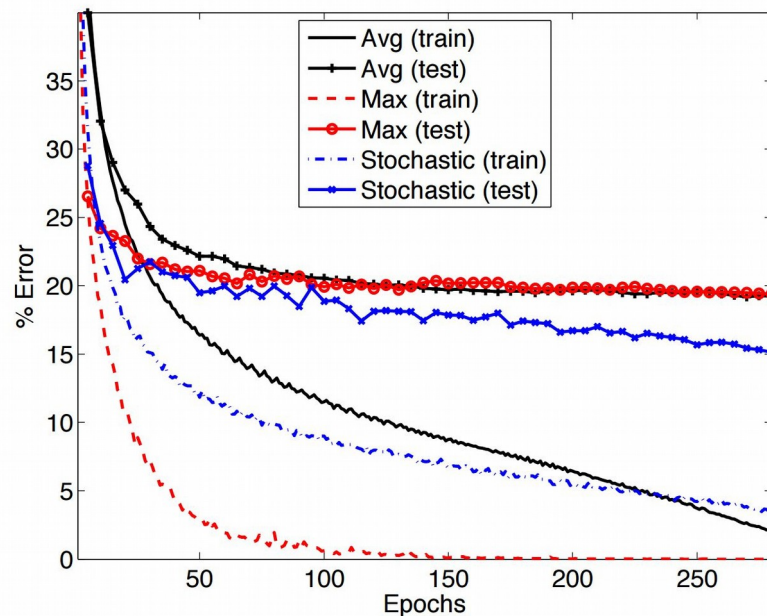
CIFAR-10 results

Stochastic generalizes the best

Max and mean pooling work similarly

Max generalizes better earlier

Randomization clearly helps !



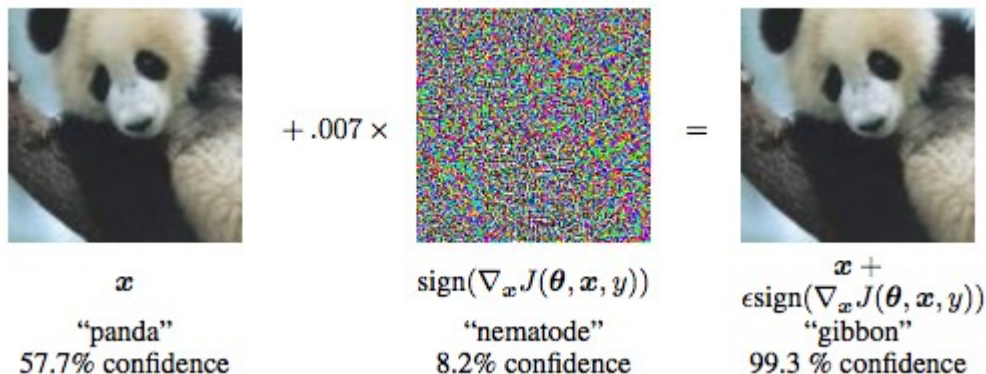
# Deeper into deep learning: Perception Solved ?

CNNs have pushed the state-of-the-art for most vision tasks.

However, they have limitations.

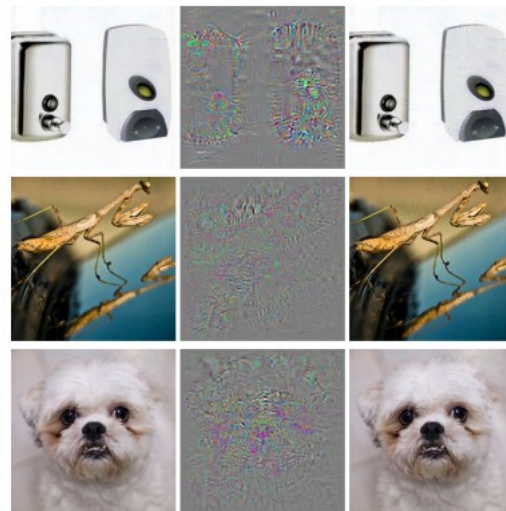
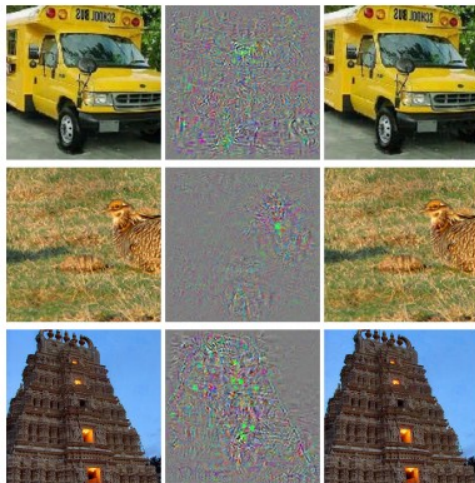
The mapping from images to the label space is not “stable”.

Specifically, a small carefully chosen distortion (imperceptible to the human eye) can make the CNN misclassify deterministically.



# Deeper into deep learning: Perception Solved ?

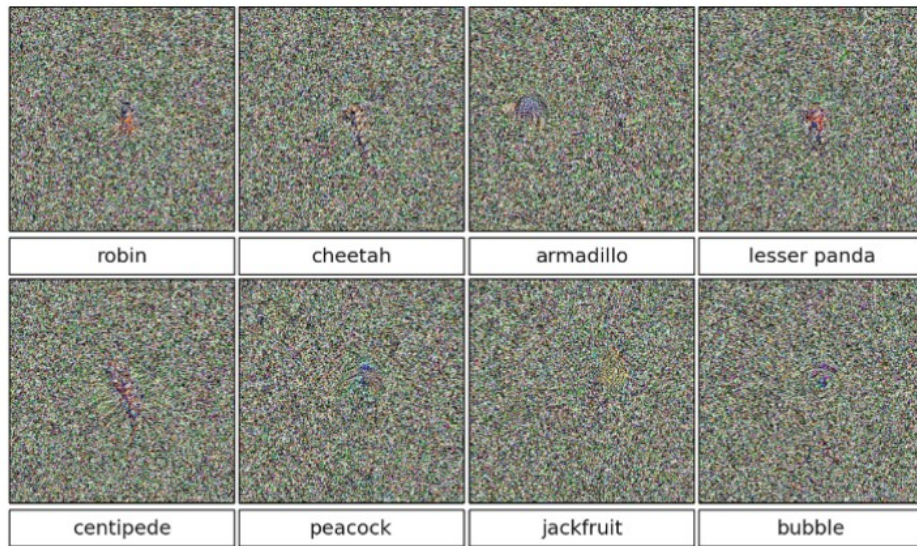
Left column original image, center is the difference image, right column classified as Ostrich (for both image blocks).



“Intriguing properties of neural networks”, Szegedy et. al., 2014

# Deeper into deep learning: Perception Solved ?

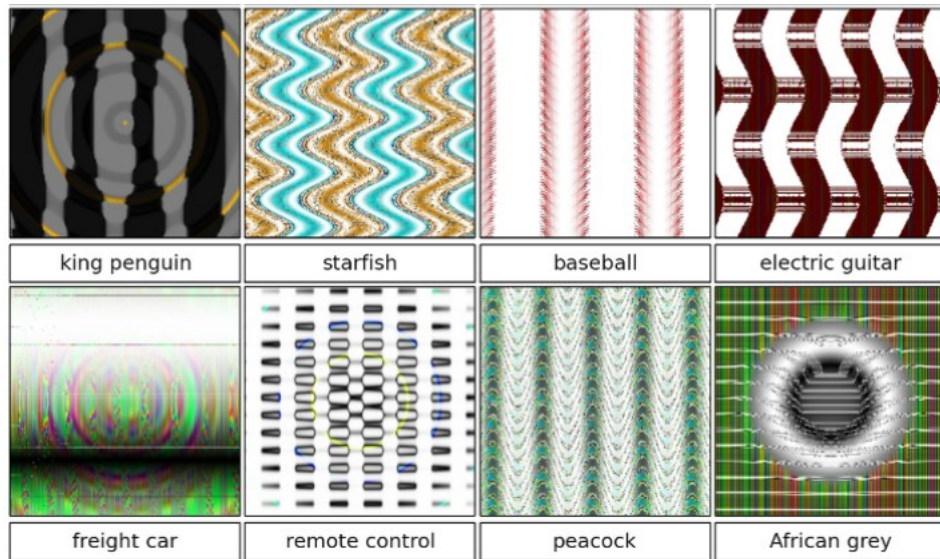
More images can be evolved (genetic algorithms) or found **directly (back propagation)** that fool CNNs.



“Deep neural networks are easily fooled: High confidence predictions for unrecognizable images”, Nguyen et. al., CVPR 2015

# Deeper into deep learning: Perception Solved ?

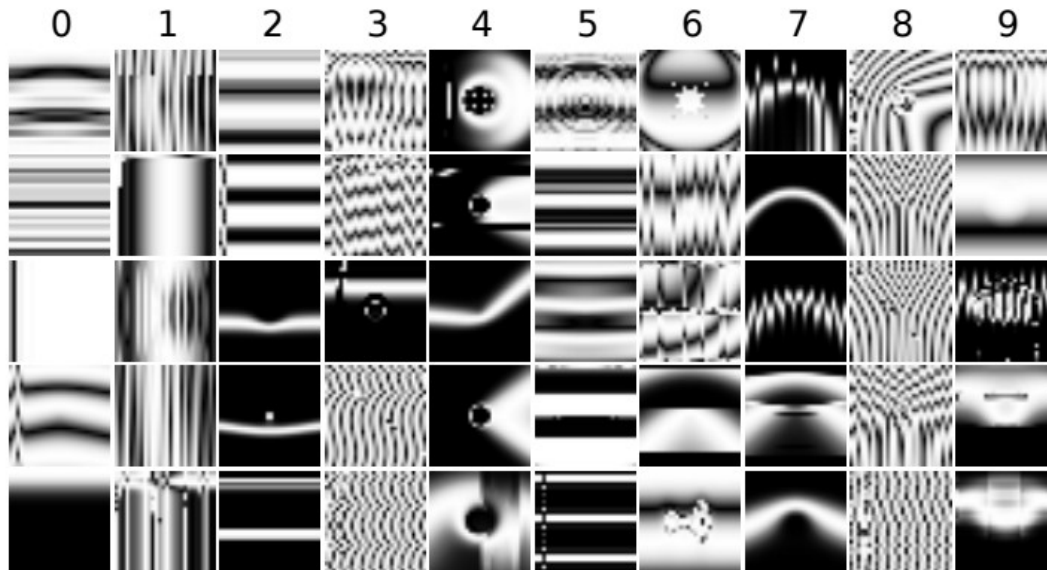
More images can be **evolved** (genetic algorithms) or found directly (back propagation) that fool CNNs.



“Deep neural networks are easily fooled: High confidence predictions for unrecognizable images”, Nguyen et. al., CVPR 2015

# Deeper into deep learning: Perception Solved ?

Images that appear to the CNN (LeNet) to be from MNIST.



“Deep neural networks are easily fooled: High confidence predictions for unrecognizable images”, Nguyen et. al., CVPR 2015

# Applications of CNNs

Video Classification

Image Segmentation

Object Detection

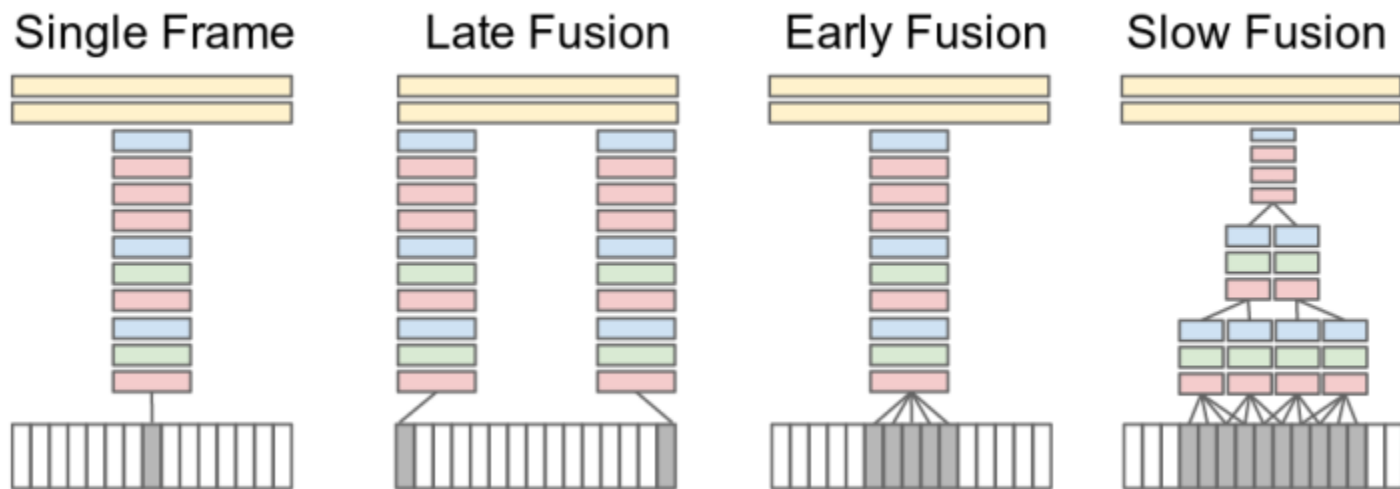
Face Landmarking

Face Recognition

Face Detection

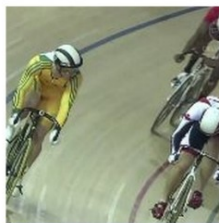


# Applications of CNNs: Video Classification





# Applications of CNNs: Video Classification



track cycling  
cycling  
track cycling  
road bicycle racing  
marathon  
ultramarathon



ultramarathon  
ultramarathon  
half marathon  
running  
marathon  
inline speed skating



heptathlon  
heptathlon  
decathlon  
hurdles  
pentathlon  
sprint (running)



bikejoring  
mushing  
bikejoring  
harness racing  
skijoring  
carting



longboarding  
longboarding  
aggressive inline skating  
freestyle scootering  
freboard (skateboard)  
sandboarding



ultimate (sport)  
ultimate (sport)  
hurling  
flag football  
association football  
rugby sevens



demolition derby  
demolition derby  
monster truck  
mud bogging  
motocross  
grand prix motorcycle racing



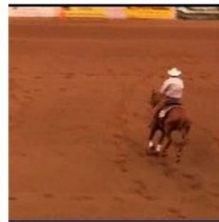
telemark skiing  
snowboarding  
telemark skiing  
nordic skiing  
ski touring  
skijoring



whitewater kayaking  
whitewater kayaking  
rafting  
kayaking  
canoeing  
adventure racing



arena football  
indoor american football  
arena football  
canadian football  
american football  
women's lacrosse

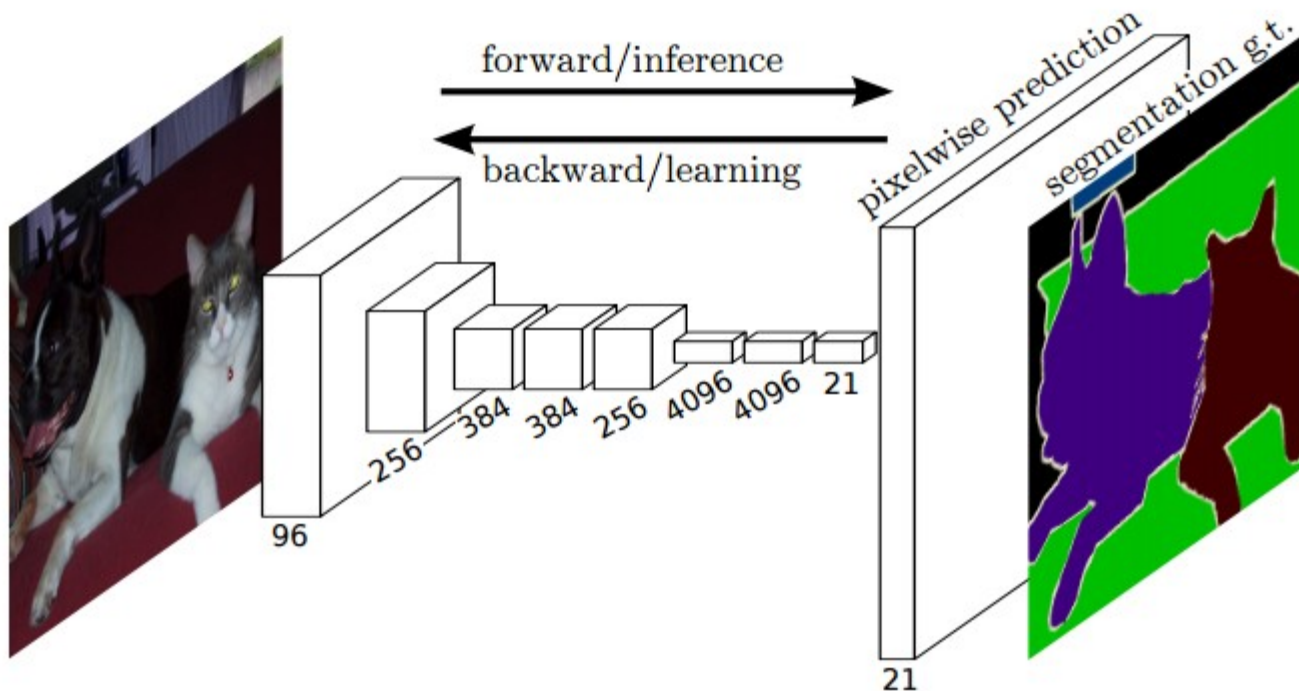


reining  
barrel racing  
rodeo  
reining  
cowboy action shooting  
bull riding

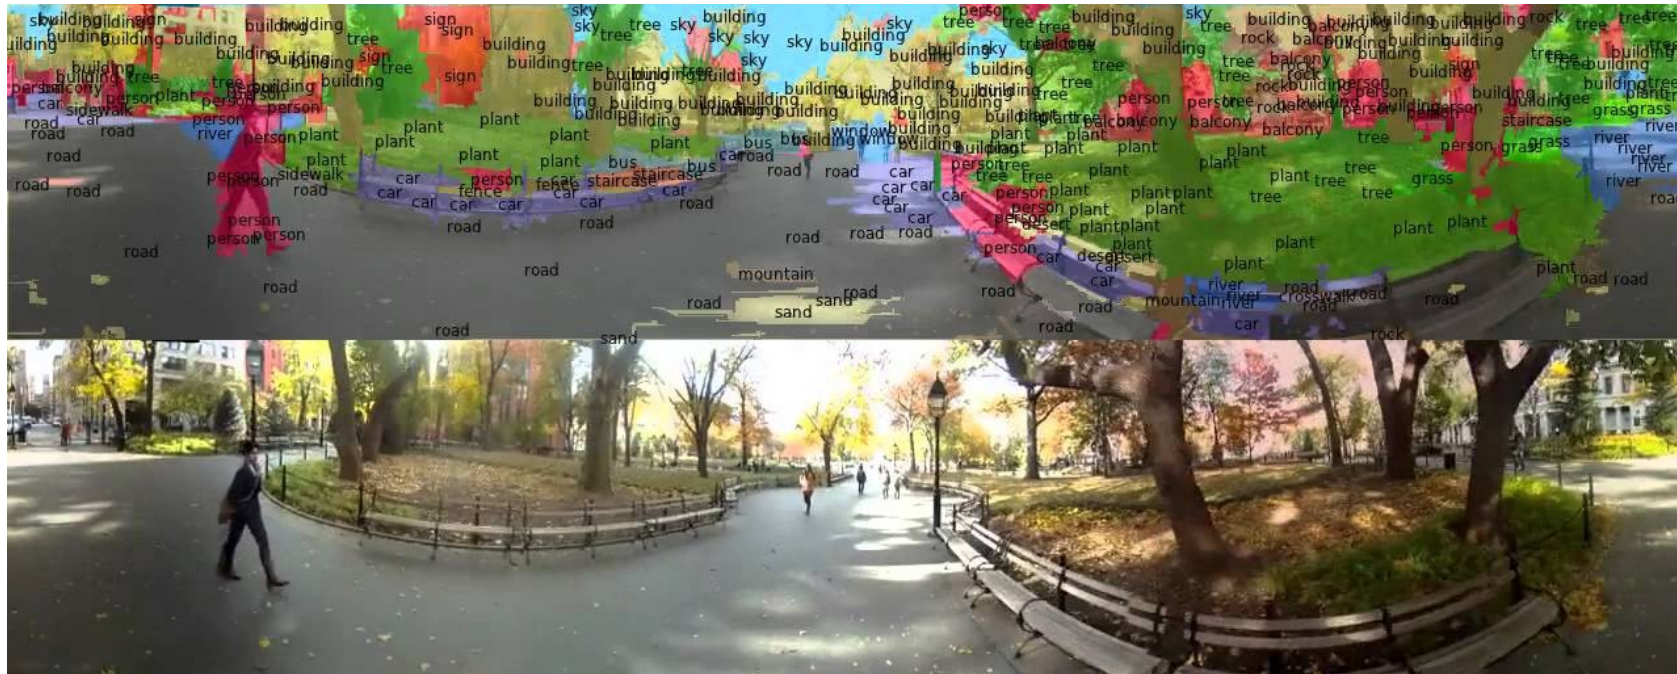


eight-ball  
nine-ball  
blackball (pool)  
trick shot  
eight-ball  
straight pool

# Applications of CNNs: Image Segmentation



# Applications of CNNs: Image Segmentation



# Applications of CNNs: Face Recognition

